Follow-up data with the `Lexis` functions in `Epi`

Bendix Carstensen   Steno Diabetes Center Copenhagen, Gentofte, Denmark
& Department of Biostatistics, University of Copenhagen
`b@bxc.dk`
<http://BendixCarstensen.com>

Contents

Chapter 1

Introduction

This is an introduction to the `Lexis` machinery in the `Epi` package. It is intended for representation of follow-up data (event history data) from studies where exact dates of events are known. It accommodates follow-up through multiple states and on multiple time scales.

This vignette uses an example from the *Epi* package to illustrate the set-up of a simple `Lexis` object (a data frame of follow-up intervals), as well as the subdivision of follow-up intervals needed for multistate representation and analysis of transition rates.

The first chapter is exclusively on manipulation of the follow-up representation, but it points to the subsequent chapter where analysis is based on a `Lexis` representation with very small follow-up intervals.

The next chapter demonstrates analysis of mortality rates among Danish diabetes patients currently on insulin treatment or not.

Chapter 2

Representation of follow-up data in the `Epi` package

In the `Epi`-package, follow-up data is represented by adding some extra variables to a data frame. Such a data frame is called a `Lexis` object. The tools for handling follow-up data then use the structure of this for special plots, tabulations and modeling.

Follow-up data basically consists of a time of entry, a time of exit and an indication of the status at exit (normally either "alive" or "dead") for each person. Implicitly is also assumed a status *during* the follow-up (usually "alive").
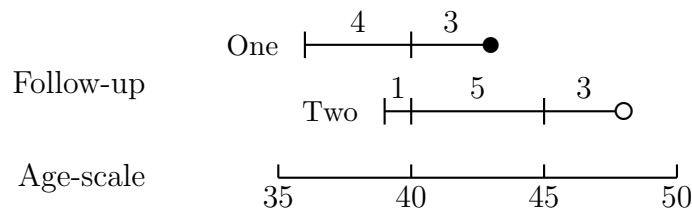
Figure 2.1: *Follow-up of two persons*

2.1 Timescales

A timescale is a variable that varies deterministicly *within* each person during follow-up, *e.g.*:

- Age

- Calendar time

- Time since start of treatment

- Time since relapse

All timescales advance at the same pace, so the time followed is the same on all timescales. Therefore, it will suffice to use only the entry point on each of the time scales, for example:

- Age at entry

- Date of entry

- Time at treatment (*at* treatment this is 0)

- Time at relapse (*at* relapse this is 0)

For illustration we need to load the `Epi` package:

```
> library(Epi)
> print( sessionInfo(), l=F )
R version 3.5.2 (2018-12-20)
Platform: x86_64-pc-linux-gnu (64-bit)
Running under: Ubuntu 14.04.5 LTS

Matrix products: default
BLAS: /usr/lib/openblas-base/libopenblas.so.0
LAPACK: /usr/lib/lapack/liblapack.so.3.0

attached base packages:
[1] utils     datasets  graphics  grDevices stats     methods   base

other attached packages:
[1] Epi_2.35

loaded via a namespace (and not attached):
 [1] Rcpp_1.0.0        lattice_0.20-38   zoo_1.8-4         MASS_7.3-51.1
 [5] grid_3.5.2        plyr_1.8.4        nlme_3.1-137      etm_1.0.4
 [9] data.table_1.12.0 Matrix_1.2-15     splines_3.5.2     tools_3.5.2
[13] cmprsk_2.2-7      numDeriv_2016.8-1 survival_2.43-3   parallel_3.5.2
[17] compiler_3.5.2    mgcv_1.8-27
```

In the `Epi` package, follow-up in a cohort is represented in a `Lexis` object. As mentioned, a `Lexis` object is a data frame with some extra structure representing the follow-up. For the `DMlate` data — follow-up of diabetes patients in Denmark recording date of birth, date of diabetes, date of insulin use, date of first oral drug use and date of death — we can construct a `Lexis` object by:

```
> data( DMlate )
> dmL <- Lexis( entry = list( per=dodm,
+                             age=dodm-dobth,
+                             tfD=0 ),
+                exit = list( per=dox ),
+         exit.status = factor( !is.na(dodth), labels=c("DM","Dead") ),
+                data = DMlate )
NOTE: entry.status has been set to "DM" for all.
> timeScales(dmL)
```

```
[1] "per" "age" "tfD"
```

(The excluded persons are persons with date of diabetes equal to date of death.)

The `entry` argument is a *named* list with the entry points on each of the timescales we want to use. It defines the names of the timescales and the entry points of the follow-up of each person. The `exit` argument gives the exit time on *one* of the timescales, so the name of the element in this list must match one of the names of the `entry` list. This is sufficient, because the follow-up time on all time scales is the same, in this case `dox-dodm`.

The `exit.status` is a categorical variable (a *factor*) that indicates the exit status — in this case whether the person (still) is in state `DM` or exits to `Dead` at the end of follow-up.

In principle we should also indicate the `entry.status`, but the default is to assume that all persons enter in the `first` of the mentioned `exit.state`s — in this case `DM`, because `FALSE < TRUE`.

Now take a look at the result:

```
> str( DMlate )
'data.frame':         10000 obs. of  7 variables:
 $ sex  : Factor w/ 2 levels "M","F": 2 1 2 2 1 2 1 1 2 1 ...
 $ dobth: num  1940 1939 1918 1965 1933 ...
 $ dodm : num  1999 2003 2005 2009 2009 ...
 $ dodth: num  NA NA NA NA NA ...
 $ dooad: num  NA 2007 NA NA NA ...
 $ doins: num  NA NA NA NA NA NA NA NA NA NA ...
 $ dox  : num  2010 2010 2010 2010 2010 ...
> str( dmL )
Classes 'Lexis' and 'data.frame':          9996 obs. of  14 variables:
 $ per    : num  1999 2003 2005 2009 2009 ...
 $ age    : num  58.7 64.1 86.3 44 75.8 ...
 $ tfD    : num  0 0 0 0 0 0 0 0 0 0 ...
 $ lex.dur: num  11.08 6.689 5.446 0.736 1.344 ...
 $ lex.Cst: Factor w/ 2 levels "DM","Dead": 1 1 1 1 1 1 1 1 1 1 ...
 $ lex.Xst: Factor w/ 2 levels "DM","Dead": 1 1 1 1 1 2 1 1 2 1 ...
 $ lex.id : int  1 2 3 4 5 6 7 8 9 10 ...
 $ sex    : Factor w/ 2 levels "M","F": 2 1 2 2 1 2 1 1 2 1 ...
 $ dobth  : num  1940 1939 1918 1965 1933 ...
 $ dodm   : num  1999 2003 2005 2009 2009 ...
 $ dodth  : num  NA NA NA NA NA ...
 $ dooad  : num  NA 2007 NA NA NA ...
 $ doins  : num  NA NA NA NA NA NA NA NA NA NA ...
 $ dox    : num  2010 2010 2010 2010 2010 ...
 - attr(*, "time.scales")= chr  "per" "age" "tfD"
 - attr(*, "time.since")= chr  "" "" ""
 - attr(*, "breaks")=List of 3
  ..$ per: NULL
  ..$ age: NULL
  ..$ tfD: NULL
> head( dmL )[,1:10]
             per       age tfD    lex.dur lex.Cst lex.Xst lex.id sex      dobth      dodm
50185   1998.917 58.66119   0 11.0800821      DM      DM       1   F 1940.256 1998.917
307563 2003.309 64.09035   0  6.6885695      DM      DM       2   M 1939.218 2003.309
294104 2004.552 86.25051   0  5.4455852      DM      DM       3   F 1918.301 2004.552
336439 2009.261 44.03559   0  0.7364819      DM      DM       4   F 1965.225 2009.261
245651 2008.653 75.77550   0  1.3442847      DM      DM       5   M 1932.877 2008.653
216824 2007.886 80.01643   0  2.0369610      DM    Dead       6   F 1927.870 2007.886
```

The `Lexis` object `dmL` has a variable for each timescale which is the entry point on this timescale. The follow-up time is in the variable `lex.dur` (duration). Note that the exit status is in the variable `lex.Xst` (eXit state. The variable `lex.Cst` is the state where the follow-up takes place (Current state), in this case `DM` (alive with diabetes) for all persons. This implies that *censored* observations are characterized by having `lex.Cst = lex.Xst`.

There is a `summary` function for `Lexis` objects that lists the number of transitions and records as well as the total amount of follow-up time; it also (optionally) prints information about the names of the variables that constitute the timescales:

```
> summary.Lexis( dmL, timeScales=TRUE )
Transitions:
     To
From   DM Dead  Records:  Events: Risk time:  Persons:
  DM 7497 2499      9996     2499   54273.27      9996

Timescales:
per age tfD
 ""  ""  ""
```

It is possible to get a visualization of the follow-up along the timescales chosen by using the **plot** method for **Lexis** objects. **dmL** is an object of *class* **Lexis**, so using the function **plot()** on it means that R will look for the function **plot.Lexis** and use this function.

```
> plot( dmL )
```

The function allows quite a bit of control over the output, and a **points.Lexis** function allows plotting of the endpoints of follow-up:

```
> par( mar=c(3,3,1,1), mgp=c(3,1,0)/1.6 )
> plot( dmL, 1:2, lwd=1, col=c("blue","red")[dmL$sex],
+       grid=TRUE, lty.grid=1, col.grid=gray(0.7),
+       xlim=1960+c(0,60), xaxs="i",
+       ylim=  40+c(0,60), yaxs="i", las=1 )
> points( dmL, 1:2, pch=c(NA,3)[dmL$lex.Xst],
+         col="lightgray", lwd=3, cex=0.3 )
> points( dmL, 1:2, pch=c(NA,3)[dmL$lex.Xst],
+         col=c("blue","red")[dmL$sex], lwd=1, cex=0.3 )
> box(bty='o')
```

In the above code you will note that the values of the arguments **col** and **pch** are indexed by factors, using the convention R that the index is taken as *number of the level* of the supplied factor. Thus **c("blue","red")[dmL$sex]** is **"blue"** when **sex** is **M** (the first level).

The results of these two plotting commands are in figure 2.2.

2.2   Splitting the follow-up time along a timescale

In next chapter we shall conduct statistical analysis of mortality rates, and a prerequisite for parametric analysis of rates is that follow-up time is subdivided in smaller intervals, where we can reasonably assume that rates are constant.

The follow-up time in a cohort can be subdivided ("split") along a time scale, for example current age. This is achieved by the **splitLexis** (note that it is *not* called **split.Lexis**). This requires that the timescale and the breakpoints on this timescale are supplied. Try:

```
> dmS1 <- splitLexis( dmL, "age", breaks=seq(0,100,5) )
> summary( dmL )
Transitions:
     To
From   DM Dead  Records:  Events: Risk time:  Persons:
  DM 7497 2499      9996     2499   54273.27      9996

> summary( dmS1 )
```

Figure 2.2: *Lexis diagram of the* `DMlate` *dataset; left panel is the default version, right panel: plot with some bells and whistles. The red lines are for women, blue for men, crosses indicate deaths.*

```
Transitions:
     To
From    DM Dead  Records:  Events: Risk time:  Persons:
  DM 18328 2499     20827     2499    54273.27      9996
```

We see that the number of persons and events and the amount of follow-up is the same in the two data sets; only the number of records differ — the extra records all have `lex.Cst=DM` and `lex.Xst=DM`.

To see how records are split for each individual, it is useful to list the results for a few individuals (whom we selected with a view to the illustrative usefulness):

```
> wh.id <- c(9,27,52,484)
> subset( dmL , lex.id %in% wh.id )[,1:10]
             per      age tfD    lex.dur lex.Cst lex.Xst lex.id sex    dobth      dodm
430048 1998.956 61.87269   0   9.508556      DM    Dead      9   F 1937.083 1998.956
22671  2000.042 52.71184   0   9.954825      DM      DM     27   M 1947.331 2000.042
338459 1998.249 61.85626   0  11.748118      DM      DM     52   F 1936.393 1998.249
274124 1998.260 62.37919   0  10.929500      DM    Dead    484   F 1935.881 1998.260

> subset( dmS1, lex.id %in% wh.id )[,1:10]
   lex.id      per      age      tfD  lex.dur lex.Cst lex.Xst sex    dobth      dodm
14      9 1998.956 61.87269 0.000000 3.127310      DM      DM   F 1937.083 1998.956
15      9 2002.083 65.00000 3.127310 5.000000      DM      DM   F 1937.083 1998.956
16      9 2007.083 70.00000 8.127310 1.381246      DM    Dead   F 1937.083 1998.956
54     27 2000.042 52.71184 0.000000 2.288159      DM      DM   M 1947.331 2000.042
55     27 2002.331 55.00000 2.288159 5.000000      DM      DM   M 1947.331 2000.042
```

```
56        27 2007.331 60.00000 7.288159 2.666667       DM       DM   M 1947.331 2000.042
108       52 1998.249 61.85626 0.000000 3.143737       DM       DM   F 1936.393 1998.249
109       52 2001.393 65.00000 3.143737 5.000000       DM       DM   F 1936.393 1998.249
110       52 2006.393 70.00000 8.143737 3.604381       DM       DM   F 1936.393 1998.249
1004     484 1998.260 62.37919 0.000000 2.620808       DM       DM   F 1935.881 1998.260
1005     484 2000.881 65.00000 2.620808 5.000000       DM       DM   F 1935.881 1998.260
1006     484 2005.881 70.00000 7.620808 3.308693       DM     Dead   F 1935.881 1998.260
```

The resulting object, `dmS1`, is again a `Lexis` object, and the follow-up may be split further along another timescale, for example diabetes duration, `tfD`. Subsequently we list the results for the chosen individuals:

```
> dmS2 <- splitLexis( dmS1, "tfD", breaks=c(0,1,2,5,10,20,30,40) )
> subset( dmS2, lex.id %in% wh.id )[,1:10]
     lex.id       per      age       tfD   lex.dur lex.Cst lex.Xst sex     dobth      dodm
31        9 1998.956 61.87269  0.000000 1.0000000      DM      DM   F 1937.083 1998.956
32        9 1999.956 62.87269  1.000000 1.0000000      DM      DM   F 1937.083 1998.956
33        9 2000.956 63.87269  2.000000 1.1273101      DM      DM   F 1937.083 1998.956
34        9 2002.083 65.00000  3.127310 1.8726899      DM      DM   F 1937.083 1998.956
35        9 2003.956 66.87269  5.000000 3.1273101      DM      DM   F 1937.083 1998.956
36        9 2007.083 70.00000  8.127310 1.3812457      DM    Dead   F 1937.083 1998.956
111      27 2000.042 52.71184  0.000000 1.0000000      DM      DM   M 1947.331 2000.042
112      27 2001.042 53.71184  1.000000 1.0000000      DM      DM   M 1947.331 2000.042
113      27 2002.042 54.71184  2.000000 0.2881588      DM      DM   M 1947.331 2000.042
114      27 2002.331 55.00000  2.288159 2.7118412      DM      DM   M 1947.331 2000.042
115      27 2005.042 57.71184  5.000000 2.2881588      DM      DM   M 1947.331 2000.042
116      27 2007.331 60.00000  7.288159 2.6666667      DM      DM   M 1947.331 2000.042
229      52 1998.249 61.85626  0.000000 1.0000000      DM      DM   F 1936.393 1998.249
230      52 1999.249 62.85626  1.000000 1.0000000      DM      DM   F 1936.393 1998.249
231      52 2000.249 63.85626  2.000000 1.1437372      DM      DM   F 1936.393 1998.249
232      52 2001.393 65.00000  3.143737 1.8562628      DM      DM   F 1936.393 1998.249
233      52 2003.249 66.85626  5.000000 3.1437372      DM      DM   F 1936.393 1998.249
234      52 2006.393 70.00000  8.143737 1.8562628      DM      DM   F 1936.393 1998.249
235      52 2008.249 71.85626 10.000000 1.7481177      DM      DM   F 1936.393 1998.249
2084    484 1998.260 62.37919  0.000000 1.0000000      DM      DM   F 1935.881 1998.260
2085    484 1999.260 63.37919  1.000000 1.0000000      DM      DM   F 1935.881 1998.260
2086    484 2000.260 64.37919  2.000000 0.6208077      DM      DM   F 1935.881 1998.260
2087    484 2000.881 65.00000  2.620808 2.3791923      DM      DM   F 1935.881 1998.260
2088    484 2003.260 67.37919  5.000000 2.6208077      DM      DM   F 1935.881 1998.260
2089    484 2005.881 70.00000  7.620808 2.3791923      DM      DM   F 1935.881 1998.260
2090    484 2008.260 72.37919 10.000000 0.9295003      DM    Dead   F 1935.881 1998.260
```

A more efficient (and more intuitive) way of making this double split is to use the `splitMulti` function from the `popEpi` package:

```
> library( popEpi )
> dmM <- splitMulti( dmL, age = seq(0,100,5),
+                    tfD = c(0,1,2,5,10,20,30,40),
+              drop = FALSE )
> summary( dmS2 )
Transitions:
     To
From    DM Dead   Records:   Events: Risk time:  Persons:
  DM 40897 2499      43396      2499   54273.27      9996

> summary( dmM )
```

```
Transitions:
     To
From     DM Dead   Records:   Events: Risk time:  Persons:
  DM 40897 2499      43396      2499   54273.27      9996
```

Note we used the argument `drop=FALSE` which will retain follow-up also outside the window defined by the breaks. Otherwise the default for `splitMulti` would be to drop follow-up outside `age` [0,100] and `tfD` [0,40]. This clipping behaviour is not available in `splitLexis`, nevertheless this may be exactly what we want in some situations.

So we see that the two ways of splitting data yields the same amount of follow-up, but the results are not identical:

```
> identical( dmS2, dmM )
[1] FALSE
> class( dmS2 )
[1] "Lexis"      "data.frame"
> class( dmM )
[1] "Lexis"      "data.table" "data.frame"
```

As we see, this is because the `dmM` object also is a `data.table` object; the `splitMulti` uses the `data.table` machinery which makes the splitting substantially faster — this is of particular interest if you operate on large data sets ($> 100,000$ records).

Thus the recommended way of splitting follow-up time is by `splitMulti`. But you should be aware that the result is a `data.table` object, which in some circumstances behaves slightly different from `data.frame`s. See the manual for `data.table`.

2.3   Cutting follow up time at a dates of intermediate events

If we have a recording of the date of a specific event as for example recovery or relapse, we may classify follow-up time as being before or after this intermediate event, but it requires that follow-up records that straddle the event be cut in two and placed in separate records, one representing follow-up *before* the intermediate event, and another representing follow-up *after* the intermediate event. This is achieved with the function `cutLexis`, which takes three arguments: the time point of the intermediate event, the timescale that this point refers to, and the value of the (new) state following the date. Optionally, we may also define a new time scale with the argument `new.scale=`.

We are interested in the time before and after inception of insulin use, which occurs at the date `doins`

```
> subset( dmL, lex.id %in% wh.id )[,1:10]
            per       age tfD   lex.dur lex.Cst lex.Xst lex.id sex     dobth     dodm
430048 1998.956 61.87269   0  9.508556      DM    Dead      9   F 1937.083 1998.956
22671  2000.042 52.71184   0  9.954825      DM      DM     27   M 1947.331 2000.042
338459 1998.249 61.85626   0 11.748118      DM      DM     52   F 1936.393 1998.249
274124 1998.260 62.37919   0 10.929500      DM    Dead    484   F 1935.881 1998.260
```

```
> dmC <- cutLexis( data = dmL,
+                   cut = dmL$doins,
+            timescale = "per",
+            new.state = "Ins",
+            new.scale = "tfI",
+       precursor.states = "DM" )
> subset( dmC, lex.id %in% wh.id )[,1:10]
            per      age       tfD tfI  lex.dur lex.Cst lex.Xst lex.id sex    dobth
9      1998.956 61.87269 0.000000  NA 9.508556      DM    Dead      9   F 1937.083
27     2000.042 52.71184 0.000000  NA 9.954825      DM      DM     27   M 1947.331
52     1998.249 61.85626 0.000000  NA 6.554415      DM     Ins     52   F 1936.393
10048  2004.804 68.41068 6.554415   0 5.193703     Ins     Ins     52   F 1936.393
484    1998.260 62.37919 0.000000  NA 5.700205      DM     Ins    484   F 1935.881
10480  2003.960 68.07940 5.700205   0 5.229295     Ins    Dead    484   F 1935.881
```

(The `precursor.states=` argument is explained below). Note that individual 52 has had his follow-up cut at 6.55 years from diabetes diagnosis and individual 484 at 5.70 years from diabetes diagnosis. This dataset could then be split along the timescales as we did before with `dmL`.

We can see which of the time scales that are defined as time since entry into an intermediate state:

```
> timeSince( dmC )
  per   age   tfD   tfI
   ""    ""    "" "Ins"
```

The names of the vector are the time scales; each element is the name of the state entry into which defines the origin of the time scale.

However the result of this can also be achieved by cutting the split dataset `dmS2` instead of `dmL`:

```
> dmS2C <- cutLexis( data = dmS2,
+                     cut = dmS2$doins,
+              timescale = "per",
+              new.state = "Ins",
+              new.scale = "tfI",
+         precursor.states = "DM" )
> subset( dmS2C, lex.id %in% wh.id )[,1:10]
          per      age       tfD tfI lex.id   lex.dur lex.Cst lex.Xst sex    dobth
31   1998.956 61.87269 0.000000  NA      9 1.0000000      DM      DM   F 1937.083
32   1999.956 62.87269 1.000000  NA      9 1.0000000      DM      DM   F 1937.083
33   2000.956 63.87269 2.000000  NA      9 1.1273101      DM      DM   F 1937.083
34   2002.083 65.00000 3.127310  NA      9 1.8726899      DM      DM   F 1937.083
35   2003.956 66.87269 5.000000  NA      9 3.1273101      DM      DM   F 1937.083
36   2007.083 70.00000 8.127310  NA      9 1.3812457      DM    Dead   F 1937.083
111  2000.042 52.71184 0.000000  NA     27 1.0000000      DM      DM   M 1947.331
112  2001.042 53.71184 1.000000  NA     27 1.0000000      DM      DM   M 1947.331
113  2002.042 54.71184 2.000000  NA     27 0.2881588      DM      DM   M 1947.331
114  2002.331 55.00000 2.288159  NA     27 2.7118412      DM      DM   M 1947.331
115  2005.042 57.71184 5.000000  NA     27 2.2881588      DM      DM   M 1947.331
116  2007.331 60.00000 7.288159  NA     27 2.6666667      DM      DM   M 1947.331
229  1998.249 61.85626 0.000000  NA     52 1.0000000      DM      DM   F 1936.393
230  1999.249 62.85626 1.000000  NA     52 1.0000000      DM      DM   F 1936.393
231  2000.249 63.85626 2.000000  NA     52 1.1437372      DM      DM   F 1936.393
232  2001.393 65.00000 3.143737  NA     52 1.8562628      DM      DM   F 1936.393
```

```
233    2003.249 66.85626  5.000000        NA    52 1.5544148    DM    Ins   F 1936.393
43629 2004.804 68.41068  6.554415 0.000000    52 1.5893224    Ins   Ins   F 1936.393
43630 2006.393 70.00000  8.143737 1.589322    52 1.8562628    Ins   Ins   F 1936.393
43631 2008.249 71.85626 10.000000 3.445585    52 1.7481177    Ins   Ins   F 1936.393
2084  1998.260 62.37919  0.000000        NA   484 1.0000000    DM    DM   F 1935.881
2085  1999.260 63.37919  1.000000        NA   484 1.0000000    DM    DM   F 1935.881
2086  2000.260 64.37919  2.000000        NA   484 0.6208077    DM    DM   F 1935.881
2087  2000.881 65.00000  2.620808        NA   484 2.3791923    DM    DM   F 1935.881
2088  2003.260 67.37919  5.000000        NA   484 0.7002053    DM    Ins   F 1935.881
45484 2003.960 68.07940  5.700205 0.000000   484 1.9206023    Ins   Ins   F 1935.881
45485 2005.881 70.00000  7.620808 1.920602   484 2.3791923    Ins   Ins   F 1935.881
45486 2008.260 72.37919 10.000000 4.299795   484 0.9295003    Ins  Dead   F 1935.881
```

Thus it does not matter in which order we use `splitLexis` and `cutLexis`. Mathematicians would say that `splitLexis` and `cutLexis` are commutative.

Note in `lex.id`=484, that follow-up subsequent to the event is classified as being in state `Ins`, but that the final transition to state `Dead` is preserved. This is the point of the `precursor.states=` argument. It names the states (in this case `DM`) that will be over-written by `new.state` (in this case `Ins`), while the state `Dead` should not be updated even if it is after the time where the persons moves to state `Ins`. In other words, only state `DM` is a precursor to state `Ins`, state `Dead` is always subsequent to state `Ins`.

Note that we defined a new timescale, `tfI`, using the argument `new.scale=tfI`. This has a special status relative to the other three timescales, it is defined as time since entry into a state, namely `Ins`, this is noted in the timescale part of the summary of `Lexis` object (the information sits in the attribute `time.since` of the `Lexis` object):

```
> summary( dmS2C, timeScales=TRUE )
Transitions:
     To
From     DM  Ins Dead  Records:  Events: Risk time:  Persons:
  DM  35135 1694 2048     38877     3742    45885.49      9899
  Ins     0 5762  451      6213      451     8387.77      1791
  Sum 35135 7456 2499     45090     4193    54273.27      9996

Timescales:
  per   age   tfD   tfI
   ""    ""    "" "Ins"
```

Finally we can get a quick overview of the states and transitions by using `boxes` — `scale.R` scales transition rates to rates per 1000 PY:

```
> boxes( dmC, boxpos=TRUE, scale.R=1000, show.BE=TRUE )
```
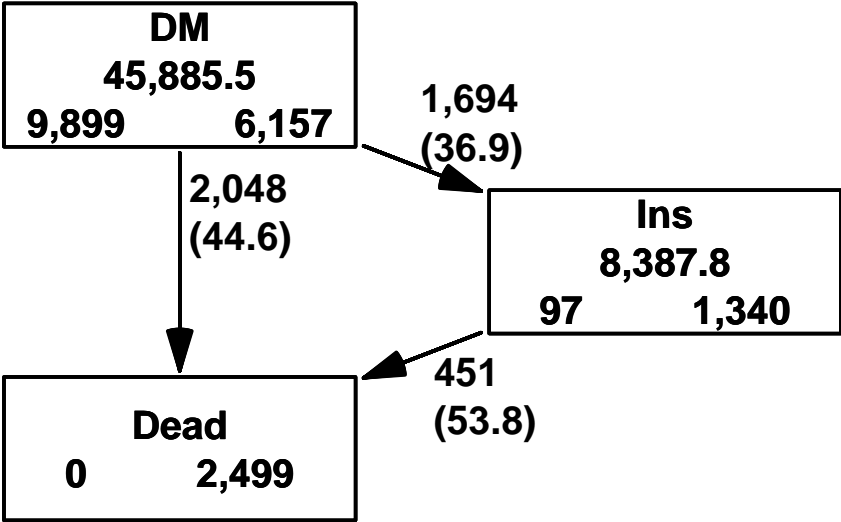
Figure 2.3: *States, person years, transitions and rates in the cut dataset. The numbers* in
*the boxes are person-years and the number of persons Beginning, resp. Ending their follow-up
in each state (triggered by* `show.BE=TRUE`*). The numbers at the arrows are the number of
transitions and transition rates per 1000 (triggered by* `scale.R=1000`*).*   `./flup-box1`

Chapter 3

Modeling rates from `Lexis` objects

3.1 Covariates

In the dataset `dmS2C` there are three types of covariates that can be used to describe mortality rates:

1. time-dependent covariates

2. time scales

3. fixed covariates

There is only one time-dependent here, namely `lex.Cst`, the current state of the person's follow up; it takes the values `DM` and `Ins` according to whether the person has ever purchased insulin at a give time of follow-up.

The time-scales are obvious candidates for explanatory variables for the rates, notably age and time from diagnosis (duration of diabetes) and insulin.

3.1.1 Time scales as covariates

If we want to model the effect of the timescale variables on occurrence rates, we will for each interval use either the value of the left endpoint in each interval or the middle. There is a function `timeBand` which returns either of these:

```
> timeBand( dmS2C, "age", "middle" )[1:20]
 [1] 57.5 57.5 62.5 62.5 62.5 67.5 67.5 62.5 67.5 67.5 67.5 67.5 72.5 87.5 87.5 87.5 92.5
[18] 92.5 42.5 77.5
> # For nice printing and column labelling we use the data.frame() function:
> data.frame( dmS2C[,c("per","age","tfD","lex.dur")],
+             mid.age=timeBand( dmS2C, "age", "middle" ),
+               mid.t=timeBand( dmS2C, "tfD", "middle" ),
+              left.t=timeBand( dmS2C, "tfD", "left"   ),
+             right.t=timeBand( dmS2C, "tfD", "right"  ),
+              fact.t=timeBand( dmS2C, "tfD", "factor" ) )[1:15,]
       per      age       tfD    lex.dur mid.age mid.t left.t right.t  fact.t
1  1998.917 58.66119 0.0000000 1.00000000    57.5   0.5      0       1   (0,1]
2  1999.917 59.66119 1.0000000 0.33880903    57.5   1.5      1       2   (1,2]
3  2000.256 60.00000 1.3388090 0.66119097    62.5   1.5      1       2   (1,2]
4  2000.917 60.66119 2.0000000 3.00000000    62.5   3.5      2       5   (2,5]
```

```
5   2003.917 63.66119  5.0000000 1.33880903     62.5   7.5     5     10  (5,10]
6   2005.256 65.00000  6.3388090 3.66119097     67.5   7.5     5     10  (5,10]
7   2008.917 68.66119 10.0000000 1.08008214     67.5  15.0    10     20 (10,20]
8   2003.309 64.09035  0.0000000 0.90965092     62.5   0.5     0      1   (0,1]
9   2004.218 65.00000  0.9096509 0.09034908     67.5   0.5     0      1   (0,1]
10  2004.309 65.09035  1.0000000 1.00000000     67.5   1.5     1      2   (1,2]
11  2005.309 66.09035  2.0000000 3.00000000     67.5   3.5     2      5   (2,5]
12  2008.309 69.09035  5.0000000 0.90965092     67.5   7.5     5     10  (5,10]
13  2009.218 70.00000  5.9096509 0.77891855     72.5   7.5     5     10  (5,10]
14  2004.552 86.25051  0.0000000 1.00000000     87.5   0.5     0      1   (0,1]
15  2005.552 87.25051  1.0000000 1.00000000     87.5   1.5     1      2   (1,2]
```

Note that the values of these functions are characteristics of the intervals defined by
`breaks=`, *not* the midpoints nor left or right endpoints of the actual follow-up intervals
(which would be `tfD` and `tfD+lex.dur`, respectively).

These functions are intended for modeling timescale variables as factors (categorical
variables) in which case the coding must be independent of the censoring and mortality
pattern — it should only depend on the chosen grouping of the timescale. Modeling
timescales as *quantitative* should not be based on these codings but directly on the values
of the time-scale variables, notably the left endpoints of the intervals.

3.1.2   Differences between time scales

Apparently, the only fixed variable is `sex`, but formally the dates of birth (`dobth`),
diagnosis (`dodom`) and first insulin purchase (`doins`) are fixed covariates too. They can be
constructed as origins of time scales referred to the calendar time scale. Likewise, and
possibly of greater interest, we can consider these origins on the age scale, calculated as the
difference between age and another time scale.

These would then be age at birth (hardly relevant), age at diabetes diagnosis and age at
insulin treatment.

3.1.3   Keeping the relation between time scales

The midpoint (as well as the right interval endpoint) should be used with caution if the
variable age at diagnosis `dodm-dobth` is modeled too; the age at diabetes is logically equal
to the difference between current age (`age`) and time since diabetes diagnosis (`tfD`):

```
> summary( (dmS2$age-dmS2$tfD) - (dmS2$dodm-dmS2$dobth) )
   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
      0       0       0       0       0       0
```

This calculation refer to the *start* of each interval — the time scale variables in the `Lexis`
object. But when using the middle of the intervals, this relationship is not preserved:

```
> summary( timeBand( dmS2, "age", "middle" ) -
+          timeBand( dmS2, "tfD", "middle" ) - (dmS2$dodm-dmS2$dobth) )
   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
-7.4870 -2.0862 -0.3765     Inf  1.3641     Inf
```

If all three variables are to be included in a model, we must make sure that the *substantial*
relationship between the variables be maintained. One way is to recompute age at diabetes

diagnosis from the two midpoint variables, but more straightforward would be to use the left endpoint of the intervals, that is the time scale variables in the `Lexis` object.

If we dissolve the relationship between the variables `age`, `tfD` and age at diagnosis by grouping we may obtain identifiability of the three separate effects, but it will be at the price of an arbitrary allocation of a linear trend between them.

For the sake of clarity, consider current age, a, duration of disease, d and age at diagnosis e, so that $a = e + d$ or $e + d - a = 0$. If we model log-rates by three functions f, g and h: $\log(\lambda) = f(a) + g(d) + h(e)$ then for any κ:

$$
\begin{aligned}
\log(\lambda) &= f(a) + g(d) + h(e) + \kappa(e + d - a) \\
&= \big(f(a) - \kappa a\big) + \big(g(d) + \kappa d\big) + \big(h(e) + \kappa e\big) \\
&= \tilde{f}(a) + \tilde{g}(d) + \tilde{h}(e)
\end{aligned}
$$

In practical modeling this will turn op as a singular model matrix with one parameter aliased. This is well known from age-period-cohort models.

Thus we see that we can move any slope around between the three terms, so if we achieve identifiability by using grouping of one of the variables we will in reality have settled for a particular value of κ, without known why we chose just that.

3.2   Modeling of rates

As mentioned, the purpose of subdividing follow-up data in smaller intervals is to be able to model effects of time scale variables as parametric functions. When we split along a timescale we can get intervals that are as small as we want; if they are sufficiently small, an assumption of constant rates in each interval becomes reasonable.

In a model that assumes a constant occurrence rate in each of the intervals the likelihood contribution from each interval is the same as the likelihood contribution from a Poisson variate D, say, with mean $\lambda\ell$ where λ is the rate and ℓ is the interval length, and where the value of the variate D is 1 or 0 according to whether an event has occurred or not. Moreover, the likelihood contributions from all follow-up intervals from a single person are *conditionally* independent (conditional on having survived till the start of the interval in question). This implies that the total contribution to the likelihood from a single person is a product of terms, and hence the same as the likelihood of a number of independent Poisson terms, one from each interval.

Note that variables are neither Poisson distributed (*e.g.* they can only ever assume values 0 or 1) nor independent — it is only the likelihood for the follow-up data that happens to be the same as the likelihood from independent Poisson variates. Different models can have the same likelihood, a model cannot be inferred from the likelihood.

Parametric modeling of the rates is obtained by using the *values* of the time scales for each interval as quantitative explanatory variables, using for example splines. And of course also the values of the fixed covariates and the time-dependent variables for each interval. Thus the model will be one where the rate is assumed constant in each (small) interval, but where a parametric form of the *size* of the rate in each interval is imposed by the model, using the timescale as a quantitative covariate.

3.2.1   Interval length

In the first chapter we illustrated cutting and splitting by listing the results for a few individuals across a number of intervals. For illustrational purposes we used 5-year age bands to avoid excessive listings, but since the doubling time for mortality on the age scale is only slightly larger than 5 years, the assumption about constant rates in each interval would be pretty far fetched if we were to use 5 year intervals.

Thus, for modeling purposes we split the follow-up in 3 month intervals. When we use intervals of 3 month's length it is superfluous to split along multiple time scales — the precise location of tightly spaced splits will be irrelevant from any practical point of view. `splitLexis` and `splitMulti` will allocate the actual split times for all of the timescale variables, so these can be used directly in modeling.

So we split the cut dataset in 3-month intervals along the age scale:

```
> dmCs <- splitMulti( dmC, age = seq(0,110,1/4) )
> summary( dmCs, t=T )
Transitions:
     To
From       DM   Ins Dead  Records:  Events: Risk time:  Persons:
  DM   189669  1694 2048    193411     3742   45885.49      9899
  Ins       0 34886  451     35337      451    8387.77      1791
  Sum  189669 36580 2499    228748     4193   54273.27      9996

Timescales:
  per   age   tfD   tfI
   ""    ""    "" "Ins"
```

We see that we have 228,748 records and 9996 persons, so about 23 records per person. The total risk time is 54,275 years, a bit less than 3 months per record as expected.

3.2.2   Practicalities for splines

In this study we want to look at how mortality depend on age (`age`) and time since start of insulin use (`tfI`). If we want to use splines in the description we must allocate knots for anchoring the splines at each of the time scales, either by some *ad hoc* method or by using some sort of penalized splines. The latter will not be treated here; it belongs in the realm of the `mgcv` package.

Here we shall use the former approach and allocate 5 knots on each of the time-scales. We allocate knots so that we have the events evenly distributed between the knots. Since the insulin state starts at 0 for all we include 0 as the first knot, such that any set of natural splines with these knots will have the value 0 at 0 on the time scale.

```
> ( a.kn <- with( subset( dmCs, lex.Xst=="Dead" ),
+               quantile( age+lex.dur, (1:5-0.5)/5 ) ) )
     10%      30%      50%      70%      90%
60.29350 71.31937 77.72758 82.72745 89.86393

> ( i.kn <- c( 0,
+         with( subset( dmCs, lex.Xst=="Dead" & lex.Cst=="Ins" ),
+               quantile( tfI+lex.dur, (1:4)/5 ) ) ) )
              20%       40%       60%       80%
0.0000000 0.3093771 1.1307324 2.5489391 4.9117043
```

In the `Epi` package there is a convenience wrapper, `Ns`, for the `natural` `spline` generator `ns`, that takes the smallest and the largest of a set of supplied knots to be the boundary knots, so the explicit definition of the boundary knots becomes superfluous.

Note that it is a feature of the `Ns` (via the features of `ns`) that any generated spline function is 0 at the leftmost knot.

3.2.3   Poisson model

A model that describes mortality rates as only a function of age would then be:

```
> ma <- glm( (lex.Xst=="Dead") ~ Ns(age,knots=a.kn),
+            family = poisson,
+            offset = log(lex.dur),
+             data = dmCs )
> summary( ma )
Call:
glm(formula = (lex.Xst == "Dead") ~ Ns(age, knots = a.kn), family = poisson,
    data = dmCs, offset = log(lex.dur))

Deviance Residuals:
    Min      1Q   Median       3Q      Max
-0.5883  -0.1688  -0.1144  -0.0766   4.5958

Coefficients:
                        Estimate Std. Error z value Pr(>|z|)
(Intercept)             -3.82830    0.03861  -99.16   <2e-16
Ns(age, knots = a.kn)1   1.36254    0.08723   15.62   <2e-16
Ns(age, knots = a.kn)2   1.49446    0.06845   21.83   <2e-16
Ns(age, knots = a.kn)3   2.63557    0.07058   37.34   <2e-16
Ns(age, knots = a.kn)4   1.94173    0.05769   33.66   <2e-16

(Dispersion parameter for poisson family taken to be 1)

    Null deviance: 27719  on 228747  degrees of freedom
Residual deviance: 25423  on 228743  degrees of freedom
AIC: 30431

Number of Fisher Scoring iterations: 8
```

The offset, `log(lex.dur)` comes from the fact that the likelihood for the follow-up data during ℓ time is the same as that for independent Poisson variates with mean $\lambda\ell$, and that the default link function for the Poisson family is the log, so that we are using a linear model for the log-mean, it is for $\log(\lambda) + \log(\ell)$. But when we want a model for the log-rate $(\log(\lambda))$, the term $\log(\ell)$ must still be included as a covariate, but with regression coefficient fixed to 1; a so-called *offset*. This is however a technicality; it just exploits that the likelihood of the Poisson model and the model for rates is the same.

In the `Epi` package is a `glm` family, `poisreg` that has a more intuitive interface, where the response is a 2-column matrix of events and person-time, respectively. This is in concert with the fact that the outcome variable in follow-up studies is bivariate: (event, risk time).

```
> Ma <- glm( cbind(lex.Xst=="Dead",lex.dur) ~ Ns(age,knots=a.kn),
+            family = poisreg, data = dmCs )
> summary( Ma )
```

```
Call:
glm(formula = cbind(lex.Xst == "Dead", lex.dur) ~ Ns(age, knots = a.kn),
    family = poisreg, data = dmCs)

Deviance Residuals:
    Min       1Q    Median       3Q      Max
-0.5883  -0.1688  -0.1144  -0.0766   4.5958

Coefficients:
                         Estimate Std. Error z value Pr(>|z|)
(Intercept)              -3.82830    0.03861  -99.15   <2e-16
Ns(age, knots = a.kn)1    1.36254    0.08723   15.62   <2e-16
Ns(age, knots = a.kn)2    1.49446    0.06845   21.83   <2e-16
Ns(age, knots = a.kn)3    2.63557    0.07058   37.34   <2e-16
Ns(age, knots = a.kn)4    1.94173    0.05769   33.66   <2e-16

(Dispersion parameter for poisson family taken to be 1)

    Null deviance: 27719  on 228747  degrees of freedom
Residual deviance: 25423  on 228743  degrees of freedom
AIC: 30431

Number of Fisher Scoring iterations: 7
```

Exploiting the multistate structure in the `Lexis` object there is a multistate convenience wrapper for `glm` with the `poisreg` family, that just requires specification of the transitions in terms of `from` and `to`. Although it is called `glm.Lexis` it is *not* as S3 method for `Lexis` objects:

```
> Xa <- glm.Lexis( dmCs, from="DM", to="Dead",
+                  formula = ~ Ns(age,knots=a.kn) )
stats::glm Poisson analysis of Lexis object dmCs with log link:
Rates for the transition DM->Dead
```

There are similar wrappers for `gam` and `coxph` models, `gam.Lexis` and `coxph.Lexis`, but not elaborated here.

The `from=` and `to=` can even be omitted, in which case all possible transitions *into* any of the absorbing states is modeled:

```
> xa <- glm.Lexis( dmCs, formula = ~ Ns(age,knots=a.kn) )
stats::glm Poisson analysis of Lexis object dmCs with log link:
Rates for transitions DM->Dead, Ins->Dead
```

We can check if the four models fitted are the same:

```
> c( deviance(ma), deviance(Ma), deviance(Xa), deviance(xa) )
[1] 25422.92 25422.92 20902.31 25422.92
```

Oops! the model `Xa` is apparently not the same as the other three? This is because the explicit specification `from="DM", to="Dead"`, omits modeling contributions from the Ins $\rightarrow$ Dead transition. The other three models all use both transitions — and assume that the two transition rates are the same, *i.e.* that start of insulin has no effect on mortality. We shall relax this assumption later.

The parameters from the model do not have any direct interpretation *per se*, but we can compute the estimated mortality rates for a range of ages using `ci.pred` with a suitably defined prediction data frame.

Note that if we use the `poisson` family of models, we must specify all covariates in the model, including the variable in the offset, `lex.dur` (remember that this was a covariate with coefficient fixed at 1). We set the latter to 1000, because we want the mortality rates per 1000 person-years. Using the `poisreg` family, the prediction will ignore any value of `lex.dur` specified in the prediction data frame, the returned rates will be per unit in which `lex.dur` is recorded.

```
> nd <- data.frame( age=40:85 )
> pr.a <- ci.pred( Ma, newdata = nd )*1000 # mortality per 100 PY
> matshade( nd$age, pr.a, plot=TRUE,
+          type="l", lty=1,
+          log="y", xlab="Age (years)",
+          ylab="DM mortality per 1000 PY")
```


Figure 3.1: *Mortality among Danish diabetes patients by age with 95% CI as shaded area. We see that the rates increase linearly on the log-scale, that is exponentially by age.* `./flup-pr-a`

3.3   Time dependent covariate

A Poisson model approach to mortality by insulin status, would be to assume that the rate-ratio between patients on insulin and not on insulin is a fixed quantity, independent of time since start of insulin, independent of age. This is commonly termed a proportional hazards assumption, because the rates (hazards) in the two groups are proportional along the age (baseline time) scale.

```
> pm <- glm( cbind(lex.Xst=="Dead",lex.dur) ~ Ns(age,knots=a.kn)
+                                            + lex.Cst + sex,
+            family=poisreg, data = dmCs )
> round( ci.exp( pm ), 3 )
                        exp(Est.)   2.5%  97.5%
(Intercept)                 0.022  0.021  0.024
Ns(age, knots = a.kn)1      4.248  3.581  5.040
Ns(age, knots = a.kn)2      5.008  4.376  5.731
Ns(age, knots = a.kn)3     16.832 14.624 19.373
Ns(age, knots = a.kn)4      7.994  7.126  8.968
lex.CstIns                  1.985  1.791  2.200
sexF                        0.668  0.617  0.724
```

So we see that persons on insulin have about twice the mortality of persons not on insulin and that women have 2/3 the mortality of men.

3.3.1   Time since insulin start

If we want to test whether the excess mortality depend in the time since start if insulin treatment, we can add a spline terms in `tfI`. But since `tfI` is a timescale defined as time since entry into a new state (`Ins`), the variable `tfI` will be missing for those in the `DM` state, so before modeling we must set the `NA`s to 0, which we do with `tsNA20` (acronym for timescale `NA`s to zero):

```
> pm <- glm( cbind(lex.Xst=="Dead",lex.dur) ~ Ns(age,knots=a.kn)
+                                            + Ns(tfI,knots=i.kn)
+                                            + lex.Cst + sex,
+            family=poisreg, data = tsNA20(dmCs) )
```

As noted before we could do this simpler with `glm.Lexis`, even without the `from=` and `to=` arguments, because we are modeling all transitions *into* the absorbing state (`Dead`):

```
> Pm <- glm.Lexis( tsNA20(dmCs),
+                  form = ~ Ns(age,knots=a.kn)
+                         + Ns(tfI,knots=i.kn)
+                         + lex.Cst + sex )
stats::glm Poisson analysis of Lexis object tsNA20(dmCs) with log link:
Rates for transitions DM->Dead, Ins->Dead
> c( deviance(Pm), deviance(pm) )
[1] 25096.33 25096.33
> identical( model.matrix(Pm), model.matrix(pm) )
[1] TRUE
```

The coding of the effect of `tfI` is so that the value is 0 at 0, so the meaning of the estimate of the effect of `lex.Cst` is the RR between persons with and without insulin, immediately after start of insulin:

```
> round( ci.exp( Pm, subset="ex" ), 3 )
            exp(Est.)   2.5% 97.5%
lex.CstIns      5.632 4.430  7.16
sexF            0.674 0.622  0.73
```

We see that the effect of sex is pretty much the same as before, but the effect of `lex.Cst` is much larger, it now refers to a different quantity, namely the RR at `tfI`=0. If we want to see the effect of time since insulin, it is best viewed jointly with the effect of age:

```
> ndI <- data.frame( expand.grid( tfI=c(NA,seq(0,15,0.1)),
+                                  ai=seq(40,80,10) ),
+                  sex="M",
+                  lex.Cst="Ins" )
> ndI <- transform( ndI, age=ai+tfI )
> head( ndI )
  tfI ai sex lex.Cst  age
1  NA 40   M     Ins   NA
2 0.0 40   M     Ins 40.0
3 0.1 40   M     Ins 40.1
4 0.2 40   M     Ins 40.2
5 0.3 40   M     Ins 40.3
6 0.4 40   M     Ins 40.4

> ndA <- data.frame( age= seq(40,100,0.1), tfI=0,  lex.Cst="DM", sex="M" )
> pri <- ci.pred( Pm, ndI )
> pra <- ci.pred( Pm, ndA )
> matshade( ndI$age, pri, plot=TRUE, las=1,
+           log="y", lty=1, col="blue" )
> matshade( ndA$age, pra )
```

We see that mortality is high just after insulin start, but falls by almost a factor 3 during the first year. Also we see that there is a tendency that mortality in a given age is smallest for those with the longest duration of insulin use.

3.4   The Cox model

Note that in the Cox-model the age is used as response variable, slightly counter-intuitive. Hence the age part of the linear predictors is not in that model:

```
> library( survival )
> cm <- coxph( Surv(age,age+lex.dur,lex.Xst=="Dead") ~
+              Ns(tfI,knots=i.kn) + lex.Cst + sex,
+              data = tsNA20(dmCs) )
```

There is also a multistate wrapper for Cox models, requiring a l.h.s. side for the `formula=` argument:

```
> Cm <- coxph.Lexis( tsNA20(dmCs),
+                    form= age ~ Ns(tfI,knots=i.kn) + lex.Cst + sex )
model survival::coxph analysis of Lexis object tsNA20(dmCs):
Rates for transitions DM->Dead, Ins->Dead

> cbind( ci.exp( cm ), ci.exp( Cm ) )
```

Figure 3.2: *Mortality rates of persons on insulin, starting insulin at ages 40,45,...,75 (red), compared with persons not on insulin (black curve). Shaded areas are 95% CI.*`../flup-ins-time`

	exp(Est.)	2.5%	97.5%	exp(Est.)	2.5%	97.5%
Ns(tfI, knots = i.kn)1	0.2984062	0.19417148	0.4585960	0.2984062	0.19417148	0.4585960
Ns(tfI, knots = i.kn)2	0.3871151	0.29011380	0.5165495	0.3871151	0.29011380	0.5165495
Ns(tfI, knots = i.kn)3	0.1239128	0.06287008	0.2442238	0.1239128	0.06287008	0.2442238
Ns(tfI, knots = i.kn)4	0.4405121	0.34839015	0.5569932	0.4405121	0.34839015	0.5569932
lex.CstIns	5.6700284	4.45011220	7.2243623	5.6700284	4.45011220	7.2243623
lex.CstDead	1.0000000	1.00000000	1.0000000	1.0000000	1.00000000	1.0000000
sexF	0.6753202	0.62316569	0.7318397	0.6753202	0.62316569	0.7318397

We can compare the estimates from the Cox model with those from the Poisson model — we must add `NA`s because the Cox-model does not give the parameters for the baseline timescale (`age`), but also remove one of the parameters, because `coxph` parametrizes factors (in this case `lex.Cst`) by all defined levels and not only by the levels present in the dataset at hand (not the line of `1.0000000`s in the print above):

```
> round( cbind( ci.exp( Pm ),
+        rbind( matrix(NA,5,3),
+           ci.exp( cm )[-6,] ) ), 3 )
               exp(Est.)  2.5%  97.5% exp(Est.)  2.5% 97.5%
(Intercept)        0.022 0.021  0.024        NA    NA    NA
```

```
Ns(age, knots = a.kn)1     4.208  3.546  4.993          NA    NA    NA
Ns(age, knots = a.kn)2     5.012  4.380  5.736          NA    NA    NA
Ns(age, knots = a.kn)3    16.560 14.386 19.063          NA    NA    NA
Ns(age, knots = a.kn)4     7.921  7.061  8.885          NA    NA    NA
Ns(tfI, knots = i.kn)1     0.298  0.194  0.458      0.298 0.194 0.459
Ns(tfI, knots = i.kn)2     0.385  0.289  0.514      0.387 0.290 0.517
Ns(tfI, knots = i.kn)3     0.125  0.064  0.246      0.124 0.063 0.244
Ns(tfI, knots = i.kn)4     0.438  0.346  0.553      0.441 0.348 0.557
lex.CstIns                 5.632  4.430  7.160      5.670 4.450 7.224
sexF                       0.674  0.622  0.730      0.675 0.623 0.732
```

Thus we see that the Poisson and Cox gives pretty much the same results. You may argue that Cox requires a smaller dataset, because there is no need to subdivide data in small intervals *before* insulin use. But certainly the time *after* insulin inception need to be if the effect of this time should be modeled.

The drawback of the Cox-modeling is that it is not possible to show the absolute rates as we did in the graph above.

3.5   Marginal effect of time since insulin

When we plot the marginal effect of `tfI` from the two models we get pretty much the same; here we plot the RR relative to `tfI`=2 years. Note that we are deriving the RR as the ratio of two sets of predictions, from the data frames `nd` and `nr` — for further details consult the help page for `ci.lin`, specifically the use of a list as the `ctr.mat` argument:

```
> nd <- data.frame( tfI=seq(0,15,,151), lex.Cst="Ins", sex="M" )
> nr <- data.frame( tfI=    2         , lex.Cst="Ins", sex="M" )
> ppr <- ci.exp( pm, list(nd,nr), xvars="age" )
> cpr <- ci.exp( cm, list(nd,nr) )
> par( mar=c(3,3,1,1), mgp=c(3,1,0)/1.6, las=1, bty="n" )
> matshade( nd$tfI, cbind(ppr,cpr), lty=c(1,2), log="y", plot=T )
```

So the duration effect is essentially the same from the two modeling approaches.

We will also want the RR relative to the non-insulin users — recall these are coded 0 on the `tfI` variable:

```
> nd <- data.frame( tfI=seq(0,15,,151), lex.Cst="Ins", sex="M" )
> nr <- data.frame( tfI=    0         , lex.Cst="DM" , sex="M" )
> ppr <- ci.exp( pm, list(nd,nr), xvars="age" )
> cpr <- ci.exp( cm, list(nd,nr) )
> par( mar=c(3,3,1,1), mgp=c(3,1,0)/1.6, las=1, bty="n" )
> matshade( nd$tfI, cbind(ppr,cpr), lty=c(1,2), log="y", plot=T )
```

This is the effect of increasing duration of insulin use *for a fixed age* which is a bit artificial, so we would like to see the *joint* effects of age and insulin duration. What we cannot see is how the duration affects mortality *relative* to current age.

Another way of interpreting this curve is as the rate ratio relative to a person not on insulin, so we see that the RR (or hazard ratio, HR as some call it) is over 5 at the start of insulin (the `lex.Cst` estimate), and decreases to about 1.5 in the long term.

Both figure 3.3 and 3.4 indicates a declining RR by insulin duration, but only from figure 3.2 it is visible that mortality actually is *in*creasing by age some 2 years after insulin start. This point would not be available if we had only fitted a Cox model where we did not have access to the baseline hazard as a function of age.

Figure 3.3: *The naked duration effects relative to 2 years of duration, black from Poisson model, red from Cox model. The two sets of estimates are identical.*   `./flup-Ieff`

3.6   Age×duration interaction

The model we fitted assumes that the RR is the same regardless of the age at start of insulin — the hazards are multiplicative. Sometimes this is termed the proportional hazards assumption: For *any* fixed age the HR is the same as a function of time since insulin, and vice versa.

   A more correct term would be "main effects model" — there is no interaction between age (the baseline time scale) and other covariates. Sp there is no need for the term "proportional hazards"; well defined and precise statistical terms for it has existed for aeons.

3.6.1   Age at insulin start

In order to check the consistency of the multiplicativity assumption across the spectrum of age at insulin inception, we can fit an interaction model. One approach to this would be using a non-linear effect of age at insulin use (for convenience we use the same knots as for age) — note that the prediction data frames are the same as we used above, because we do not compute age at insulin use as a separate variable, but rather enter it as the difference

Figure 3.4: *Insulin duration effect (state* `Ins`*) relative to no insulin (state* `DM`*), black from Poisson model, red from Cox model. The* shape *is the same as the previous figure, but the RR is now relative to non-insulin, instead of relative to insulin users at 2 years duration.*
`./flup-IeffR`

between current age (`age`) and insulin duration (`tfI`).

At first glance we might think of doing:

```
> imx <- glm.Lexis( tsNA20(dmCs),
+                   formula = ~ Ns(age    ,knots=a.kn)
+                             + Ns(    tfI,knots=i.kn)
+                             + Ns(age-tfI,knots=a.kn)
+                             + lex.Cst + sex )
stats::glm Poisson analysis of Lexis object tsNA20(dmCs) with log link:
Rates for transitions DM->Dead, Ins->Dead
```

But this will fit a model with a rate-ratio between persons with and without insulin that depends only on age at insulin start for the time *after* insulin start, the RR at `tfI`=0 will be the same at any age, which really is not the type of interaction we wanted.

If we want the `age-tfI` term to be specific for the insulin exposed so we may use one of two other approaches, that are conceptually alike but mathematically different:

```
> Im <- glm.Lexis( tsNA20(dmCs),
+                   formula = ~ Ns(age    ,knots=a.kn)
+                             + Ns(   tfI,knots=i.kn)
+                             + Ns((age-tfI)*(lex.Cst=="Ins"),knots=a.kn)
+                             + lex.Cst + sex )
stats::glm Poisson analysis of Lexis object tsNA20(dmCs) with log link:
Rates for transitions DM->Dead, Ins->Dead
> im <- glm.Lexis( tsNA20(dmCs),
+                   formula = ~ Ns(age    ,knots=a.kn)
+                             + Ns(   tfI,knots=i.kn)
+                             + lex.Cst:Ns(age-tfI,knots=a.kn)
+                             + lex.Cst + sex )
stats::glm Poisson analysis of Lexis object tsNA20(dmCs) with log link:
Rates for transitions DM->Dead, Ins->Dead
```

The first model (`Im`) has a common age-effect (`Ns(age,...)`) for persons with and without diabetes and a RR depending on insulin duration `tfI` and age at insulin (`age-tfI`). Since the linear effect of these two terms are in the model as well, a linear trend in the RR by current age (`age`) is accommodated as well.

The second model allows age-effects that differ non-linearly between person with and without insulin, because the interaction term `lex.Cst:Ns(age-tfI...` for persons not on insulin is merely an age term (since `tfI` is coded 0 for all follow-up not on insulin).

We can compare the models fitted:

```
> anova( imx, Im, im, test='Chisq')
Analysis of Deviance Table

Model 1: cbind(trt(Lx$lex.Cst, Lx$lex.Xst) %in% trnam, Lx$lex.dur) ~ Ns(age,
    knots = a.kn) + Ns(tfI, knots = i.kn) + Ns(age - tfI, knots = a.kn) +
    lex.Cst + sex
Model 2: cbind(trt(Lx$lex.Cst, Lx$lex.Xst) %in% trnam, Lx$lex.dur) ~ Ns(age,
    knots = a.kn) + Ns(tfI, knots = i.kn) + Ns((age - tfI) *
    (lex.Cst == "Ins"), knots = a.kn) + lex.Cst + sex
Model 3: cbind(trt(Lx$lex.Cst, Lx$lex.Xst) %in% trnam, Lx$lex.dur) ~ Ns(age,
    knots = a.kn) + Ns(tfI, knots = i.kn) + lex.Cst:Ns(age -
    tfI, knots = a.kn) + lex.Cst + sex
  Resid. Df Resid. Dev Df Deviance Pr(>Chi)
1    228734      25096
2    228733      25087  1   8.9631 0.002755
3    228730      25082  3   4.6804 0.196749
```

so we see that the models indeed are different, and moreover that the last model does not provide substantial further improvement, by allowing non-linear RR along the age-scale.

We can illustrate the different estimated rates from the three models:

```
> pxi <- ci.pred( imx, ndI )
> pxa <- ci.pred( imx, ndA )
> pii <- ci.pred( im , ndI )
> pia <- ci.pred( im , ndA )
> pIi <- ci.pred( Im , ndI )
> pIa <- ci.pred( Im , ndA )
> par( mar=c(3,3,1,1), mgp=c(3,1,0)/1.6, las=1, bty="n" )
> matshade( ndI$age, cbind( pxi, pii, pIi)*1000, plot=T, log="y",
+           xlab="Age", ylab="Mortality per 1000 PY",
+           lty=1, lwd=2, col=c("blue","forestgreen","red"), alpha=0.1 )
> matshade( ndA$age, cbind( pxa, pia, pIa)*1000,
+           lty=1, lwd=2, col=c("blue","forestgreen","red"), alpha=0.1 )
```

Figure 3.5: *Age at insulin as interaction between age and duration. Blue curves are from the naive interaction model (*imx*) with identical RR at* tfi*=0 at any age; green curves are from the interaction model with age at insulin, from the model (*im*) with only linear differences by age, and red lines from the full interaction model (*Im*).*   ./flup-dur-int

We can also plot the RRs only from these models; for this we need the reference frames, and the machinery from `ci.exp` allowing a list of two data frames:

```
> ndR <- transform( ndI, tfI=0, lex.Cst="DM" )
> cbind( head(ndI), head(ndR) )
  tfI ai sex lex.Cst  age tfI ai sex lex.Cst  age
1  NA 40   M     Ins   NA   0 40   M      DM   NA
2 0.0 40   M     Ins 40.0   0 40   M      DM 40.0
3 0.1 40   M     Ins 40.1   0 40   M      DM 40.1
4 0.2 40   M     Ins 40.2   0 40   M      DM 40.2
5 0.3 40   M     Ins 40.3   0 40   M      DM 40.3
6 0.4 40   M     Ins 40.4   0 40   M      DM 40.4

> Rxi <- ci.exp( imx, list(ndI,ndR) )
> Rii <- ci.exp( im , list(ndI,ndR) )
> RIi <- ci.exp( Im , list(ndI,ndR) )
> par( mar=c(3,3,1,1), mgp=c(3,1,0)/1.6, las=1, bty="n" )
> matshade( ndI$age, cbind( Rxi, Rii, RIi), plot=T, log="y",
+          lty=1, lwd=2, col=c("blue","forestgreen","red"), alpha=0.1 )
```

```
> abline( h=1 )
> abline( h=ci.exp(imx,subset="lex.Cst")[,1], lty="25", col="blue" )
```


Figure 3.6: *RR from three different interaction models. The horizontal dotted line is at the estimated effect of* `lex.Cst`, *to illustrate that the first model (blue) constrains this initial HR to be constant over age. The green curves are the extended interaction model, and the blue the full one.*   `./flup-dur-int-RR`

3.6.2 General interaction

As a final illustration we may want to explore a different kind of interaction, not defined from the duration — here we simplify the interaction by not using the second-last knot in the interaction terms. Note again that the prediction code is the same:

```
> gm <- glm.Lexis( tsNA20(dmCs),
+                  formula = ~ Ns(age,knots=a.kn)
+                           + Ns(tfI,knots=i.kn)
+                           + lex.Cst:Ns(age,knots=a.kn):Ns(tfI,knots=i.kn)
+                           + lex.Cst + sex )
stats::glm Poisson analysis of Lexis object tsNA20(dmCs) with log link:
Rates for transitions DM->Dead, Ins->Dead

> pgi <- ci.pred( gm, ndI )
> pga <- ci.pred( gm, ndA )
> par( mar=c(3,3,1,1), mgp=c(3,1,0)/1.6, las=1, bty="n" )
> matshade( ndI$age, cbind( pgi, pIi )*1000, lty=1, lwd=2, log="y", plot=T,
+           alpha=c(0.1,0.2), col=c("black","red") )
> matshade( ndA$age, cbind( pga, pIa )*1000, lty=1, lwd=2,
+           alpha=c(0.1,0.2), col=c("black","red") )
```

3.6.3 Evaluating interactions

Here we see that the interaction effect is such that in the older ages the length of insulin use has an increasing effect on mortality.

Even though there is no statistically significant interaction between age and time since start of insulin, it would be illustrative to show the RR as a function of age at insulin and age at follow-up:

```
> ndR <- transform( ndI, lex.Cst="DM", tfI=0 )
> iRR <- ci.exp( Im, ctr.mat=list(ndI,ndR) )
> gRR <- ci.exp( gm, ctr.mat=list(ndI,ndR) )
> par( mar=c(3,3,1,1), mgp=c(3,1,0)/1.6, las=1, bty="n" )
> matshade( ndI$age, cbind(gRR,iRR), lty=1, log="y", plot=TRUE,
+           col=c("black","red") )
> abline( h=1 )
```

The advantage of the parametric modeling (be that with age at insulin or general spline interaction) is that it is straight-forward to *test* whether we have an interaction.

3.7 Separate models

In the above we insisted on making a joint model for the DM→Dead and the Ins→Dead transitions, but with the complications demonstrated it would actually have been more sensible to model the two transitions separately:

```
> dmd <- glm.Lexis( tsNA20(dmCs),
+                   from="DM", to="Dead",
+                   formula = ~ Ns(age,knots=a.kn)
+                             + sex )
stats::glm Poisson analysis of Lexis object tsNA20(dmCs) with log link:
Rates for the transition DM->Dead
```

Figure 3.7: *Spline-by-spline interaction between age and duration (black), and the interaction through age at entry (red).*   `./flup-splint`

```
> ind <- glm.Lexis( tsNA20(dmCs),
+                   from="Ins", to="Dead",
+                   formula = ~ Ns(age,knots=a.kn)
+                             + Ns(tfI,knots=i.kn)
+                             + Ns(age-tfI,knots=a.kn)
+                             + sex )
```

```
stats::glm Poisson analysis of Lexis object tsNA20(dmCs) with log link:
Rates for the transition Ins->Dead
```

```
> ini <- ci.pred( ind, ndI )
> dmi <- ci.pred( dmd, ndI )
> dma <- ci.pred( dmd, ndA )
```

The estimated mortality rates:

```
> par(mar=c(3,3,1,1),mgp=c(3,1,0)/1.6,las=1,bty="n")
> matshade( ndI$age, ini*1000, plot=TRUE, log="y",
+           lwd=2, col="red" )
> matshade( ndA$age, dma*1000,
+           lwd=2, col="black" )
```

Figure 3.8: *The effect of duration of insulin use at different ages of follow-up (and age at insulin start). Estimates are from the model with an interaction term using a non-linear effect of age at insulin start (red) and using a general spline interactions (black). It appears that the general interaction over-models a bit.*   `./flup-RR-int`

The estimated RRs are computed using that the estimates from the two models are uncorrelated, and hence qualify for `ci.ratio`:

```
> par(mar=c(3,3,1,1),mgp=c(3,1,0)/1.6,las=1,bty="n")
> matshade( ndI$age, ci.ratio(ini,dmi), plot=TRUE, log="y",
+          lwd=2, col="red" )
> abline( h=1 )
```

Figure 3.9: *Left panel: Mortality rates from separate models for the two mortality transitions; the* DM→Dead *transition modeled by age alone;* Ins→Dead *transition modeled with spline effects of current age, time since insulin and and age at insulin.*
Right panel: Mortality HR of insulin vs. no insulin.

Chapter 4

More states

4.1 Subdividing states

It may be of interest to subdivide the states following the intermediate event according to whether the event has occurred or not. This will enable us to address the question of the fraction of the patients that ever go on insulin.

This is done by the argument `split.states=TRUE`.

```
> dmCs <- cutLexis( data = dmS2,
+                    cut = dmS2$doins,
+              timescale = "per",
+              new.state = "Ins",
+              new.scale = "tfI",
+        precursor.states = "DM",
+            split.states = TRUE )
> summary( dmCs )
Transitions:
     To
From     DM   Ins Dead Dead(Ins)  Records:  Events: Risk time:  Persons:
  DM   35135 1694 2048         0     38877     3742   45885.49      9899
  Ins      0 5762    0       451      6213      451    8387.77      1791
  Sum  35135 7456 2048       451     45090     4193   54273.27      9996
```

We can illustrate the numbers and the transitions

```
> boxes( dmCs, boxpos=list(x=c(15,15,85,85),
+                          y=c(85,15,85,15)),
+        scale.R=1000, show.BE=TRUE )
```

Note that it is only the mortality rates that we have been modeling, namely the transitions DM→Dead and Ins→Dead(Ins). If we were to model the cumulative risk of using insulin we would also need a model for the DM→Ins transition. Subsequent to that we would then compute the probability of being in each state conditional on suitable starting conditions. With models where transition rates depend on several time scales this is not a trivial task. This is treated in more detail in the vignette on `simLexis`.

4.2 Multiple intermediate events

We may be interested in starting either insulin or OAD (oral anti-diabetic drugs), thus giving rise to more states and more timescales. This can be accomplished by the

Figure 4.1: *Transitions between 4 states: the numbers* in *the boxes are person-years (middle), and below the number of persons who start, respectively end their follow-up in each of the states.* ./flup-box4

mcutLexis function, that generalizes cutLexis:

```
> dmM <- mcutLexis( dmL,
+            timescale = "per",
+                   wh = c("doins","dooad"),
+           new.states = c("Ins","OAD"),
+           new.scales = c("tfI","tfO"),
+      precursor.states = "DM",
+         ties.resolve = TRUE )
NOTE:  9996 records with tied events times resolved.
 Results only reproducible if the seed for the random number generator is set.
> summary( dmM, t=T )
Transitions:
     To
From      DM Dead  OAD  Ins OAD-Ins Ins-OAD  Records:  Events: Risk time:  Persons:
   DM    2830 1056 2957  689       0       0      7532     4702   22920.31      7532
   OAD      0  992 3327    0    1005       0      5324     1997   22965.24      5324
   Ins      0  152    0  462       0     172       786      324    3883.07       786
   OAD-Ins  0  266    0    0     739       0      1005      266    3770.54      1005
```

```
    Ins-OAD     0   33    0    0          0    139        172         33     734.10            172
    Sum      2830 2499 6284 1151      1744    311      14819       7322   54273.27           9996
```

```
Timescales:
  per   age   tfD    tfI    tfO
   ""    ""     "" "Ins" "OAD"
```

We see that we now have two timescales defined as entry into states. We can also illustrate the transitions to the different states:

```
> wh <- c(subset(dmM,lex.Cst=="Ins-OAD")$lex.id[1:2],
+         subset(dmM,lex.Cst=="OAD-Ins")$lex.id[1:2])
> options( width=110 )
> subset( dmM, lex.id %in% wh )[,c('lex.id',names(dmM[1:8]),c("doins","dooad"))]
      lex.id         tfI      tfO      per      age         tfD     lex.dur lex.Cst lex.Xst
11675     18          NA       NA 1996.746 61.72211 0.00000000 1.169062286      DM     OAD
11676     18          NA 0.000000 1997.915 62.89117 1.16906229 8.079397673     OAD OAD-Ins
11677     18 0.000000000 8.079398 2005.995 70.97057 9.24845996 4.002737851 OAD-Ins OAD-Ins
11678     25          NA       NA 2003.689 60.34223 0.00000000 1.880903491      DM     OAD
11679     25          NA 0.000000 2005.570 62.22313 1.88090349 3.069130732     OAD OAD-Ins
11680     25 0.000000000 3.069131 2008.639 65.29227 4.95003422 1.357973990 OAD-Ins OAD-Ins
14324     20          NA       NA 2009.247 53.21834 0.00000000 0.038735436      DM     Ins
14325     20 0.000000000       NA 2009.286 53.25708 0.03873544 0.001897682     Ins Ins-OAD
14326     20 0.001897682 0.000000 2009.288 53.25898 0.04063312 0.709537998 Ins-OAD Ins-OAD
14327     38          NA       NA 2008.366 63.93155 0.00000000 0.093086927      DM     Ins
14328     38 0.000000000       NA 2008.459 64.02464 0.09308693 0.213552361     Ins Ins-OAD
14329     38 0.213552361 0.000000 2008.672 64.23819 0.30663929 1.325119781 Ins-OAD    Dead
```

```
> boxes( dmM, boxpos=list(x=c(15,80,40,40,85,85),
+                         y=c(50,50,90,10,90,10)),
+             scale.R=1000, show.BE=TRUE )
```
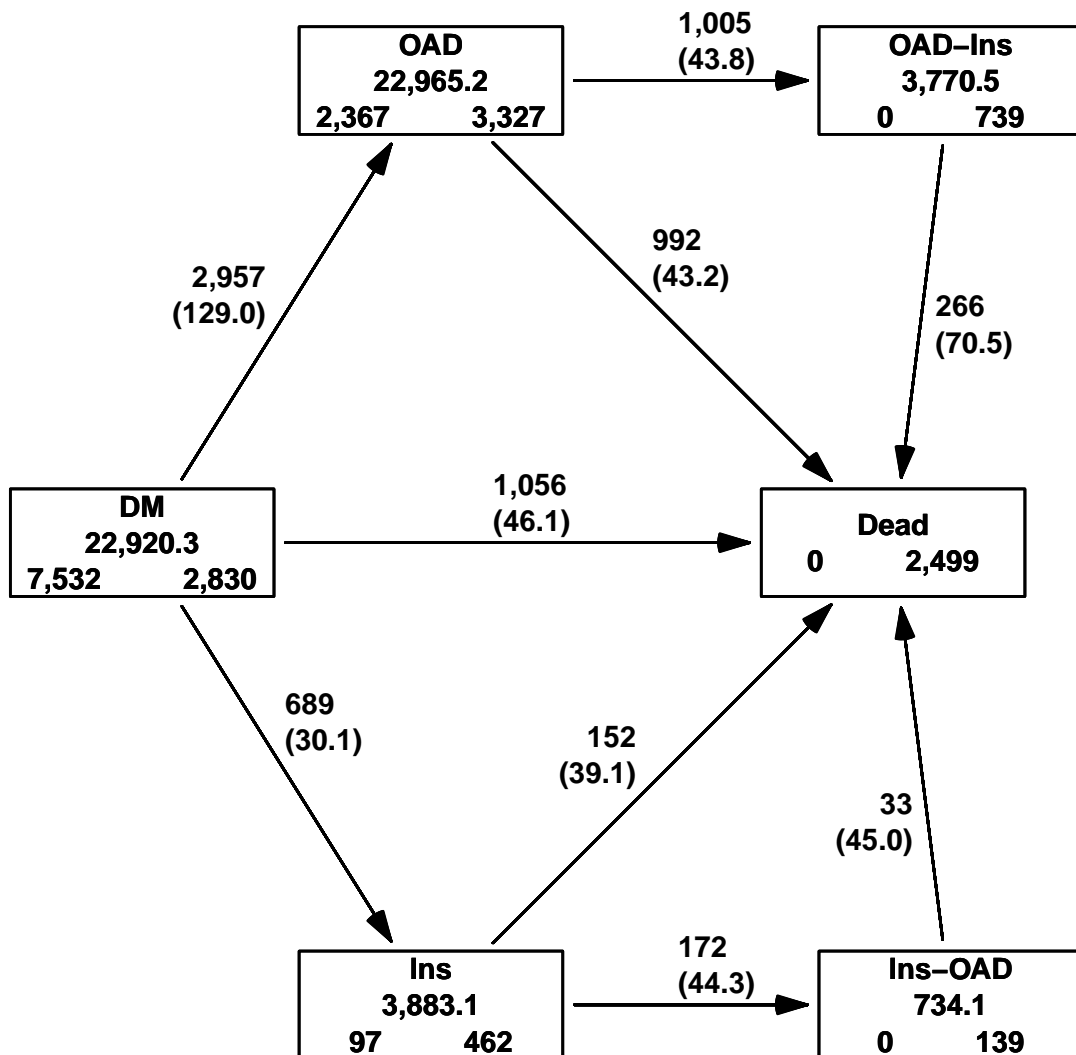
Figure 4.2: *Boxes for the* dmM *object.* ./flup-mbox