

Follow-up data with R and Epi

SDCC

December 2017

<http://bendixcarstensen.com/>

Compiled Tuesday 12th December, 2017, 05:25
from: /home/bendix/stat/R/lib.src/Epi/pkg/vignettes/flup.tex

Bendix Carstensen Steno Diabetes Center, Gentofte, Denmark
& Department of Biostatistics, University of Copenhagen
b@bxc.dk
<http://BendixCarstensen.com>

Contents

1	Follow-up data in the Epi package	1
1.1	Timescales	1
2	Subdividing follow-up for analysis	5
2.1	Splitting the follow-up time along a timescale	5
2.1.1	Time scales as covariates	7
2.1.2	Differences between time scales	8
2.2	Cutting follow up time at a specific date	8
3	Modeling rates	11
3.1	Background	11
3.2	Practicalities	11
3.3	Models for rates	12
3.3.1	One time scale	12
3.3.2	More time scales	13
3.3.3	Difference between time scales	16
3.3.4	The complete picture — exercise	17
4	Competing risks — multiple types of events	20
4.1	Subdividing states	22

Chapter 1

Follow-up data in the Epi package

In the `Epi`-package, follow-up data is represented by adding some extra variables to a data frame. Such a data frame is called a `Lexis` object. The tools for handling follow-up data then use the structure of this for special plots, tabulations etc.

Follow-up data basically consists of a time of entry, a time of exit and an indication of the status at exit (normally either “alive” or “dead”). Implicitly is also assumed a status *during* the follow-up (usually “alive”).




Figure 1.1: *Follow-up of two persons*

1.1 Timescales

A timescale is a variable that varies deterministically *within* each person during follow-up, *e.g.*:

- Age
- Calendar time
- Time since treatment
- Time since relapse

All timescales advance at the same pace, so the time followed is the same on all timescales. Therefore, it suffices to use only the entry point on each of the time scale, for example:

- Age at entry.
- Date of entry.
- Time since treatment (*at* treatment this is 0).

- Time since relapse (*at* relapse this is 0)..

For illustration we need to load the `Epi` package:

```
> library(Epi)
> print( sessionInfo(), l=F )
R version 3.4.3 (2017-11-30)
Platform: x86_64-pc-linux-gnu (64-bit)
Running under: Ubuntu 14.04.5 LTS

Matrix products: default
BLAS: /usr/lib/openblas-base/libopenblas.so.0
LAPACK: /usr/lib/lapack/liblapack.so.3.0

attached base packages:
[1] utils      datasets  graphics  grDevices  stats      methods    base

other attached packages:
[1] Epi_2.24

loaded via a namespace (and not attached):
[1] cmpsrsk_2.2-7      zoo_1.8-0          MASS_7.3-47        compiler_3.4.3
[5] Matrix_1.2-11     plyr_1.8.4         parallel_3.4.3     survival_2.41-3
[9] etm_0.6-2         Rcpp_0.12.12       splines_3.4.3      grid_3.4.3
[13] numDeriv_2016.8-1 lattice_0.20-35
```

In the `Epi` package, follow-up in a cohort is represented in a `Lexis` object. A `Lexis` object is a data frame with a bit of extra structure representing the follow-up. For the `nickel` data we would construct a `Lexis` object by:

```
> data( nickel )
> nicL <- Lexis( entry = list( per=agein+dob,
+                             age=agein,
+                             tfh=agein-age1st ),
+               exit = list( age=ageout ),
+               exit.status = ( icd %in% c(162,163) ) * 1,
+               data = nickel )
```

The `entry` argument is a *named* list with the entry points on each of the timescales we want to use. It defines the names of the timescales and the entry points of the follow-up of each person. The `exit` argument gives the exit time on *one* of the timescales, so the name of the element in this list must match one of the names of the `entry` list. This is sufficient, because the follow-up time on all time scales is the same, in this case `ageout - agein`. Now take a look at the result:

```
> str( nickel )
'data.frame':      679 obs. of  7 variables:
 $ id      : num  3 4 6 8 9 10 15 16 17 18 ...
 $ icd     : num  0 162 163 527 150 163 334 160 420 12 ...
 $ exposure: num  5 5 10 9 0 2 0 0.5 0 0 ...
 $ dob     : num  1889 1886 1881 1886 1880 ...
 $ age1st  : num  17.5 23.2 25.2 24.7 30 ...
 $ agein   : num  45.2 48.3 53 47.9 54.7 ...
 $ ageout  : num  93 63.3 54.2 69.7 76.8 ...

> str( nicL )
```

```

Classes 'Lexis' and 'data.frame':      679 obs. of  14 variables:
 $ per      : num  1934 1934 1934 1934 1934 ...
 $ age      : num  45.2 48.3 53 47.9 54.7 ...
 $ tfh      : num  27.7 25.1 27.7 23.2 24.8 ...
 $ lex.dur  : num  47.75 15 1.17 21.77 22.1 ...
 $ lex.Cst  : num  0 0 0 0 0 0 0 0 0 0 ...
 $ lex.Xst  : num  0 1 1 0 0 1 0 0 0 0 ...
 $ lex.id   : int  1 2 3 4 5 6 7 8 9 10 ...
 $ id       : num  3 4 6 8 9 10 15 16 17 18 ...
 $ icd      : num  0 162 163 527 150 163 334 160 420 12 ...
 $ exposure: num  5 5 10 9 0 2 0 0.5 0 0 ...
 $ dob      : num  1889 1886 1881 1886 1880 ...
 $ age1st   : num  17.5 23.2 25.2 24.7 30 ...
 $ agein    : num  45.2 48.3 53 47.9 54.7 ...
 $ ageout   : num  93 63.3 54.2 69.7 76.8 ...
 - attr(*, "time.scales")= chr  "per" "age" "tfh"
 - attr(*, "time.since")= chr  "" "" ""
 - attr(*, "breaks")=List of 3
 ..$ per: NULL
 ..$ age: NULL
 ..$ tfh: NULL
> head( nicL )
      per      age      tfh lex.dur lex.Cst lex.Xst lex.id id icd exposure      dob
1 1934.246 45.2273 27.7465 47.7535      0      0      1 3 0      5 1889.019
2 1934.246 48.2684 25.0820 15.0028      0      1      2 4 162      5 1885.978
3 1934.246 52.9917 27.7465 1.1727      0      1      3 6 163     10 1881.255
4 1934.246 47.9067 23.1861 21.7727      0      0      4 8 527      9 1886.340
5 1934.246 54.7465 24.7890 22.0977      0      0      5 9 150      0 1879.500
6 1934.246 44.3314 23.0437 18.2099      0      1      6 10 163      2 1889.915
      age1st agein ageout
1 17.4808 45.2273 92.9808
2 23.1864 48.2684 63.2712
3 25.2452 52.9917 54.1644
4 24.7206 47.9067 69.6794
5 29.9575 54.7465 76.8442
6 21.2877 44.3314 62.5413

```

The `Lexis` object `nicL` has a variable for each timescale which is the entry point on this timescale. The follow-up time is in the variable `lex.dur` (**d**uration).

There is a `summary` function for `Lexis` objects that list the number of transitions and records as well as the total amount of follow-up time:

```

> summary( nicL )
Transitions:
      To
From  0   1 Records:  Events: Risk time:  Persons:
      0 542 137      679      137  15348.06      679

```

We defined the exit status to be death from lung cancer (ICD7 162,163), i.e. this variable is 1 if follow-up ended with a death from this cause. If follow-up ended alive or by death from another cause, the exit status is coded 0, i.e. as a censoring.

Note that the exit status is in the variable `lex.Xst` (e**X**it status). The variable `lex.Cst` is the state where the follow-up takes place (C**u**rrent status), in this case 0 (alive).

It is possible to get a visualization of the follow-up along the timescales chosen by using the `plot` method for `Lexis` objects. `nicL` is an object of *class* `Lexis`, so using the function `plot()` on it means that R will look for the function `plot.Lexis` and use this function.




Figure 1.2: *Lexis diagram of the nickel dataset; left panel the default version, right panel with bells and whistles. The red lines are for persons with exposure > 0, so it is pretty evident that the oldest ones are the exposed part of the cohort.*

```
> plot( nicL )
```

The function allows quite a bit of control over the output, and a `points.Lexis` function allows plotting of the endpoints of follow-up:

```
> par( mar=c(3,3,1,1), mgp=c(3,1,0)/1.6 )
> plot( nicL, 1:2, lwd=1, col=c("blue","red")[(nicL$exp>0)+1],
+       grid=TRUE, lty.grid=1, col.grid=gray(0.7),
+       xlim=1900+c(0,90), xaxs="i",
+       ylim= 10+c(0,90), yaxs="i", las=1 )
> points( nicL, 1:2, pch=c(NA,3)[nicL$lex.Xst+1],
+         col="lightgray", lwd=3, cex=1.5 )
> points( nicL, 1:2, pch=c(NA,3)[nicL$lex.Xst+1],
+         col=c("blue","red")[(nicL$exp>0)+1], lwd=1, cex=1.5 )
```

The results of these two plotting commands are in figure ??.

Chapter 2

Subdividing follow-up for analysis

2.1 Splitting the follow-up time along a timescale

The follow-up time in a cohort can be subdivided by for example current age. This is achieved by the `splitLexis` (note that it is *not* called `split.Lexis`). This requires that the timescale and the breakpoints on this timescale are supplied. Try:

```
> nicS1 <- splitLexis( nicL, "age", breaks=seq(0,100,10) )
> summary( nicL )
```

Transitions:

	From	To	0	1	Records:	Events:	Risk time:	Persons:
	0	542	137	679	137	15348.06	679	

```
> summary( nicS1 )
```

Transitions:

	From	To	0	1	Records:	Events:	Risk time:	Persons:
	0	2073	137	2210	137	15348.06	679	

So we see that the number of events and the amount of follow-up is the same in the two data sets; only the number of records differ — the extra records all have `lex.Cst=0` and `lex.Xst=0`.

To see how records are split for each individual, it is useful to list the results for a few individuals:

```
> round( subset( nicS1, id %in% 8:10 ), 2 )
```

	lex.id	per	age	tfh	lex.dur	lex.Cst	lex.Xst	id	icd	exposure	dob	age1st
11	4	1934.25	47.91	23.19	2.09	0	0	8	527	9	1886.34	24.72
12	4	1936.34	50.00	25.28	10.00	0	0	8	527	9	1886.34	24.72
13	4	1946.34	60.00	35.28	9.68	0	0	8	527	9	1886.34	24.72
14	5	1934.25	54.75	24.79	5.25	0	0	9	150	0	1879.50	29.96
15	5	1939.50	60.00	30.04	10.00	0	0	9	150	0	1879.50	29.96
16	5	1949.50	70.00	40.04	6.84	0	0	9	150	0	1879.50	29.96
17	6	1934.25	44.33	23.04	5.67	0	0	10	163	2	1889.91	21.29
18	6	1939.91	50.00	28.71	10.00	0	0	10	163	2	1889.91	21.29
19	6	1949.91	60.00	38.71	2.54	0	1	10	163	2	1889.91	21.29
	agein	ageout										
11	47.91	69.68										
12	47.91	69.68										
13	47.91	69.68										

```

14 54.75 76.84
15 54.75 76.84
16 54.75 76.84
17 44.33 62.54
18 44.33 62.54
19 44.33 62.54

```

The resulting object, `nicS1`, is again a `Lexis` object, and so follow-up may be split further along another timescale. Subsequently we list the results for individuals 8, 9 and 10 again:

```

> nicS2 <- splitLexis( nicS1, "tfh", breaks=c(0,1,5,10,20,30,100) )
> round( subset( nicS2, id %in% 8:10 ), 2 )

```

	lex.id	per	age	tfh	lex.dur	lex.Cst	lex.Xst	id	icd	exposure	dob	age1st
13	4	1934.25	47.91	23.19	2.09	0	0	8	527	9	1886.34	24.72
14	4	1936.34	50.00	25.28	4.72	0	0	8	527	9	1886.34	24.72
15	4	1941.06	54.72	30.00	5.28	0	0	8	527	9	1886.34	24.72
16	4	1946.34	60.00	35.28	9.68	0	0	8	527	9	1886.34	24.72
17	5	1934.25	54.75	24.79	5.21	0	0	9	150	0	1879.50	29.96
18	5	1939.46	59.96	30.00	0.04	0	0	9	150	0	1879.50	29.96
19	5	1939.50	60.00	30.04	10.00	0	0	9	150	0	1879.50	29.96
20	5	1949.50	70.00	40.04	6.84	0	0	9	150	0	1879.50	29.96
21	6	1934.25	44.33	23.04	5.67	0	0	10	163	2	1889.91	21.29
22	6	1939.91	50.00	28.71	1.29	0	0	10	163	2	1889.91	21.29
23	6	1941.20	51.29	30.00	8.71	0	0	10	163	2	1889.91	21.29
24	6	1949.91	60.00	38.71	2.54	0	1	10	163	2	1889.91	21.29

```

      agein ageout
13 47.91 69.68
14 47.91 69.68
15 47.91 69.68
16 47.91 69.68
17 54.75 76.84
18 54.75 76.84
19 54.75 76.84
20 54.75 76.84
21 44.33 62.54
22 44.33 62.54
23 44.33 62.54
24 44.33 62.54

```

A more efficient (and more intuitive) way of making this double split is to use the `splitMulti` function from the `popEpi` package:

```

> library( popEpi )
> nicM <- splitMulti( nicL, age = seq(0,100,10),
+                    tfh = c(0,1,5,10,20,30,100) )
> summary( nicS2 )

```

Transitions:

```

      To
From  0  1 Records: Events: Risk time: Persons:
    0 2992 137      3129      137  15348.06      679

```

```

> summary( nicM )

```

Transitions:

```

      To
From  0  1 Records: Events: Risk time: Persons:
    0 2992 137      3129      137  15348.06      679

```


So we see that the two ways of splitting data yields the same amount of follow-up, but the results are not identical:

```
> identical( nicS2, nicM )
[1] FALSE
> class( nicS2 )
[1] "Lexis"          "data.frame"
> class( nicM )
[1] "Lexis"          "data.table" "data.frame"
```

As we see, this is because the `nicM` object also is a `data.table` object; the `splitMulti` uses the `data.table` machinery which makes the splitting substantially faster — this is of particular interest if you operate on large data sets ($> 1,000,000$ records).

Thus the recommended way of splitting follow-up time is by `splitMulti`. But you should be aware that the result is a `data.table` object, which in some circumstances behaves slightly different from `data.frames`. See the manual for `data.table`.

2.1.1 Time scales as covariates

If we want to model the effect of these timescale variables on occurrence rates, we will for each interval use either the value of the left endpoint in each interval or the middle. There is a function `timeBand` which returns either of these:

```
> timeBand( nicM, "age", "middle" )[1:20]
[1] 45 45 55 65 75 85 95 45 55 55 65 55 45 55 55 65 55 55 65 75
> # For nice printing and column labelling use the data.frame() function:
> data.frame( nicS2[,c("lex.id","per","age","tfh","lex.dur")],
+           mid.age=timeBand( nicS2, "age", "middle" ),
+           mid.t=timeBand( nicS2, "tfh", "middle" ),
+           left.t=timeBand( nicS2, "tfh", "left" ),
+           right.t=timeBand( nicS2, "tfh", "right" ),
+           fact.t=timeBand( nicS2, "tfh", "factor" ) )[1:20,]
```

	lex.id	per	age	tfh	lex.dur	mid.age	mid.t	left.t	right.t	fact.t
1	1	1934.246	45.2273	27.7465	2.2535	45	25	20	30	(20,30]
2	1	1936.500	47.4808	30.0000	2.5192	45	65	30	100	(30,100]
3	1	1939.019	50.0000	32.5192	10.0000	55	65	30	100	(30,100]
4	1	1949.019	60.0000	42.5192	10.0000	65	65	30	100	(30,100]
5	1	1959.019	70.0000	52.5192	10.0000	75	65	30	100	(30,100]
6	1	1969.019	80.0000	62.5192	10.0000	85	65	30	100	(30,100]
7	1	1979.019	90.0000	72.5192	2.9808	95	65	30	100	(30,100]
8	2	1934.246	48.2684	25.0820	1.7316	45	25	20	30	(20,30]
9	2	1935.978	50.0000	26.8136	3.1864	55	25	20	30	(20,30]
10	2	1939.164	53.1864	30.0000	6.8136	55	65	30	100	(30,100]
11	2	1945.978	60.0000	36.8136	3.2712	65	65	30	100	(30,100]
12	3	1934.246	52.9917	27.7465	1.1727	55	25	20	30	(20,30]
13	4	1934.246	47.9067	23.1861	2.0933	45	25	20	30	(20,30]
14	4	1936.340	50.0000	25.2794	4.7206	55	25	20	30	(20,30]
15	4	1941.060	54.7206	30.0000	5.2794	55	65	30	100	(30,100]
16	4	1946.340	60.0000	35.2794	9.6794	65	65	30	100	(30,100]
17	5	1934.246	54.7465	24.7890	5.2110	55	25	20	30	(20,30]
18	5	1939.457	59.9575	30.0000	0.0425	55	65	30	100	(30,100]
19	5	1939.500	60.0000	30.0425	10.0000	65	65	30	100	(30,100]
20	5	1949.500	70.0000	40.0425	6.8442	75	65	30	100	(30,100]

Note that these are characteristics of the intervals defined by `breaks=`, *not* the midpoints nor left or right endpoints of the actual follow-up intervals (which would be `tfh` and `tfh+lex.dur`, respectively).

These functions are intended for modeling timescale variables as factors (categorical variables) in which case the coding must be independent of the censoring and mortality pattern — it should only depend on the chosen grouping of the timescale. Modeling timescales as *quantitative* should not be based on these codings but directly on the values of the time-scale variables.

2.1.2 Differences between time scales

The midpoint (as well as the left and right interval endpoint) should be used with caution if the variable `age1st` is modeled too; the age at hire is logically equal to the difference between current age (`age`) and time since hire (`thf`):

```
> summary( (nicS2$age-nicS2$tfh) - nicS2$age1st )
      Min.      1st Qu.      Median      Mean      3rd Qu.      Max.
-7.105e-15  0.000e+00  0.000e+00  2.214e-17  0.000e+00  7.105e-15
```

This calculation refer to the *start* of each interval — the time scale variables in the `Lexis` object. But when using the middle of the intervals, this relationship is not preserved:

```
> summary( timeBand( nicS2, "age", "middle" ) -
+          timeBand( nicS2, "tfh", "middle" ) - nicS2$age1st )
      Min. 1st Qu. Median      Mean 3rd Qu.      Max.
-39.958 -24.178  -5.103 -10.129   2.575  12.519
```

If all three variable are to be included in a model, you must make sure that the *substantial* relationship between the variables be maintained. One way is to recompute age at first hire from the two midpoint variables, but more straightforward would be to use the left endpoint of the intervals, that is the time scale variables in the `Lexis` object. The latter approach however requires that the follow-up is split in fairly small chunks.

2.2 Cutting follow up time at a specific date

If we have a recording of the date of a specific event as for example recovery or relapse, we may classify follow-up time as being before or after this intermediate event, but it requires that follow-up records that straddle the event be cut into two record. This is achieved with the function `cutLexis`, which takes three arguments: the time point, the timescale, and the value of the (new) state following the date.

Now we define the age for the nickel workers where the cumulative exposure exceeds 50 exposure years:

```
> subset( nicL, id %in% 8:10 )
      per      age      tfh lex.dur lex.Cst lex.Xst lex.id id icd exposure      dob
4 1934.246 47.9067 23.1861 21.7727      0      0      4  8 527      9 1886.340
5 1934.246 54.7465 24.7890 22.0977      0      0      5  9 150      0 1879.500
6 1934.246 44.3314 23.0437 18.2099      0      1      6 10 163      2 1889.915
      age1st agein ageout
4 24.7206 47.9067 69.6794
5 29.9575 54.7465 76.8442
6 21.2877 44.3314 62.5413
```

```
> agehi <- nicL$age1st + 50 / nicL$exposure
> nicC <- cutLexis( data = nicL,
+                   cut = agehi,
+                   timescale = "age",
+                   new.state = 2,
+                   precursor.states = 0 )
> subset( nicC, id %in% 8:10 )
```

	per	age	tfh	lex.dur	lex.Cst	lex.Xst	lex.id	id	icd	exposure	dob
683	1934.246	47.9067	23.1861	21.7727	2	2	4	8	527	9	1886.340
5	1934.246	54.7465	24.7890	22.0977	0	0	5	9	150	0	1879.500
6	1934.246	44.3314	23.0437	1.9563	0	2	6	10	163	2	1889.915
685	1936.203	46.2877	25.0000	16.2536	2	1	6	10	163	2	1889.915
	age1st	agein	ageout								
683	24.7206	47.9067	69.6794								
5	29.9575	54.7465	76.8442								
6	21.2877	44.3314	62.5413								
685	21.2877	44.3314	62.5413								

(The `precursor.states=` argument is explained below). Note that individual 6 has had his follow-up split at 25 years since hire where 50 exposure-years were attained. This could also have been achieved in the split dataset `nicS2` instead of `nicL`, try:

```
> subset( nicS2, id %in% 8:10 )
```

	lex.id	per	age	tfh	lex.dur	lex.Cst	lex.Xst	id	icd	exposure	dob
13	4	1934.246	47.9067	23.1861	2.0933	0	0	8	527	9	1886.340
14	4	1936.340	50.0000	25.2794	4.7206	0	0	8	527	9	1886.340
15	4	1941.060	54.7206	30.0000	5.2794	0	0	8	527	9	1886.340
16	4	1946.340	60.0000	35.2794	9.6794	0	0	8	527	9	1886.340
17	5	1934.246	54.7465	24.7890	5.2110	0	0	9	150	0	1879.500
18	5	1939.457	59.9575	30.0000	0.0425	0	0	9	150	0	1879.500
19	5	1939.500	60.0000	30.0425	10.0000	0	0	9	150	0	1879.500
20	5	1949.500	70.0000	40.0425	6.8442	0	0	9	150	0	1879.500
21	6	1934.246	44.3314	23.0437	5.6686	0	0	10	163	2	1889.915
22	6	1939.915	50.0000	28.7123	1.2877	0	0	10	163	2	1889.915
23	6	1941.203	51.2877	30.0000	8.7123	0	0	10	163	2	1889.915
24	6	1949.915	60.0000	38.7123	2.5413	0	1	10	163	2	1889.915
	age1st	agein	ageout								
13	24.7206	47.9067	69.6794								
14	24.7206	47.9067	69.6794								
15	24.7206	47.9067	69.6794								
16	24.7206	47.9067	69.6794								
17	29.9575	54.7465	76.8442								
18	29.9575	54.7465	76.8442								
19	29.9575	54.7465	76.8442								
20	29.9575	54.7465	76.8442								
21	21.2877	44.3314	62.5413								
22	21.2877	44.3314	62.5413								
23	21.2877	44.3314	62.5413								
24	21.2877	44.3314	62.5413								

```
> agehi <- nicS2$age1st + 50 / nicS2$exposure
> nicS2C <- cutLexis( data = nicS2,
+                     cut = agehi,
+                     timescale = "age",
+                     new.state = 2,
+                     precursor.states = 0 )
> subset( nicS2C, id %in% 8:10 )
```

	lex.id	per	age	tfh	lex.dur	lex.Cst	lex.Xst	id	icd	exposure	dob
3142	4	1934.246	47.9067	23.1861	2.0933	2	2	8	527	9	1886.340
3143	4	1936.340	50.0000	25.2794	4.7206	2	2	8	527	9	1886.340
3144	4	1941.060	54.7206	30.0000	5.2794	2	2	8	527	9	1886.340
3145	4	1946.340	60.0000	35.2794	9.6794	2	2	8	527	9	1886.340
17	5	1934.246	54.7465	24.7890	5.2110	0	0	9	150	0	1879.500
18	5	1939.457	59.9575	30.0000	0.0425	0	0	9	150	0	1879.500
19	5	1939.500	60.0000	30.0425	10.0000	0	0	9	150	0	1879.500
20	5	1949.500	70.0000	40.0425	6.8442	0	0	9	150	0	1879.500
21	6	1934.246	44.3314	23.0437	1.9563	0	2	10	163	2	1889.915
3150	6	1936.203	46.2877	25.0000	3.7123	2	2	10	163	2	1889.915
3151	6	1939.915	50.0000	28.7123	1.2877	2	2	10	163	2	1889.915
3152	6	1941.203	51.2877	30.0000	8.7123	2	2	10	163	2	1889.915
3153	6	1949.915	60.0000	38.7123	2.5413	2	1	10	163	2	1889.915
	age1st	agein	ageout								
3142	24.7206	47.9067	69.6794								
3143	24.7206	47.9067	69.6794								
3144	24.7206	47.9067	69.6794								
3145	24.7206	47.9067	69.6794								
17	29.9575	54.7465	76.8442								
18	29.9575	54.7465	76.8442								
19	29.9575	54.7465	76.8442								
20	29.9575	54.7465	76.8442								
21	21.2877	44.3314	62.5413								
3150	21.2877	44.3314	62.5413								
3151	21.2877	44.3314	62.5413								
3152	21.2877	44.3314	62.5413								
3153	21.2877	44.3314	62.5413								

The same results would have emerged if we had used the `nicM` dataset (the `data.table` object). Mathematicians would say that `splitLexis` and `cutLexis` are commutative.

Note that follow-up subsequent to the event is classified as being in state 2, but that the final transition to state 1 (death from lung cancer) is preserved. This is the point of the `precursor.states=` argument. It names the states (in this case 0, “Alive”) that will be over-written by `new.state` (in this case state 2, “High exposure”), while state 1 (“Dead”) should not be updated even if it is after the time where the persons moves to state 2. In other words, only state 0 is a precursor to state 2, state 1 is always subsequent to state 2. Even if you at a high exposure level, death is pretty final.

If the intermediate event is to be used as a time-dependent variable in a Cox-model, then `lex.Cst` should be used as the time-dependent variable, and `lex.Xst==1` as the event.

Chapter 3

Modeling rates

3.1 Background

The purpose of subdividing follow-up data is to be able to model the effects of the time scale variables as parametric functions.

In a model that assumes a constant occurrence rate in each of the intervals the likelihood contribution from each interval is the same as the likelihood contribution from a Poisson variate D , say, with mean $\lambda\ell$ where λ is the rate and ℓ is the interval length, and where the value of the variate D is 1 or 0 according to whether an event has occurred or not.

Moreover, the likelihood contributions from all follow-up intervals from a single person are *conditionally* independent (conditional on having survived till the start of the interval in question). This implies that the total contribution to the likelihood from a single person is a product of terms, and hence the same as the likelihood of a number of independent Poisson terms, one from each interval.

Parametric modeling of the rates is obtained by using the *value* of the timescale for each interval as quantitative explanatory variables, using for example splines. Thus the model will be one where the rate is assumed constant in each interval, but where a parametric form of the *size* of the rate in each interval is imposed by the model, using the timescale as a covariate.

3.2 Practicalities

In the nickel worker study we might want to look at the effects of age and time since hire. If we want to use splines we must allocate knots for anchoring the splines at each of the time scales, either by some *ad hoc* method or by using some sort of penalized splines. The latter will not be treated here.

Here we shall use the former approach and allocate 5 knots on each of the two time-scales. We allocate knots so that we have the event evenly distributed between the knots:

```
> ( a.kn <- with( subset( nicM, lex.Xst==1 ), quantile( age+lex.dur, (1:5-0.5)/5 ) ) )
      10%      30%      50%      70%      90%
50.11874 55.61674 61.09590 64.88704 73.32220
> ( t.kn <- with( subset( nicM, lex.Xst==1 ), quantile( tfh+lex.dur, (1:5-0.5)/5 ) ) )
      10%      30%      50%      70%      90%
24.25572 30.02202 34.00440 39.84592 45.95512
```

In the `Epi` package there is a convenience wrapper for the natural spline generator `ns`, `Ns`, that takes the smallest and the largest of a set of supplied knots to be the boundary knots.

3.3 Models for rates

3.3.1 One time scale

A model that only models lung cancer mortality rates as a function of age would then be:

```
> ma <- glm( (lex.Xst==1) ~ Ns(age,knots=a.kn),
+           family = poisson,
+           offset = log(lex.dur),
+           data = nicM )
> summary( ma )
```

Call:
`glm(formula = (lex.Xst == 1) ~ Ns(age, knots = a.kn), family = poisson, data = nicM, offset = log(lex.dur))`

Deviance Residuals:

Min	1Q	Median	3Q	Max
-0.5074	-0.3896	-0.2143	-0.1203	3.7904

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-4.6591	0.1324	-35.199	< 2e-16
Ns(age, knots = a.kn)1	0.1671	0.2970	0.563	0.57371
Ns(age, knots = a.kn)2	-0.1315	0.3727	-0.353	0.72411
Ns(age, knots = a.kn)3	0.7827	0.2885	2.713	0.00667
Ns(age, knots = a.kn)4	-0.3717	0.2780	-1.337	0.18125

(Dispersion parameter for poisson family taken to be 1)

Null deviance: 1024.38 on 3128 degrees of freedom
 Residual deviance: 979.16 on 3124 degrees of freedom
 AIC: 1263.2

Number of Fisher Scoring iterations: 7

The offset, `log(lex.dur)` comes from the fact that the likelihood for the follow-up data is the same as that for independent Poisson variates with mean $\lambda\ell$, and that the default link function for the Poisson family is the log, so that we are using a linear model for the log-mean, that is $\log(\lambda) + \log(\ell)$. But when we want a model for the log-rate ($\log(\lambda)$), the term $\log(\ell)$ must be included as a covariate with regression coefficient fixed to 1; a so-called offset.

The parameters from the model do not have any direct interpretation *per se*, but we can compute the estimated lung cancer incidence rates for a range of ages using `ci.pred` with a suitably defined prediction data frame. Note that we must specify all covariates in the model, also the variable in the offset, `lex.dur`. We set the latter to 1000, because we want the lung cancer mortality rates per 1000 PY. By default `ci.pred` yields a prediction on the response-scale, that is the rate-scale:

```
> nd <- data.frame( age=40:85, lex.dur=1000 )
> pr.a <- ci.pred( ma, newdata = nd )
```

```
> matplot( nd$age, pr.a,
+          type="l", lty=1, col=1, lwd=c(3,1,1),
+          log="y", xlab="Age (years)",
+          ylab="Lunng cancer mortality per 1000 PY")
```




Figure 3.1: *Lung cancer mortality among Nickel smelter workers by age. We see that the rates increase till about 55 years, and from then on is approximately flat.* `./flup-pr-a`

3.3.2 More time scales

There may however also be an effect of time since hire too, so we can add this term to the model:

```
> mat <- update( ma, . ~ . + Ns(tfh,knots=t.kn) )
> summary( mat )
```

Call:

```
glm(formula = (lex.Xst == 1) ~ Ns(age, knots = a.kn) + Ns(tfh,
  knots = t.kn), family = poisson, data = nicM, offset = log(lex.dur))
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-----	----	--------	----	-----

```
-0.6308 -0.3730 -0.2170 -0.1180 3.8903
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-4.65125	0.14844	-31.335	<2e-16
Ns(age, knots = a.kn)1	0.19029	0.32601	0.584	0.5594
Ns(age, knots = a.kn)2	0.04239	0.40857	0.104	0.9174
Ns(age, knots = a.kn)3	0.87848	0.37395	2.349	0.0188
Ns(age, knots = a.kn)4	0.08124	0.37567	0.216	0.8288
Ns(tfh, knots = t.kn)1	0.05961	0.45702	0.130	0.8962
Ns(tfh, knots = t.kn)2	-0.30254	0.39214	-0.771	0.4404
Ns(tfh, knots = t.kn)3	-0.08144	0.37493	-0.217	0.8281
Ns(tfh, knots = t.kn)4	-0.63400	0.34055	-1.862	0.0626

(Dispersion parameter for poisson family taken to be 1)

```
Null deviance: 1024.4 on 3128 degrees of freedom
Residual deviance: 970.7 on 3120 degrees of freedom
AIC: 1262.7
```

Number of Fisher Scoring iterations: 7

This model has two time scales, age and time since hire, so it makes little sense to report the effect of age for a *fixed* value of time since hire — the time since hire increases by age. Instead we can show the mortality rates for persons hired at different ages, and report the *joint* effect of increasing age and time since hire.

In order to get a feeling for the values that can be use we look at `age1st`

```
> summary( nickel$age1st )
   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
  10.78  21.80   26.16   26.74  30.63   52.19
```

Thus we shall show mortality rates in ages 20–90 for persons hired in ages 15, 25, 35 and 45:

```
> nd <- data.frame( expand.grid( age=c(20:90,NA), age1st=seq(15,45,10) ) )
> nd <- transform( nd, tfh = ifelse( age > age1st, age-age1st, NA ),
+                   lex.dur = 1000 )
> # makes no sense to have age < age1st
> nd <- transform( nd, age = ifelse( age > age1st, age, NA ) )
> head( nd )
  age age1st tfh lex.dur
1  20     15   5    1000
2  21     15   6    1000
3  22     15   7    1000
4  23     15   8    1000
5  24     15   9    1000
6  25     15  10    1000
```

With this in place we can plot the estimated rates as before, only now we will get 4 separate lines. The purpose of inserting an NA on the age-scale in the `expand.grid` is that the different instances of `age1st` be separated by NAs, and hence will not be connected when we plot the curves. The downside of this trick is that lines cannot be plotted with different colors or symbols.


```
> pr.at <- ci.pred( mat, newdata = nd )
> matplot( nd$age, pr.at,
+         type="l", lty=1, col=1, lwd=c(3,1,1),
+         log="y", xlab="Age (years)",
+         ylab="Lunng cancer mortality per 1000 PY")
```




Figure 3.2: *Lung cancer mortality among Nickel smelter workers by age and age at hire 15, 25, 35 and 45. Each line (except the first) starts at the age of hire; we see that the later in life you are hired, the smaller the initial risk, but the higher the eventual risk of lung cancer death.*

./flup-pr-at

We can check whether the effect of time since hire is actually improving the model:

```
> anova( ma, mat, test="Chisq" )
Analysis of Deviance Table

Model 1: (lex.Xst == 1) ~ Ns(age, knots = a.kn)
Model 2: (lex.Xst == 1) ~ Ns(age, knots = a.kn) + Ns(tfh, knots = t.kn)
   Resid. Df Resid. Dev Df Deviance Pr(>Chi)
1       3124      979.16
2       3120      970.70  4    8.4626  0.07603
```

We see a pretty strong indication that this is the case.

3.3.3 Difference between time scales

However it might be the case that it really is the age at first hire that is the main determinant (recall that `age - thf = age1st`), so we could fit a model with this variable instead — a model with only 1 timescale, namely `age`.

```
> ( f.kn <- with( subset( nicM, lex.Xst==1 ), quantile( age1st, (1:5-0.5)/5 ) ) )
      10%      30%      50%      70%      90%
20.25860 22.55422 26.00000 28.36578 33.96910

> maf <- update( ma, . ~ . + Ns(age1st,knots=f.kn) )
> summary( maf )
```

Call:

```
glm(formula = (lex.Xst == 1) ~ Ns(age, knots = a.kn) + Ns(age1st,
  knots = f.kn), family = poisson, data = nicM, offset = log(lex.dur))
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-0.5696	-0.3671	-0.2257	-0.1197	3.7777

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-4.62646	0.17564	-26.340	< 2e-16
Ns(age, knots = a.kn)1	0.21589	0.29742	0.726	0.46792
Ns(age, knots = a.kn)2	-0.06427	0.37653	-0.171	0.86446
Ns(age, knots = a.kn)3	0.79456	0.29345	2.708	0.00678
Ns(age, knots = a.kn)4	-0.31305	0.27976	-1.119	0.26314
Ns(age1st, knots = f.kn)1	-0.15145	0.38279	-0.396	0.69237
Ns(age1st, knots = f.kn)2	0.04607	0.27980	0.165	0.86923
Ns(age1st, knots = f.kn)3	0.26374	0.26156	1.008	0.31331
Ns(age1st, knots = f.kn)4	-0.22878	0.23117	-0.990	0.32234

(Dispersion parameter for poisson family taken to be 1)

```
Null deviance: 1024.4 on 3128 degrees of freedom
Residual deviance: 973.2 on 3120 degrees of freedom
AIC: 1265.2
```

Number of Fisher Scoring iterations: 7

```
> anova( maf, ma, mat, test="Chisq" )
```

Analysis of Deviance Table

```
Model 1: (lex.Xst == 1) ~ Ns(age, knots = a.kn) + Ns(age1st, knots = f.kn)
Model 2: (lex.Xst == 1) ~ Ns(age, knots = a.kn)
Model 3: (lex.Xst == 1) ~ Ns(age, knots = a.kn) + Ns(tfh, knots = t.kn)
  Resid. Df Resid. Dev Df Deviance Pr(>Chi)
1      3120      973.20
2      3124      979.16 -4   -5.9624  0.20198
3      3120      970.70  4    8.4626  0.07603
```

We see that there is much less indication that the age at first hire has an effect.

For the sake of completeness we can draw the predicted values from the `maf` model on top of the ones from the `mat` model:

```

> pr.af <- ci.pred( maf, newdata = nd )
> matplot( nd$age, pr.at,
+         type="l", lty=1, col=1, lwd=c(3,1,1),
+         log="y", xlab="Age (years)",
+         ylab="Lunng cancer mortality per 1000 PY")
> matlines( nd$age, pr.af,
+         type="l", lty=1, col=2, lwd=c(3,0,0) )

```




Figure 3.3: Lung cancer mortality among Nickel smelter workers by age and age at hire 15, 25, 35 and 45. Each line (except the first) starts at the age of hire; we see that the later in life you are hired, the smaller the initial risk, but the higher the eventual risk of lung cancer death. The red lines are from the model `maf` where the lines are constrained to be parallel, and which gives a worse fit to data. ./flup-pr-at-af

3.3.4 The complete picture — exercise

We could fit the remaining models where one or more of the three variables are included, and compare all of them:

```

> maft <- update( mat, . ~ . + Ns(age1st,knots=f.kn) )
> summary( maft )

```

Call:

```
glm(formula = (lex.Xst == 1) ~ Ns(age, knots = a.kn) + Ns(tfh,
  knots = t.kn) + Ns(agemst, knots = f.kn), family = poisson,
  data = nicM, offset = log(lex.dur))
```

Deviance Residuals:

	Min	1Q	Median	3Q	Max
	-0.5899	-0.3579	-0.2224	-0.1185	3.8687

Coefficients: (1 not defined because of singularities)

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-4.71537	0.16481	-28.612	<2e-16
Ns(age, knots = a.kn)1	0.01671	0.35152	0.048	0.9621
Ns(age, knots = a.kn)2	-0.11682	0.44638	-0.262	0.7935
Ns(age, knots = a.kn)3	0.47689	0.50638	0.942	0.3463
Ns(age, knots = a.kn)4	-0.18241	0.47318	-0.385	0.6999
Ns(tfh, knots = t.kn)1	0.35272	0.51329	0.687	0.4920
Ns(tfh, knots = t.kn)2	-0.11034	0.43043	-0.256	0.7977
Ns(tfh, knots = t.kn)3	0.26874	0.49133	0.547	0.5844
Ns(tfh, knots = t.kn)4	-0.30302	0.43585	-0.695	0.4869
Ns(agemst, knots = f.kn)1	-0.10650	0.37476	-0.284	0.7763
Ns(agemst, knots = f.kn)2	0.17245	0.20063	0.860	0.3900
Ns(agemst, knots = f.kn)3	0.47357	0.24239	1.954	0.0507
Ns(agemst, knots = f.kn)4	NA	NA	NA	NA

(Dispersion parameter for poisson family taken to be 1)

Null deviance: 1024.38 on 3128 degrees of freedom
 Residual deviance: 966.31 on 3117 degrees of freedom
 AIC: 1264.3

Number of Fisher Scoring iterations: 7

```
> mft <- update( maft, . ~ . - Ns(age,knots=a.kn) )
> mf <- update( maf , . ~ . - Ns(age,knots=a.kn) )
> mt <- update( mat , . ~ . - Ns(age,knots=a.kn) )
> allp <- anova( maft, mat, ma, maf, mf, mft, mt, mat,
+               maf, maft, mft,
+               test="Chisq" )
> mall <- as.matrix( allp )
> cbind( mod = c("maft","mat","ma","maf","mf","mft","mt","mat","maf","maft","mft"),
+        round( allp[,1:5], 3 ) )
```

	mod	Resid.	Df	Resid. Dev	Df	Deviance	Pr(>Chi)
1	maft	3117		966.306	NA	NA	NA
2	mat	3120		970.697	-3	-4.391	0.222
3	ma	3124		979.160	-4	-8.463	0.076
4	maf	3120		973.197	4	5.962	0.202
5	mf	3124		1011.593	-4	-38.396	0.000
6	mft	3120		971.120	4	40.473	0.000
7	mt	3124		985.734	-4	-14.614	0.006
8	mat	3120		970.697	4	15.037	0.005
9	maf	3120		973.197	0	-2.500	NA
10	maft	3117		966.306	3	6.892	0.075
11	mft	3120		971.120	-3	-4.814	0.186

1. Explain why there are NAs among the parameters in the model maf.

2. Draw a graph (a “DAG”) with the models as nodes and the tests as vertices, put the p-values on the vertices and use the result to argue that the model with age and time since hire is actually the most sensible description in this case.

Chapter 4

Competing risks — multiple types of events

If we want to consider death from lung cancer and death from other causes as separate events we can code these as for example 1 and 2.

```
> data( nickel )
> nicL <- Lexis( entry = list( per = agein+dob,
+                             age = agein,
+                             tfh = agein-age1st ),
+               exit = list( age = ageout ),
+               exit.status = ( icd > 0 ) + ( icd %in% c(162,163) ),
+               data = nickel )
> summary( nicL )
Transitions:
      To
From 0   1   2 Records: Events: Risk time: Persons:
      0 47 495 137      679      632  15348.06      679
> subset( nicL, id %in% 8:10 )
      per      age      tfh lex.dur lex.Cst lex.Xst lex.id id icd exposure      dob
4 1934.246 47.9067 23.1861 21.7727      0      1      4  8 527      9 1886.340
5 1934.246 54.7465 24.7890 22.0977      0      1      5  9 150      0 1879.500
6 1934.246 44.3314 23.0437 18.2099      0      2      6 10 163      2 1889.915
      age1st agein ageout
4 24.7206 47.9067 69.6794
5 29.9575 54.7465 76.8442
6 21.2877 44.3314 62.5413
```

In order to have a more readable output we can label the states, we can enter the names of these in the `states` parameter, try for example:

```
> nicL <- Lexis( entry = list( per = agein+dob,
+                             age = agein,
+                             tfh = agein-age1st ),
+               exit = list( age = ageout ),
+               exit.status = ( icd > 0 ) + ( icd %in% c(162,163) ),
+               data = nickel,
+               states = c("Alive", "D.oth", "D.lung") )
> summary( nicL )
Transitions:
      To
From      Alive D.oth D.lung Records: Events: Risk time: Persons:
      Alive      47  495  137      679      632  15348.06      679
```

```
> str( nicL )

Classes 'Lexis' and 'data.frame':      679 obs. of  14 variables:
 $ per      : num  1934 1934 1934 1934 1934 ...
 $ age      : num  45.2 48.3 53 47.9 54.7 ...
 $ tfh      : num  27.7 25.1 27.7 23.2 24.8 ...
 $ lex.dur  : num  47.75 15 1.17 21.77 22.1 ...
 $ lex.Cst  : Factor w/ 3 levels "Alive","D.oth",...: 1 1 1 1 1 1 1 1 1 1 ...
 $ lex.Xst  : Factor w/ 3 levels "Alive","D.oth",...: 1 3 3 2 2 3 2 2 2 2 ...
 $ lex.id   : int   1 2 3 4 5 6 7 8 9 10 ...
 $ id       : num   3 4 6 8 9 10 15 16 17 18 ...
 $ icd      : num   0 162 163 527 150 163 334 160 420 12 ...
 $ exposure: num   5 5 10 9 0 2 0 0.5 0 0 ...
 $ dob      : num  1889 1886 1881 1886 1880 ...
 $ age1st   : num  17.5 23.2 25.2 24.7 30 ...
 $ agein    : num  45.2 48.3 53 47.9 54.7 ...
 $ ageout   : num  93 63.3 54.2 69.7 76.8 ...
 - attr(*, "time.scales")= chr  "per" "age" "tfh"
 - attr(*, "time.since")= chr  "" "" ""
 - attr(*, "breaks")=List of 3
 ..$ per: NULL
 ..$ age: NULL
 ..$ tfh: NULL
```

Note that the `Lexis` function automatically assumes that all persons enter in the first level (given in the `states=` argument), corresponding to the numerical values given in `exit.status`.

When we cut at a date as in this case, the date where cumulative exposure exceeds 50 exposure-years, we get the follow-up *after* the date classified as being in the new state if the exit (`lex.Xst`) was to a state we defined as one of the `precursor.states`:

```
> nicL$agehi <- nicL$age1st + 50 / nicL$exposure
> nicC <- cutLexis( data = nicL,
+                   cut = nicL$agehi,
+                   timescale = "age",
+                   new.state = "HiExp",
+                   precursor.states = "Alive" )
> subset( nicC, id %in% 8:10 )

   per    age    tfh lex.dur lex.Cst lex.Xst lex.id id icd exposure    dob
683 1934.246 47.9067 23.1861 21.7727  HiExp  D.oth     4  8 527         9 1886.340
5   1934.246 54.7465 24.7890 22.0977  Alive  D.oth     5  9 150         0 1879.500
6   1934.246 44.3314 23.0437  1.9563  Alive  HiExp     6 10 163         2 1889.915
685 1936.203 46.2877 25.0000 16.2536  HiExp  D.lung    6 10 163         2 1889.915
   age1st agein ageout agehi
683 24.7206 47.9067 69.6794 30.27616
5   29.9575 54.7465 76.8442      Inf
6   21.2877 44.3314 62.5413 46.28770
685 21.2877 44.3314 62.5413 46.28770

> summary( nicC, scale=1000 )

Transitions:
  To
From  Alive HiExp D.oth D.lung  Records:  Events: Risk time:  Persons:
  Alive    39   83   279    65      466     427    10.77      466
  HiExp     0    8   216    72      296     288     4.58      296
  Sum      39   91   495   137      762     715    15.35      679
```

Note that the persons-years is the same, but that the number of events has changed. This is because events are now defined as any transition, including the transitions to HiExp.

Also note that (so far) it is necessary to specify the variable with the cut points in full, using only `cut=agehi` would give an error.

4.1 Subdividing states

It may be of interest to subdivide the states following the intermediate event according to whether the event has occurred or not. That is done by the argument `split.states=TRUE`.

Moreover, it will also often be of interest to introduce a new timescale indicating the time since intermediate event. This can be done by the argument `new.scale=TRUE`, alternatively `new.scale="tfe"`, as illustrated here:

```
> nicC <- cutLexis( data = nicL,
+                   cut = nicL$agehi,
+                   timescale = "age",
+                   new.state = "HiExp",
+                   new.scale = "tfe",
+                   split.states = TRUE,
+                   precursor.states = "Alive" )
> subset( nicC, id %in% 8:10 )
```

	per	age	tfh	tfe	lex.dur	lex.Cst	lex.Xst	lex.id	id	icd
683	1934.246	47.9067	23.1861	17.63054	21.7727	HiExp	D.oth(HiExp)	4	8	527
5	1934.246	54.7465	24.7890	NA	22.0977	Alive	D.oth	5	9	150
6	1934.246	44.3314	23.0437	NA	1.9563	Alive	HiExp	6	10	163
685	1936.203	46.2877	25.0000	0.00000	16.2536	HiExp	D.lung(HiExp)	6	10	163
	exposure	dob	age1st	agein	ageout	agehi				
683	9	1886.340	24.7206	47.9067	69.6794	30.27616				
5	0	1879.500	29.9575	54.7465	76.8442	Inf				
6	2	1889.915	21.2877	44.3314	62.5413	46.28770				
685	2	1889.915	21.2877	44.3314	62.5413	46.28770				

```
> summary( nicC, scale=1000, timeScales=TRUE )
```

Transitions:

To								Records:	Events:	Risk	time:
From	Alive	HiExp	D.oth	D.lung	D.lung(HiExp)	D.oth(HiExp)					
Alive	39	83	279	65	0	0	466	427	10.77		
HiExp	0	8	0	0	72	216	296	288	4.58		
Sum	39	91	279	65	72	216	762	715	15.35		

Transitions:

To			
From	Persons:		
Alive	466		
HiExp	296		
Sum	679		

Timescales:

	time.scale	time.since
1	per	
2	age	
3	tfh	
4	tfe	HiExp

Note that the `timeScales=TRUE` to `summary` lists the timescales available in the object, and also indicates which of them that are defined as time since entry to a particular state. This facility is not used here, but it is needed when simulating follow-up data — see the vignette on `simLexis`.

With 6 different states it is quite difficult to get an overview of the transitions between states from the `summary()`. Therefore there is function that gives a graphical display of the states showing the transitions between the states:

```
> boxes( nicC, boxpos = list(x=c(10,10,80,80,80,80),
+                             y=c(75,25,87,63,13,37)),
+       scale.Y = 1000,
+       show.BE = TRUE )
```

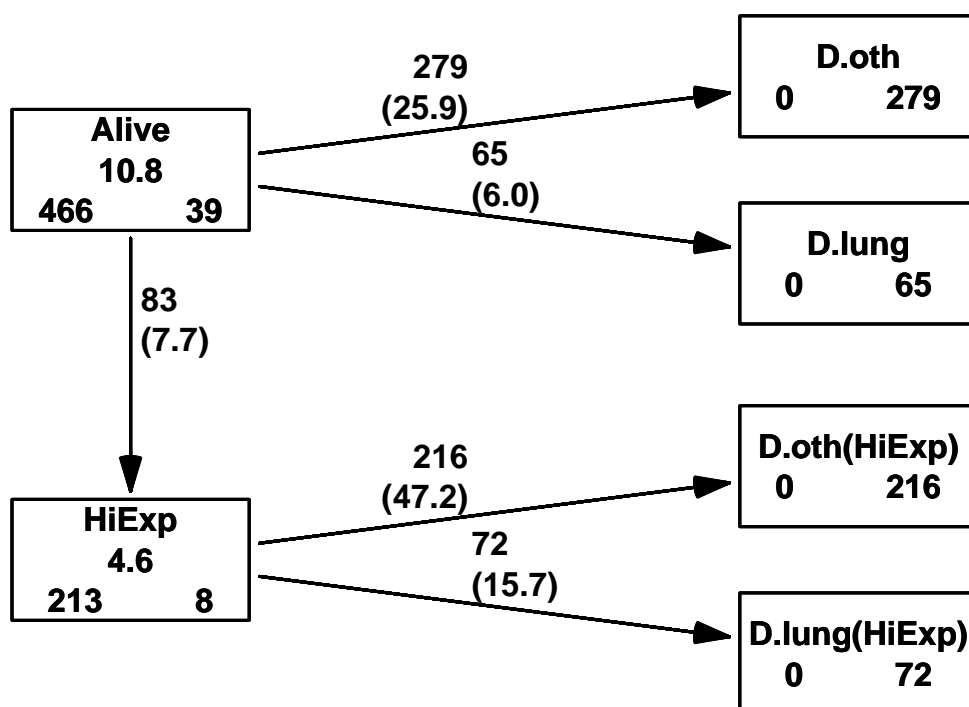


Figure 4.1: Transitions between states; the number in the middle of each box is the person-years (in 1000s — since `scale.Y=1000`), the numbers at the bottom of the boxes are the number that start, respectively end their follow-up in each state. The numbers on the arrows are the number of transitions and crude transition rates (the latter in events per 1000 PY). The function `boxes.Lexis` has a zillion arguments to fine-tune the appearance of the display in terms of colors etc.

./flup-nic-box