

Follow-up data with the Epi package

November 2017

Bendix Carstensen Steno Diabetes Center Copenhagen, Gentofte, Denmark
& Department of Biostatistics, University of Copenhagen
b@bxc.dk
<http://BendixCarstensen.com>

Contents

1	Follow-up data in the Epi package	1
2	Timescales	1
3	Splitting the follow-up time along a timescale	4
4	Splitting (cutting) time at a specific date	8
5	Competing risks — multiple types of events	11

```
> library(Epi)
> print( sessionInfo(), l=F )
```

```
R version 3.4.2 (2017-09-28)
Platform: x86_64-pc-linux-gnu (64-bit)
Running under: Ubuntu 14.04.5 LTS
```

```
Matrix products: default
BLAS: /usr/lib/openblas-base/libopenblas.so.0
LAPACK: /usr/lib/lapack/liblapack.so.3.0
```

```
attached base packages:
```

```
[1] utils      datasets  graphics  grDevices  stats      methods    base
```

```
other attached packages:
```

```
[1] Epi_2.22
```

```
loaded via a namespace (and not attached):
```

```
[1] cmprsk_2.2-7      zoo_1.8-0          MASS_7.3-47        compiler_3.4.2      Matrix_1.2-10
[6] plyr_1.8.4        parallel_3.4.2     survival_2.41-3    etm_0.6-2           Rcpp_0.12.0
[11] splines_3.4.2     grid_3.4.2         numDeriv_2016.8-1  lattice_0.20-35
```

1 Follow-up data in the Epi package

In the *Epi*-package, follow-up data is represented by adding some extra variables to a data frame. Such a data frame is called a *Lexis* object. The tools for handling follow-up data then use the structure of this for special plots, tabulations etc.

Follow-up data basically consists of a time of entry, a time of exit and an indication of the status at exit (normally either “alive” or “dead”). Implicitly is also assumed a status *during* the follow-up (usually “alive”).

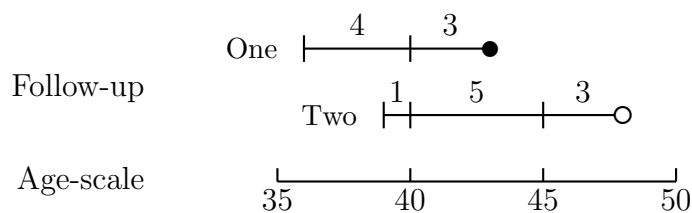


Figure 1: *Follow-up of two persons*

2 Timescales

A timescale is a variable that varies deterministically *within* each person during follow-up, *e.g.*:

- Age
- Calendar time
- Time since treatment
- Time since relapse

All timescales advance at the same pace, so the time followed is the same on all timescales. Therefore, it suffices to use only the entry point on each of the time scale, for example:

- Age at entry.
- Date of entry.
- Time since treatment (*at* treatment this is 0).
- Time since relapse (*at* relapse this is 0)..

In the `Epi` package, follow-up in a cohort is represented in a `Lexis` object. A `Lexis` object is a data frame with a bit of extra structure representing the follow-up. For the `nickel` data we would construct a `Lexis` object by:

```
> data( nickel )
> nicL <- Lexis( entry = list( per=agein+dob,
+                               age=agein,
+                               tfh=agein-age1st ),
+               exit = list( age=ageout ),
+               exit.status = ( icd %in% c(162,163) )*1,
+               data = nickel )
```

The `entry` argument is a *named* list with the entry points on each of the timescales we want to use. It defines the names of the timescales and the entry points. The `exit` argument gives the exit time on *one* of the timescales, so the name of the element in this list must match one of the names of the `entry` list. This is sufficient, because the follow-up time on all time scales is the same, in this case `ageout - agein`. Now take a look at the result:

```
> str( nickel )

'data.frame':      679 obs. of  7 variables:
 $ id      : num  3 4 6 8 9 10 15 16 17 18 ...
 $ icd     : num  0 162 163 527 150 163 334 160 420 12 ...
 $ exposure: num  5 5 10 9 0 2 0 0.5 0 0 ...
 $ dob     : num  1889 1886 1881 1886 1880 ...
 $ age1st  : num  17.5 23.2 25.2 24.7 30 ...
 $ agein   : num  45.2 48.3 53 47.9 54.7 ...
 $ ageout  : num  93 63.3 54.2 69.7 76.8 ...
```

```
> str( nicL )
```

```
Classes 'Lexis' and 'data.frame':      679 obs. of  14 variables:
```

```
$ per      : num  1934 1934 1934 1934 1934 ...
$ age      : num  45.2 48.3 53 47.9 54.7 ...
$ tfh      : num  27.7 25.1 27.7 23.2 24.8 ...
$ lex.dur  : num  47.75 15 1.17 21.77 22.1 ...
$ lex.Cst  : num  0 0 0 0 0 0 0 0 0 0 ...
$ lex.Xst  : num  0 1 1 0 0 1 0 0 0 0 ...
$ lex.id   : int   1 2 3 4 5 6 7 8 9 10 ...
$ id       : num   3 4 6 8 9 10 15 16 17 18 ...
$ icd      : num   0 162 163 527 150 163 334 160 420 12 ...
$ exposure: num   5 5 10 9 0 2 0 0.5 0 0 ...
$ dob      : num  1889 1886 1881 1886 1880 ...
$ age1st   : num  17.5 23.2 25.2 24.7 30 ...
$ agein    : num  45.2 48.3 53 47.9 54.7 ...
$ ageout   : num   93 63.3 54.2 69.7 76.8 ...
- attr(*, "time.scales")= chr  "per" "age" "tfh"
- attr(*, "time.since")= chr  "" "" ""
- attr(*, "breaks")=List of 3
..$ per: NULL
..$ age: NULL
..$ tfh: NULL
```

```
> head( nicL )
```

	per	age	tfh	lex.dur	lex.Cst	lex.Xst	lex.id	id	icd	exposure	dob	a
1	1934.246	45.2273	27.7465	47.7535	0	0	1	3	0	5	1889.019	17
2	1934.246	48.2684	25.0820	15.0028	0	1	2	4	162	5	1885.978	23
3	1934.246	52.9917	27.7465	1.1727	0	1	3	6	163	10	1881.255	25
4	1934.246	47.9067	23.1861	21.7727	0	0	4	8	527	9	1886.340	24
5	1934.246	54.7465	24.7890	22.0977	0	0	5	9	150	0	1879.500	29
6	1934.246	44.3314	23.0437	18.2099	0	1	6	10	163	2	1889.915	21
	ageout											
1	92.9808											
2	63.2712											
3	54.1644											
4	69.6794											
5	76.8442											
6	62.5413											

The *Lexis* object *nicL* has a variable for each timescale which is the entry point on this timescale. The follow-up time is in the variable *lex.dur* (**d**uration).

There is a *summary* function for *Lexis* objects that list the number of transitions and records as well as the total follow-up time:

```
> summary( nicL )
```

Transitions:

	To					
From	0	1	Records:	Events:	Risk time:	Persons:
	0	542 137	679	137	15348.06	679

We defined the exit status to be death from lung cancer (ICD7 162,163), i.e. this variable is 1 if follow-up ended with a death from this cause. If follow-up ended alive or by death from another cause, the exit status is coded 0, i.e. as a censoring.

Note that the exit status is in the variable `lex.Xst` (`eXit status`). The variable `lex.Cst` is the state where the follow-up takes place (`C`urrent `s`tatus), in this case 0 (alive).

It is possible to get a visualization of the follow-up along the timescales chosen by using the `plot` method for `Lexis` objects. `nicL` is an object of *class* `Lexis`, so using the function `plot()` on it means that **R** will look for the function `plot.Lexis` and use this function.

```
> plot( nicL )
```

The function allows a lot of control over the output, and a `points.Lexis` function allows plotting of the endpoints of follow-up:

```
> par( mar=c(3,3,1,1), mgp=c(3,1,0)/1.6 )
> plot( nicL, 1:2, lwd=1, col=c("blue","red")[(nicL$exp>0)+1],
+       grid=TRUE, lty.grid=1, col.grid=gray(0.7),
+       xlim=1900+c(0,90), xaxs="i",
+       ylim= 10+c(0,90), yaxs="i", las=1 )
> points( nicL, 1:2, pch=c(NA,3)[nicL$lex.Xst+1],
+         col="lightgray", lwd=3, cex=1.5 )
> points( nicL, 1:2, pch=c(NA,3)[nicL$lex.Xst+1],
+         col=c("blue","red")[(nicL$exp>0)+1], lwd=1, cex=1.5 )
```

The results of these two plotting commands are in figure 2.

3 Splitting the follow-up time along a timescale

The follow-up time in a cohort can be subdivided by for example current age. This is achieved by the `splitLexis` (note that it is *not* called `split.Lexis`). This requires that the timescale and the breakpoints on this timescale are supplied. Try:

```
> nicS1 <- splitLexis( nicL, "age", breaks=seq(0,100,10) )
> summary( nicL )
```

Transitions:

	To					
From	0	1	Records:	Events:	Risk time:	Persons:
	0	542 137	679	137	15348.06	679

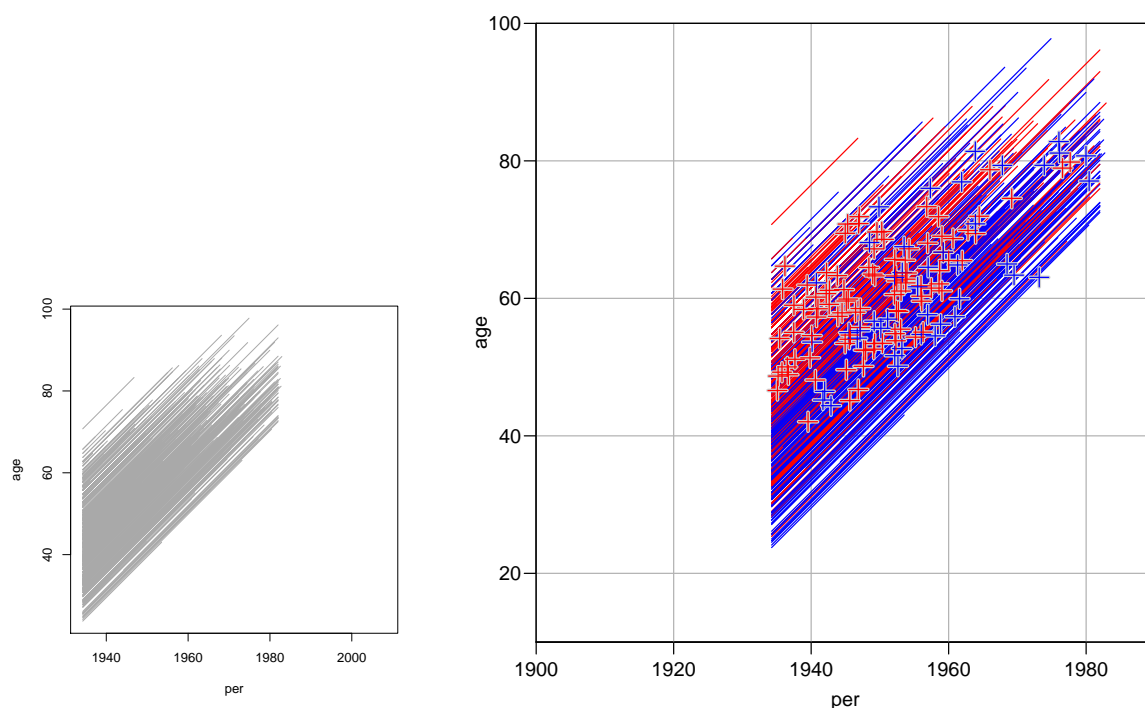


Figure 2: Lexis diagram of the *nickel* dataset, left panel the default version, the right one with bells and whistles. The red lines are for persons with $\text{exposure} > 0$, so it is pretty evident that the oldest ones are the exposed part of the cohort.

```
> summary( nicS1 )
```

Transitions:

To

```
From 0 1 Records: Events: Risk time: Persons:
0 2073 137 2210 137 15348.06 679
```

So we see that the number of events and the amount of follow-up is the same in the two data sets; only the number of records differ.

To see how records are split for each individual, it is useful to list the results for a few individuals:

```
> round( subset( nicS1, id %in% 8:10 ), 2 )
```

	lex.id	per	age	tfh	lex.dur	lex.Cst	lex.Xst	id	icd	exposure	dob	age1st
11	4	1934.25	47.91	23.19	2.09	0	0	8	527	9	1886.34	24.72
12	4	1936.34	50.00	25.28	10.00	0	0	8	527	9	1886.34	24.72
13	4	1946.34	60.00	35.28	9.68	0	0	8	527	9	1886.34	24.72
14	5	1934.25	54.75	24.79	5.25	0	0	9	150	0	1879.50	29.96
15	5	1939.50	60.00	30.04	10.00	0	0	9	150	0	1879.50	29.96
16	5	1949.50	70.00	40.04	6.84	0	0	9	150	0	1879.50	29.96
17	6	1934.25	44.33	23.04	5.67	0	0	10	163	2	1889.91	21.29

```

18      6 1939.91 50.00 28.71   10.00      0      0 10 163      2 1889.91 21.29
19      6 1949.91 60.00 38.71    2.54      0      1 10 163      2 1889.91 21.29

```

The resulting object, `nicS1`, is again a `Lexis` object, and so follow-up may be split further along another timescale. Try this and list the results for individuals 8, 9 and 10 again:

```

> nicS2 <- splitLexis( nicS1, "tfh", breaks=c(0,1,5,10,20,30,100) )
> round( subset( nicS2, id %in% 8:10 ), 2 )

```

	lex.id	per	age	tfh	lex.dur	lex.Cst	lex.Xst	id	icd	exposure	dob	age1st
13	4	1934.25	47.91	23.19	2.09	0	0	8	527	9	1886.34	24.72
14	4	1936.34	50.00	25.28	4.72	0	0	8	527	9	1886.34	24.72
15	4	1941.06	54.72	30.00	5.28	0	0	8	527	9	1886.34	24.72
16	4	1946.34	60.00	35.28	9.68	0	0	8	527	9	1886.34	24.72
17	5	1934.25	54.75	24.79	5.21	0	0	9	150	0	1879.50	29.96
18	5	1939.46	59.96	30.00	0.04	0	0	9	150	0	1879.50	29.96
19	5	1939.50	60.00	30.04	10.00	0	0	9	150	0	1879.50	29.96
20	5	1949.50	70.00	40.04	6.84	0	0	9	150	0	1879.50	29.96
21	6	1934.25	44.33	23.04	5.67	0	0	10	163	2	1889.91	21.29
22	6	1939.91	50.00	28.71	1.29	0	0	10	163	2	1889.91	21.29
23	6	1941.20	51.29	30.00	8.71	0	0	10	163	2	1889.91	21.29
24	6	1949.91	60.00	38.71	2.54	0	1	10	163	2	1889.91	21.29

A more efficient way of making this double split is to use the `splitMulti` function from the `popEpi` package:

```

> library( popEpi )
> nicM <- splitMulti( nicL, age = seq(0,100,10),
+                    tfh = c(0,1,5,10,20,30,100) )
> summary( nicS2 )

```

Transitions:

```

      To
From    0    1  Records:  Events: Risk time:  Persons:
      0 2992 137      3129      137  15348.06      679

```

```

> summary( nicM )

```

Transitions:

```

      To
From    0    1  Records:  Events: Risk time:  Persons:
      0 2992 137      3129      137  15348.06      679

```

So we see that the two ways of splitting data yields the same amount of follow-up, but the results are not identical:

```

> identical( nicS2, nicM )

```



```
[1] FALSE
> class( nicS2 )
[1] "Lexis"      "data.frame"
> class( nicM )
[1] "Lexis"      "data.table" "data.frame"
```

As we see, this is because the `nicM` object also is a `data.table` object; the `splitMulti` uses the `data.table` machinery which makes the splitting substantially more efficient — this is of particular interest if you operate on large data sets ($> 1,000,000$ records).

Time scales as covariates

If we want to model the effect of these timescales we will for each interval use either the value of the left endpoint in each interval or the middle. There is a function `timeBand` which returns these. Try:

```
> timeBand( nicM, "age", "middle" )[1:20]
[1] 45 45 55 65 75 85 95 45 55 55 65 55 45 55 55 65 55 55 65 75
> # For nice printing and column labelling use the data.frame() function:
> data.frame( nicS2[,c("id","lex.id","per","age","tfh","lex.dur")],
+           mid.age=timeBand( nicS2, "age", "middle" ),
+           mid.tfh=timeBand( nicS2, "tfh", "middle" ) )[1:20,]
```

	id	lex.id	per	age	tfh	lex.dur	mid.age	mid.tfh
1	3	1	1934.246	45.2273	27.7465	2.2535	45	25
2	3	1	1936.500	47.4808	30.0000	2.5192	45	65
3	3	1	1939.019	50.0000	32.5192	10.0000	55	65
4	3	1	1949.019	60.0000	42.5192	10.0000	65	65
5	3	1	1959.019	70.0000	52.5192	10.0000	75	65
6	3	1	1969.019	80.0000	62.5192	10.0000	85	65
7	3	1	1979.019	90.0000	72.5192	2.9808	95	65
8	4	2	1934.246	48.2684	25.0820	1.7316	45	25
9	4	2	1935.978	50.0000	26.8136	3.1864	55	25
10	4	2	1939.164	53.1864	30.0000	6.8136	55	65
11	4	2	1945.978	60.0000	36.8136	3.2712	65	65
12	6	3	1934.246	52.9917	27.7465	1.1727	55	25
13	8	4	1934.246	47.9067	23.1861	2.0933	45	25
14	8	4	1936.340	50.0000	25.2794	4.7206	55	25
15	8	4	1941.060	54.7206	30.0000	5.2794	55	65
16	8	4	1946.340	60.0000	35.2794	9.6794	65	65
17	9	5	1934.246	54.7465	24.7890	5.2110	55	25
18	9	5	1939.457	59.9575	30.0000	0.0425	55	65
19	9	5	1939.500	60.0000	30.0425	10.0000	65	65
20	9	5	1949.500	70.0000	40.0425	6.8442	75	65

Note that these are the midpoints of the intervals defined by `breaks=`, *not* the midpoints of the actual follow-up intervals. This is because the variable to be used in modeling must be independent of the censoring and mortality pattern — it should only depend on the chosen grouping of the timescale.

Difference between time scales

However, the midpoint should be used with caution if the variable `age1st` is modeled too; the age at hire is logically equal to the difference between current age (`age`) and time since hire (`thf`):

```
> summary( (nicS2$age-nicS2$tfh) - nicS2$age1st )
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
-7.105e-15	0.000e+00	0.000e+00	2.214e-17	0.000e+00	7.105e-15

This calculation refer to the *start* of each interval. But when using the middle of the intervals, this relationship is not preserved:

```
> summary( timeBand( nicS2, "age", "middle" ) -  
+          timeBand( nicS2, "tfh", "middle" ) - nicS2$age1st )
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
-39.958	-24.178	-5.103	-10.129	2.575	12.519

If all three variable are to be included in a model, you must make sure that the *substantial* relationship between the variables be maintained. One way is to recompute age at first hire from the two midpoint variables, but more straightforward is to use the left end point of the intervals, that is the time scales in the `Lexis` object. The latter approach requires that the follow-up is split in fairly small chunks.

4 Splitting (cutting) time at a specific date

If we have a recording of the date of a specific event as for example recovery or relapse, we may classify follow-up time as being before of after this intermediate event. This is achieved with the function `cutLexis`, which takes three arguments: the time point, the timescale, and the value of the (new) state following the date.

Now we define the age for the nickel workers where the cumulative exposure exceeds 50 exposure years:

```
> subset( nicL, id %in% 8:10 )
```

	per	age	tfh	lex.dur	lex.Cst	lex.Xst	lex.id	id	icd	exposure	dob	a
4	1934.246	47.9067	23.1861	21.7727	0	0	4	8	527	9	1886.340	24
5	1934.246	54.7465	24.7890	22.0977	0	0	5	9	150	0	1879.500	29
6	1934.246	44.3314	23.0437	18.2099	0	1	6	10	163	2	1889.915	21

```

      ageout
4 69.6794
5 76.8442
6 62.5413

```

```

> agehi <- nicL$age1st + 50 / nicL$exposure
> nicC <- cutLexis( data = nicL,
+                   cut = agehi,
+                   timescale = "age",
+                   new.state = 2,
+                   precursor.states = 0 )
> subset( nicC, id %in% 8:10 )

```

	per	age	tfh	lex.dur	lex.Cst	lex.Xst	lex.id	id	icd	exposure	dob
683	1934.246	47.9067	23.1861	21.7727	2	2	4	8	527	9	1886.340
5	1934.246	54.7465	24.7890	22.0977	0	0	5	9	150	0	1879.500
6	1934.246	44.3314	23.0437	1.9563	0	2	6	10	163	2	1889.915
685	1936.203	46.2877	25.0000	16.2536	2	1	6	10	163	2	1889.915

	agein	ageout
683	47.9067	69.6794
5	54.7465	76.8442
6	44.3314	62.5413
685	44.3314	62.5413

(The `precursor.states=` argument is explained below). Note that individual 6 has had his follow-up split at 25 years since hire where 50 exposure-years were attained. This could also have been achieved in the split dataset `nicS2` instead of `nicL`, try:

```

> subset( nicS2, id %in% 8:10 )

```

	lex.id	per	age	tfh	lex.dur	lex.Cst	lex.Xst	id	icd	exposure	dob
13	4	1934.246	47.9067	23.1861	2.0933	0	0	8	527	9	1886.340
14	4	1936.340	50.0000	25.2794	4.7206	0	0	8	527	9	1886.340
15	4	1941.060	54.7206	30.0000	5.2794	0	0	8	527	9	1886.340
16	4	1946.340	60.0000	35.2794	9.6794	0	0	8	527	9	1886.340
17	5	1934.246	54.7465	24.7890	5.2110	0	0	9	150	0	1879.500
18	5	1939.457	59.9575	30.0000	0.0425	0	0	9	150	0	1879.500
19	5	1939.500	60.0000	30.0425	10.0000	0	0	9	150	0	1879.500
20	5	1949.500	70.0000	40.0425	6.8442	0	0	9	150	0	1879.500
21	6	1934.246	44.3314	23.0437	5.6686	0	0	10	163	2	1889.915
22	6	1939.915	50.0000	28.7123	1.2877	0	0	10	163	2	1889.915
23	6	1941.203	51.2877	30.0000	8.7123	0	0	10	163	2	1889.915
24	6	1949.915	60.0000	38.7123	2.5413	0	1	10	163	2	1889.915

	ageout
13	69.6794
14	69.6794

```

15 69.6794
16 69.6794
17 76.8442
18 76.8442
19 76.8442
20 76.8442
21 62.5413
22 62.5413
23 62.5413
24 62.5413

```

```

> agehi <- nicS2$age1st + 50 / nicS2$exposure
> nicS2C <- cutLexis( data=nicS2, cut=agehi, timescale="age",
+                     new.state=2, precursor.states=0 )
> subset( nicS2C, id %in% 8:10 )

```

	lex.id	per	age	tfh	lex.dur	lex.Cst	lex.Xst	id	icd	exposure	dob
3142	4	1934.246	47.9067	23.1861	2.0933	2	2	8	527	9	1886.340
3143	4	1936.340	50.0000	25.2794	4.7206	2	2	8	527	9	1886.340
3144	4	1941.060	54.7206	30.0000	5.2794	2	2	8	527	9	1886.340
3145	4	1946.340	60.0000	35.2794	9.6794	2	2	8	527	9	1886.340
17	5	1934.246	54.7465	24.7890	5.2110	0	0	9	150	0	1879.500
18	5	1939.457	59.9575	30.0000	0.0425	0	0	9	150	0	1879.500
19	5	1939.500	60.0000	30.0425	10.0000	0	0	9	150	0	1879.500
20	5	1949.500	70.0000	40.0425	6.8442	0	0	9	150	0	1879.500
21	6	1934.246	44.3314	23.0437	1.9563	0	2	10	163	2	1889.915
3150	6	1936.203	46.2877	25.0000	3.7123	2	2	10	163	2	1889.915
3151	6	1939.915	50.0000	28.7123	1.2877	2	2	10	163	2	1889.915
3152	6	1941.203	51.2877	30.0000	8.7123	2	2	10	163	2	1889.915
3153	6	1949.915	60.0000	38.7123	2.5413	2	1	10	163	2	1889.915

	agein	ageout
3142	47.9067	69.6794
3143	47.9067	69.6794
3144	47.9067	69.6794
3145	47.9067	69.6794
17	54.7465	76.8442
18	54.7465	76.8442
19	54.7465	76.8442
20	54.7465	76.8442
21	44.3314	62.5413
3150	44.3314	62.5413
3151	44.3314	62.5413
3152	44.3314	62.5413
3153	44.3314	62.5413

Note that follow-up subsequent to the event is classified as being in state 2, but that the final transition to state 1 (death from lung cancer) is preserved. This is the point of the

`precursor.states=` argument. It names the states (in this case 0, “Alive”) that will be over-written by `new.state` (in this case state 2, “High exposure”). Clearly, state 1 (“Dead”) should not be updated even if it is after the time where the persons moves to state 2. In other words, only state 0 is a precursor to state 2, state 1 is always subsequent to state 2.

Note that if the intermediate event is to be used as a time-dependent variable in a Cox-model, then `lex.Cst` should be used as the time-dependent variable, and `lex.Xst==1` as the event.

5 Competing risks — multiple types of events

If we want to consider death from lung cancer and death from other causes as separate events we can code these as for example 1 and 2.

```
> data( nickel )
> nicL <- Lexis( entry = list( per = agein+dob,
+                               age = agein,
+                               tfh = agein-age1st ),
+               exit = list( age = ageout ),
+               exit.status = ( icd > 0 ) + ( icd %in% c(162,163) ),
+               data = nickel )
> summary( nicL )
```

Transitions:

	To	
From 0	1 2	Records: Events: Risk time: Persons:
	0 47 495 137	679 632 15348.06 679

```
> subset( nicL, id %in% 8:10 )
```

	per	age	tfh	lex.dur	lex.Cst	lex.Xst	lex.id	id	icd	exposure	dob	a
4	1934.246	47.9067	23.1861	21.7727	0	1	4	8	527	9	1886.340	24
5	1934.246	54.7465	24.7890	22.0977	0	1	5	9	150	0	1879.500	29
6	1934.246	44.3314	23.0437	18.2099	0	2	6	10	163	2	1889.915	21
	ageout											
4	69.6794											
5	76.8442											
6	62.5413											

In order to have a more readable output we can label the states, we can enter the names of these in the `states` parameter, try for example:

```
> nicL <- Lexis( entry = list( per = agein+dob,
+                               age = agein,
+                               tfh = agein-age1st ),
+               exit = list( age = ageout ),
```

```
+          exit.status = ( icd > 0 ) + ( icd %in% c(162,163) ),
+          data = nickel,
+          states = c("Alive", "D.oth", "D.lung") )
> summary( nicL )
```

Transitions:

	To							
From	Alive	D.oth	D.lung	Records:	Events:	Risk time:	Persons:	
	47	495	137	679	632	15348.06	679	

Note that the Lexis function automatically assumes that all persons enter in the first level (given in the `states=` argument)

When we cut at a date as in this case, the date where cumulative exposure exceeds 50 exposure-years, we get the follow-up *after* the date classified as being in the new state if the exit (`lex.Xst`) was to a state we defined as one of the `precursor.states`:

```
> nicL$agehi <- nicL$age1st + 50 / nicL$exposure
> nicC <- cutLexis( data = nicL,
+                  cut = nicL$agehi,
+                  timescale = "age",
+                  new.state = "HiExp",
+                  precursor.states = "Alive" )
> subset( nicC, id %in% 8:10 )
```

	per	age	tfh	lex.dur	lex.Cst	lex.Xst	lex.id	id	icd	exposure	dob
683	1934.246	47.9067	23.1861	21.7727	HiExp	D.oth	4	8	527	9	1886.340
5	1934.246	54.7465	24.7890	22.0977	Alive	D.oth	5	9	150	0	1879.500
6	1934.246	44.3314	23.0437	1.9563	Alive	HiExp	6	10	163	2	1889.915
685	1936.203	46.2877	25.0000	16.2536	HiExp	D.lung	6	10	163	2	1889.915
	agein	ageout	agehi								
683	47.9067	69.6794	30.27616								
5	54.7465	76.8442	Inf								
6	44.3314	62.5413	46.28770								
685	44.3314	62.5413	46.28770								

```
> summary( nicC, scale=1000 )
```

Transitions:

	To							
From	Alive	HiExp	D.oth	D.lung	Records:	Events:	Risk time:	Persons:
Alive	39	83	279	65	466	427	10.77	466
HiExp	0	8	216	72	296	288	4.58	296
Sum	39	91	495	137	762	715	15.35	679

Note that the persons-years is the same, but that the number of events has changed. This is because events are now defined as any transition from alive, including the transitions to HiExp.

Also note that (so far) it is necessary to specify the variable with the cutpoints in full, using only `cut=agehi` would give an error.

Subdivision of existing states

It may be of interest to subdivide the states following the intermediate event according to whether the event has occurred or not. That is done by the argument `split.states=TRUE`.

Moreover, it will also often be of interest to introduce a new timescale indicating the time since intermediate event. This can be done by the argument `new.scale=TRUE`, alternatively `new.scale="tfevent"`, as illustrated here:

```
> nicC <- cutLexis( data = nicL,
+                   cut = nicL$agehi,
+                   timescale = "age",
+                   new.state = "HiExp",
+                   new.scale = TRUE,
+                   split.states = TRUE,
+                   precursor.states = "Alive" )
> subset( nicC, id %in% 8:10 )
```

	per	age	tfh	tfHiExp	lex.dur	lex.Cst	lex.Xst	lex.id	id	icd	exp
683	1934.246	47.9067	23.1861	17.63054	21.7727	HiExp	D.oth(HiExp)	4	8	527	
5	1934.246	54.7465	24.7890	NA	22.0977	Alive	D.oth	5	9	150	
6	1934.246	44.3314	23.0437	NA	1.9563	Alive	HiExp	6	10	163	
685	1936.203	46.2877	25.0000	0.00000	16.2536	HiExp	D.lung(HiExp)	6	10	163	
	age1st	agein	ageout	agehi							
683	24.7206	47.9067	69.6794	30.27616							
5	29.9575	54.7465	76.8442	Inf							
6	21.2877	44.3314	62.5413	46.28770							
685	21.2877	44.3314	62.5413	46.28770							

```
> summary( nicC, scale=1000, timeScales=TRUE )
```

Transitions:

	To										
From	Alive	HiExp	D.oth	D.lung	D.lung(HiExp)	D.oth(HiExp)	Records:	Events:	Risk	t	
Alive	39	83	279	65	0	0	466	427		1	
HiExp	0	8	0	0	72	216	296	288			
Sum	39	91	279	65	72	216	762	715		1	

Timescales:

	time.scale	time.since
1	per	
2	age	
3	tfh	
4	tfHiExp	HiExp

With 6 different states it is quite difficult to get an overview of the transitions between states from the `summary()`. There there is function that gives a graphical display of the states showing the transitions between the states:

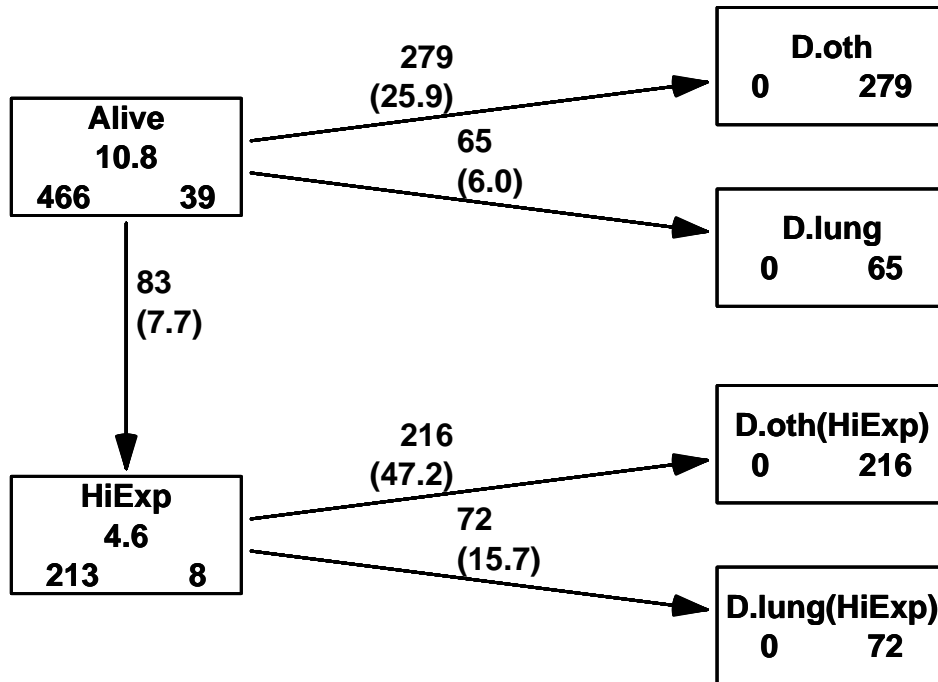


Figure 3: Transitions between states; the number in the middle of each box is the person-years (in 1000s — since `scale.Y=1000`), the numbers at the bottom of the boxes are the number that start, respectively end their follow-up in each state. The numbers on the arrows are the number of transitions and crude transition rates (the latter in events per 1000 PY).

The function `boxes.Lexis` has a zillion arguments to fine-tune the appearance of the display in terms of colors etc.

```

> boxes( nicC, boxpos = list(x=c(10,10,80,80,80,80),
+                               y=c(75,25,87,63,13,37)),
+       scale.Y = 1000,
+       show.BE = TRUE )

```