Simulation of multistate models with multiple timescales: `simLexis` in the `Epi` package

Bendix Carstensen   Steno Diabetes Center, Gentofte, Denmark
& Department of Biostatistics, University of Copenhagen
bxc@steno.dk
http://BendixCarstensen.com

Contents

Chapter 1

Using `simLexis`

1.1 Introduction

This vignette explains the machinery behind simulation of life histories through multistate models implemented in `simLexis`. In `simLexis` transition rates are allowed to depend on multiple time scales, including timescales defined as time since entry to a particular state (duration). This therefore also covers the case where time *at* entry into a state is an explanatory variable for the rates, since time at entry is merely time minus duration. Thus, the set-up here goes beyond Markov- and semi-Markov-models, and brings simulation based estimation of state-occupancy probabilities into the realm of realistic multistate models.

The basic idea is to simulate a new `Lexis` object [3, 1] as defined in the `Epi` package for R, based on 1) a multistate model defined by its states and the transition rates between them and 2) an initial population of individuals.

Thus the output will be a `Lexis` object describing the transitions of a predefined set of persons through a multistate model. Therefore, if persons are defined to be identical at start, then calculation of the probability of being in a particular state at a given time boils down to a simple enumeration of the fraction of the persons in the particular state at the given time. Bar of course the (binomial) simulation error, but this can be brought down by simulation a sufficiently large number of persons.

An observed `Lexis` object with follow-up of persons through a number of states will normally be the basis for estimation of transition rates between states, and thus will contain all information about covariates determining the occurrence rates, in particular the *timescales* [2]. Hence, the natural input to simulation from an estimated multistate model will typically be an object of the same structure as the originally observed. Since transitions and times are what is simulated, any values of `lex.Xst` and `lex.dur` in the input object will of course be ignored.

This first chapter of this vignette shows by an example how to use the function `simLexis` and display the results. The subsequent chapter discusses in more detail how the simulation machinery is implemented and is not needed for the practical use of `simLexis`.

1.2 `simLexis` in practice

This section is largely a commented walk-trough of the example from the help-page of `simLexis`, with a larger number of simulated persons in order to minimize the pure

1

simulation variation.

When we want to simulate transition times through a multistate model where transition rates may depend on time since entry to the current or a previous state, it is essential that we have a machinery to keep track of the transition time on *all* time scales, as well as a mechanism that can initiate a new time scale to 0 when a transition occurs to a state where we shall use time since entry as determinant of exit rates from that state. This is provided by `simLexis`.

1.2.1   Input for the simulation

Input for simulation of a single trajectory through a multistate model requires a representation of the *current status* of a person; the starting conditions. The object that we supply to the simulation function must contain information about all covariates and all timescales upon which transitions depend, and in particular which one(s) of the timescales that are defined as time since entry into a particular state. Hence, starting conditions should be represented as a `Lexis` object (where `lex.dur` and `lex.Xst` are ignored, since there is no follow-up yet), where the time scale information is in the attributes `time.scale` and `time.since` respectively.

Thus there are two main arguments to a function to simulate from a multistate model:

1. A `Lexis` object representing the initial states and covariates of the population to be simulated. This has to have the same structure as the original `Lexis` object representing the multistate model from which transition rates in the model were estimated. As noted above, the values for `lex.Xst` and `lex.dur` are not required (since these are the quantities that will be simulated).

2. A transition object, representing the transition intensities between states, which should be a list of lists of intensity representations. As an intensity representation we mean a function that for given a `Lexis` object that can be used to produce estimates of the transition intensities at a set of supplied time points since the state represented in the `Lexis` object.

   The names of the elements of the transition object (which are lists) will be names of the *transient* states, that is the states *from* which a transition can occur. The names of the elements of each of these lists are the names of states *to* which transitions can occur (which may be either transient or absorbing states).

   Hence, if the transition object is called `Tr` then `TR$A$B` (or `Tr[["A"]][["B"]]`) will represent the transition intensity from state `A` to the state `B`.

   The entries in the transition object can be either `glm` objects, representing Poisson models for the transitions, `coxph` objects representing an intensity model along one time scale, or simply a function that takes a `Lexis` object as input returns an estimated intensity for each row.

In addition to these two input items, there will be a couple of tuning parameters.

The output of the function will simply be a `Lexis` object with simulated transitions between states. This will be the basis for deriving sensible statistics from the `Lexis` object — see next section.

1.3 Setting up a `Lexis` object

As an example we will use the `DMlate` dataset from the `Epi` package; it is a dataset simulated to resemble a random sample of 10,000 patients from the Danish National Diabetes Register.

We start by loading the `Epi` package:

```
> options( width=90 )
> library( Epi )
> print( sessionInfo(), l=F )
R version 3.3.1 (2016-06-21)
Platform: x86_64-pc-linux-gnu (64-bit)
Running under: Ubuntu 14.04.5 LTS

attached base packages:
[1] utils     datasets  graphics  grDevices stats     methods   base

other attached packages:
[1] Epi_2.5

loaded via a namespace (and not attached):
 [1] cmprsk_2.2-7     MASS_7.3-45      Matrix_1.2-6      plyr_1.8.4
 [5] parallel_3.3.1   survival_2.39-5  etm_0.6-2         Rcpp_0.12.5
 [9] splines_3.3.1    grid_3.3.1       numDeriv_2014.2-1 lattice_0.20-33
```

First we load the diabetes data and set up a simple illness-death model:

```
> data(DMlate)
> dml <- Lexis( entry = list(Per=dodm, Age=dodm-dobth, DMdur=0 ),
+                exit = list(Per=dox),
+         exit.status = factor(!is.na(dodth),labels=c("DM","Dead")),
+                data = DMlate )
NOTE: entry.status has been set to "DM" for all.
```

This is just data for a simple survival model with states "DM" and "Dead". Now we cut the follow-up at insulin start, which for the majority of patients (T2D) is a clinical indicator of deterioration of disease regulation. We therefore also introduce a new timescale, and split the non-precursor states, so that we can address the question of ever having been on insulin:

```
> dmi <- cutLexis( dml, cut = dml$doins,
+                   pre = "DM",
+             new.state = "Ins",
+             new.scale = "t.Ins",
+           split.states = TRUE )
> summary( dmi )
Transitions:
     To
From    DM  Ins Dead Dead(Ins)  Records:  Events: Risk time:  Persons:
  DM  6157 1694 2048         0      9899     3742    45885.49      9899
  Ins    0 1340    0       451      1791      451     8387.77      1791
  Sum 6157 3034 2048       451     11690     4193    54273.27      9996
> str(dmi)
Classes 'Lexis' and 'data.frame':        11690 obs. of  15 variables:
 $ Per    : num  1999 2003 2005 2009 2009 ...
 $ Age    : num  58.7 64.1 86.3 44 75.8 ...
 $ DMdur  : num  0 0 0 0 0 0 0 0 0 0 ...
 $ t.Ins  : num  NA NA NA NA NA NA NA NA NA NA ...
 $ lex.dur: num  11.08 6.689 5.446 0.736 1.344 ...
 $ lex.Cst: Factor w/ 4 levels "DM","Ins","Dead",..: 1 1 1 1 1 1 1 1 1 1 ...
```

```
$ lex.Xst: Factor w/ 4 levels "DM","Ins","Dead",..: 1 1 1 1 1 3 1 1 3 1 ...
$ lex.id : int   1 2 3 4 5 6 7 8 9 10 ...
$ sex    : Factor w/ 2 levels "M","F": 2 1 2 2 1 2 1 1 2 1 ...
$ dobth  : num   1940 1939 1918 1965 1933 ...
$ dodm   : num   1999 2003 2005 2009 2009 ...
$ dodth  : num   NA NA NA NA NA ...
$ dooad  : num   NA 2007 NA NA NA ...
$ doins  : num   NA NA NA NA NA NA NA NA NA NA ...
$ dox    : num   2010 2010 2010 2010 2010 ...
- attr(*, "breaks")=List of 4
 ..$ Per  : NULL
 ..$ Age  : NULL
 ..$ DMdur: NULL
 ..$ t.Ins: NULL
- attr(*, "time.scales")= chr  "Per" "Age" "DMdur" "t.Ins"
- attr(*, "time.since")= chr   "" "" "" "Ins"
```

We can show how many person-years we have and show the number of transitions and transition rates (per 1000), using the `boxes.Lexis` function to display the states and the number of transitions between them:

```
> boxes( dmi, boxpos = list(x=c(20,20,80,80),
+                           y=c(80,20,80,20)),
+            scale.R = 1000, show.BE = TRUE )
```
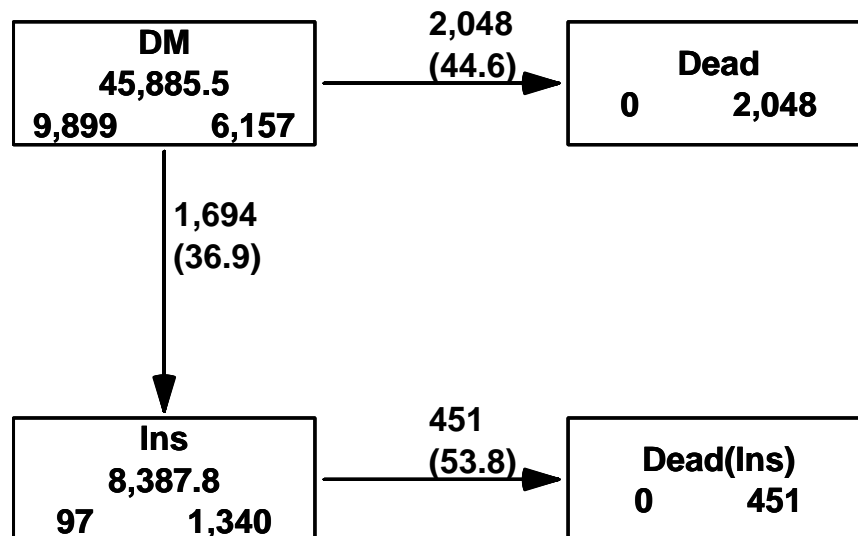

Figure 1.1: *Data overview for the* `dmi` *dataset. Numbers in the boxes are person-years and the number of persons who begin, resp. end their follow-up in each state, and numbers on the arrows are no. of transitions and rates (transition intensities) per 1000 PY.*

1.4   Analysis of rates

In the `Lexis` object (which is just a data frame) each person is represented by one record for each transient state he occupies, thus in this case either 1 or 2 records; those who have

a recorded time both without and with insulin have two records.

In order to be able to fit Poisson models with occurrence rates varying by the different time-scales, we split the follow-up in 6-month intervals for modeling:

```
> Si <- splitLexis( dmi, 0:30/2, "DMdur" )
> dim( Si )
[1] 115370     15

> print( subset( Si, lex.id==97 )[,1:10], digits=6 )

     lex.id     Per     Age   DMdur    t.Ins   lex.dur lex.Cst   lex.Xst sex   dobth
1105     97 1997.55 58.9268 0.00000       NA 0.5000000      DM        DM   F 1938.62
1106     97 1998.05 59.4268 0.50000       NA 0.5000000      DM        DM   F 1938.62
1107     97 1998.55 59.9268 1.00000       NA 0.5000000      DM        DM   F 1938.62
1108     97 1999.05 60.4268 1.50000       NA 0.5000000      DM        DM   F 1938.62
1109     97 1999.55 60.9268 2.00000       NA 0.1793292      DM       Ins   F 1938.62
1110     97 1999.72 61.1061 2.17933 0.000000 0.3206708     Ins       Ins   F 1938.62
1111     97 2000.05 61.4268 2.50000 0.320671 0.5000000     Ins       Ins   F 1938.62
1112     97 2000.55 61.9268 3.00000 0.820671 0.0116359     Ins Dead(Ins)   F 1938.62
```

Note that when we split the follow-up each person's follow up now consists of many records, each with the *current* values of the timescales at the start of the interval represented by the record. In the modelling we must necessarily assume that the rates are constant within each 6-month interval, but the *size* of these rates we model as smooth functions of the time scales (that is the values at the beginning of each interval).

The approach often used in epidemiology where one parameter is attached to each interval of time (or age) is not feasible when more than one time scale is used, because intervals are not classified the same way on all timescales.

We shall use natural splines (restricted cubic splines) for the analysis of rates, and hence we must allocate knots for the splines. This is done for each of the time-scales, and separately for the transition out of states "DM" and "Ins". For age, we place the knots so that the number of events is the same between each pair of knots, but only half of this beyond each of the boundary knots, whereas for the timescales `DMdur` and `tIns` where we have observation from a well-defined 0, we put knots at 0 and place the remaining knots so that the number of events is the same between them and beyond the last:

```
> nk <- 5
> ( ai.kn <- with( subset(Si,lex.Xst=="Ins"),
+                  quantile( Age+lex.dur  , probs=(1:nk-0.5)/nk ) ) )
      10%      30%      50%      70%      90%
23.75455 45.27279 56.62919 65.47851 77.50000

> ( ad.kn <- with( subset(Si,lex.Xst=="Dead"),
+                  quantile( Age+lex.dur  , probs=(1:nk-0.5)/nk ) ) )
      10%      30%      50%      70%      90%
61.91951 72.52731 78.43121 83.32348 90.15195

> ( di.kn <- with( subset(Si,lex.Xst=="Ins"),
+                  c(0,quantile( DMdur+lex.dur, probs=(1:(nk-1))/nk ) ) ) )
             20%      40%      60%      80%
0.000000 2.000000 4.500000 6.811225 9.500000

> ( dd.kn <- with( subset(Si,lex.Xst=="Dead"),
+                  c(0,quantile( DMdur+lex.dur, probs=(1:(nk-1))/nk ) ) ) )
              20%       40%       60%       80%
0.0000000 0.7687885 2.1327858 4.0465435 6.5232033

> ( ti.kn <- with( subset(Si,lex.Xst=="Dead(Ins)"),
+                  c(0,quantile( t.Ins+lex.dur, probs=(1:(nk-1))/nk ) ) ) )
```

```
                  20%       40%       60%       80%
0.0000000 0.3093771 1.1307324 2.5489391 4.9117043
```

We then fit Poisson models to transition rates, using the wrapper `Ns` from the `Epi` package
to simplify the specification of the rates:

```
> library( splines )
> DM.Ins <- glm( (lex.Xst=="Ins") ~ Ns( Age  , knots=ai.kn ) +
+                                   Ns( DMdur, knots=di.kn ) +
+                                   I(Per-2000) + sex,
+             family=poisson, offset=log(lex.dur),
+             data = subset(Si,lex.Cst=="DM") )
> DM.Dead <- glm( (lex.Xst=="Dead") ~ Ns( Age  , knots=ad.kn ) +
+                                   Ns( DMdur, knots=dd.kn ) +
+                                   I(Per-2000) + sex,
+             family=poisson, offset=log(lex.dur),
+             data = subset(Si,lex.Cst=="DM") )
> Ins.Dead <- glm( (lex.Xst=="Dead(Ins)") ~ Ns( Age  , knots=ad.kn ) +
+                                           Ns( DMdur, knots=dd.kn ) +
+                                           Ns( t.Ins, knots=ti.kn ) +
+                                           I(Per-2000) + sex,
+             family=poisson, offset=log(lex.dur),
+             data = subset(Si,lex.Cst=="Ins") )
```

Note the similarity of the code used to fit the three models, is is mainly redefining the
response variable ("to" state) and the subset of the data used ("from" state).

1.5 The mortality rates

This section discusses in some detail how to extract ad display the mortality rates from the
models fitted. But it is not necessary for understanding how to use `simLexis` in practice.

1.5.1 Proportionality of mortality rates

Note that we have fitted separate models for the three transitions, there is no assumption
of proportionality between the mortality rates from `DM` and `Ins`.

However, there is nothing that prevents us from testing this assumption; we can just fit a
model for the mortality rates in the entire data frame `Si`, and compare the deviance from
this with the sum of the deviances from the separate models:

```
> with( Si, table(lex.Cst) )
lex.Cst
      DM       Ins      Dead Dead(Ins)
   97039     18331         0         0
> All.Dead <- glm( (lex.Xst %in% c("Dead(Ins)","Dead")) ~
+                 Ns( Age  , knots=ad.kn ) +
+                 Ns( DMdur, knots=dd.kn ) +
+                 lex.Cst +
+                 I(Per-2000) + sex,
+                 family=poisson, offset=log(lex.dur),
+                 data = Si )
> round( ci.exp( All.Dead ), 3 )
                         exp(Est.)   2.5%   97.5%
(Intercept)                  0.049  0.043   0.056
Ns(Age, knots = ad.kn)1      4.120  3.479   4.879
Ns(Age, knots = ad.kn)2      4.652  4.054   5.338
Ns(Age, knots = ad.kn)3     15.460 13.575  17.608
```

```
Ns(Age, knots = ad.kn)4        7.529  6.711  8.447
Ns(DMdur, knots = dd.kn)1      0.520  0.429  0.629
Ns(DMdur, knots = dd.kn)2      0.707  0.622  0.803
Ns(DMdur, knots = dd.kn)3      0.319  0.238  0.428
Ns(DMdur, knots = dd.kn)4      0.829  0.742  0.926
lex.CstIns                     2.168  1.946  2.414
I(Per - 2000)                  0.965  0.954  0.977
sexF                           0.665  0.614  0.720
```

From the parameter values we would in a simple setting just claim that start of insulin-treatment was associated with a slightly more than doubling of mortality.

The model `All.dead` assumes that the age- and DM-duration effects on mortality in the "DM" and "Ins" states are the same, and moreover that there is no effect of insulin duration, but merely a mortality that is larger by a multiplicative constant not depending on insulin duration. The model `DM.Dead` has 8 parameters to describe the dependency on age and DM duration, the model `Ins.Dead` has 12 for the same plus the insulin duration (a natural spline with k knots gives $k-1$ parameters, and we chose $k=5$ above).

We can compare the fit of the proportional hazards model with the fit of the separate models for the two mortality rates, by adding up the deviances and d.f. from these:

```
> what <- c("null.deviance","df.null","deviance","df.residual")
> ( rD <- unlist(  DM.Dead[what] ) )
null.deviance          df.null       deviance   df.residual
    19957.95         97038.00       17849.90      97028.00
> ( rI <- unlist( Ins.Dead[what] ) )
null.deviance          df.null       deviance   df.residual
    4329.880        18330.000       3674.067     18316.000
> ( rA <- unlist( All.Dead[what] ) )
null.deviance          df.null       deviance   df.residual
    24300.15        115369.00       21608.79     115358.00
> round( c( dd <- rA-(rI+rD), "pVal"=1-pchisq(dd[3],dd[4]+1) ), 3 )
null.deviance          df.null       deviance   df.residual pVal.deviance
       12.314            1.000         84.822        14.000         0.000
```

Thus we see there is a substantial non-proportionality of mortality rates from the two states; but a test provides no clue whatsoever to the particular *shape* of the non-proportionality.

To this end, we shall explore the predicted mortalities under the two models quantitatively in more detail. Note that the reason that there is a difference in the null deviances (and a difference of 1 in the null d.f.) is that the null deviance of `All.Dead` refer to a model with a single intercept, that is a model with constant and *identical* mortality rates from the states "DM" and "Ins", whereas the null models for `DM.Dead` and `Ins.Dead` have constant but *different* mortality rates from the states "DM" and "Ins". This is however irrelevant for the comparison of the *residual* deviances.

1.5.2   How the mortality rates look

If we want to see how the mortality rates are modelled in `DM.Dead` and `Ins.Dead` in relation to `All.Dead`, we make a prediction of rates for say men diagnosed in different ages and going on insulin at different times after this. So we consider men diagnosed in ages 40, 50, 60 and 70, and who either never enter insulin treatment or do it 0, 2 or 5 years after diagnosis of DM.

To this end we create a prediction data frame where we have observation times from diagnosis and 12 years on (longer would not make sense as this is the extent of the data).

But we start by setting up an array to hold the predicted mortality rates, classified by diabetes duration, age at diabetes onset, time of insulin onset, and of course type of model. What we want to do is to plot the age-specific mortality rates for persons not on insulin, and for persons starting insulin at different times after DM. The mortality curves start at the age where the person gets diabetes and continues 12 years; for persons on insulin they start at the age when they initiate insulin.

```
> pr.rates <- NArray( list( DMdur = seq(0,12,0.1),
+                           DMage = 4:7*10,
+                           r.Ins = c(NA,0,2,5),
+                           model = c("DM/Ins","All"),
+                            what = c("rate","lo","hi") ) )
> str( pr.rates )
 logi [1:121, 1:4, 1:4, 1:2, 1:3] NA NA NA NA NA NA ...
 - attr(*, "dimnames")=List of 5
  ..$ DMdur: chr [1:121] "0" "0.1" "0.2" "0.3" ...
  ..$ DMage: chr [1:4] "40" "50" "60" "70"
  ..$ r.Ins: chr [1:4] NA "0" "2" "5"
  ..$ model: chr [1:2] "DM/Ins" "All"
  ..$ what : chr [1:3] "rate" "lo" "hi"
```

For convenience the `Epi` package contains a function that computes predicted (log-)rates with c.i. — it is merely a wrapper for `predict.glm`:

```
> ci.pred
function (obj, newdata, Exp = NULL, alpha = 0.05, df = Inf)
{
    if (!inherits(obj, "glm"))
        stop("Not usable for non-glm objects")
    zz <- predict(obj, newdata = newdata, se.fit = TRUE, type = "link")
    zz <- cbind(zz$fit, zz$se.fit) %*% ci.mat(alpha = alpha,
        df = df)
    if (missing(Exp)) {
        return(obj$family$linkinv(zz))
    }
    else {
        if (Exp) {
            return(exp(zz))
        }
        else if (!Exp)
            return(zz)
    }
}
<environment: namespace:Epi>
```

So set up the prediction data frame and modify it in loops over ages at onset and insulin onset. Note that we set `lex.dur` to 1000 in the prediction frame, so that we obtain rates in units of events per 1000 PY.

```
> nd <- data.frame( DMdur = as.numeric( dimnames(pr.rates)[[1]] ),
+                  lex.Cst = factor( 1, levels=1:4,
+                                    labels=levels(Si$lex.Cst) ),
+                      sex = factor( 1, levels=1:2, labels=c("M","F")),
+                  lex.dur = 1000 )
> for( ia in dimnames(pr.rates)[[2]] )
+    {
+ dnew <- transform( nd, Age = as.numeric(ia)+DMdur,
+                   Per = 1998+DMdur )
+ pr.rates[,ia,1,"DM/Ins",] <- ci.pred(  DM.Dead, newdata = dnew )
```

```
+ pr.rates[,ia,1,"All"   ,] <- ci.pred( All.Dead, newdata = dnew )
+ for( ii in dimnames(pr.rates)[[3]][-1] )
+    {
+ dnew = transform( dnew, lex.Cst = factor( 2, levels=1:4,
+                                     labels=levels(Si$lex.Cst) ),
+                     t.Ins = ifelse( (DMdur-as.numeric(ii)) >= 0,
+                                      DMdur-as.numeric(ii), NA ) )
+ pr.rates[,ia, ii ,"DM/Ins",] <- ci.pred( Ins.Dead, newdata = dnew )
+ pr.rates[,ia, ii ,"All"   ,] <- ci.pred( All.Dead, newdata = dnew )
+    }
+    }
```

So for each age at DM onset we make a plot of the mortality as function of current age both for those with no insulin treatment at those that start 1, 3 and 5 years after, thus 4 curves (with c.i.). These curves are replicated with a different color for the simplified model.

```
> par( mar=c(3,3,1,1), mgp=c(3,1,0)/1.6, las=1 )
> plot( NA, xlim=c(40,82), ylim=c(5,300), bty="n",
+       log="y", xlab="Age", ylab="Mortality rate per 1000 PY" )
> abline( v=seq(40,80,5), h=outer(1:9,10^(0:2),"*"), col=gray(0.8) )
> for( aa in 4:7*10 ) for( ii in 1:4 )
+    matlines( aa+as.numeric(dimnames(pr.rates)[[1]]),
+              cbind( pr.rates[,paste(aa),ii,"DM/Ins",],
+                      pr.rates[,paste(aa),ii,"All"   ,] ),
+              type="l", lty=1, lwd=c(3,1,1),
+              col=rep(c("red","limegreen"),each=3) )
```

From figure 1.2 we see that there is a substantial insulin-duration effect which is not accommodated by the simple model with only one time-dependent variable to describe the insulin effect. Note that the simple model (green curves) for those on insulin does not depend in insulin duration, and hence the mortality curves for those on insulin are just parallel to the mortality curves for those not on insulin, regardless of diabetes duration (or age) at the time of insulin initiation. This is the proportional hazards assumption. Thus the effect of insulin initiation is under-estimated for short duration of insulin and over-estimated for long duration of insulin.

This is the major discrepancy between the two models, and illustrates the importance of being able to accommodate different time scales, but there is also a declining overall insulin effect by age which is not accommodated by the proportional hazards approach.

Finally, this plot illustrates an important feature in reporting models with multiple timescales; all timescales must be represented in the predicted rates, only reporting the effect of one timescale, conditional on a fixed value of other timescales is misleading since all timescales by definition advance at the same pace. For example, the age-effect for a fixed value of insulin duration really is a misnomer since it does not correspond to any real person's follow-up, but to the mortality of persons in different ages but with the same duration of incuin use.

1.6   Input to the `simLexis` function

In order to simulate from the multistate model with the estimated transition rates, and get the follow-up of a hypothetical cohort, we must supply *both* the transition rates and the structure of the model *as well as* the initial cohort status to `simLexis`.
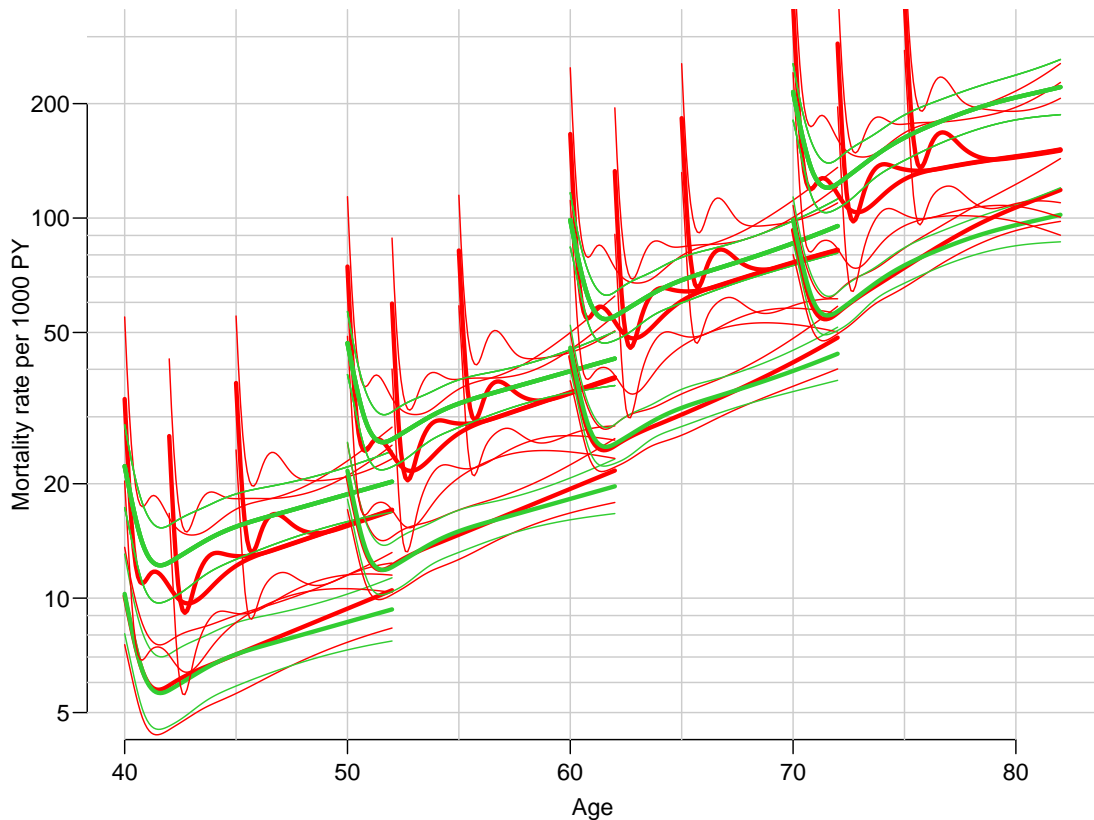
Figure 1.2: *Estimated mortality rates for male diabetes patients with no insulin (lower sets of curves) and insulin (upper curves), with DM onset in age 40, 50, 60 and 70. The red curves are from the models with separate age effects for persons with and without insulin, and a separate effect of insulin duration. The green curves are from the model with common age-effects and only a time-dependent effect of insulin, assuming no effect of insulin duration (the classical time-dependent variable approach). Hence the upper green curve is common for any time of insulin inception.*

1.6.1   The transition object

We first put the models into an object representing the transitions; note this is a list of lists, the latter having `glm` objects as elements:

```
> Tr <- list( "DM" = list( "Ins"      = DM.Ins,
+                          "Dead"     = DM.Dead  ),
+             "Ins" = list( "Dead(Ins)" = Ins.Dead ) )
```

Now we have the description of the rates and of the structure of the model. The `Tr` object defines the states and models for all transitions between them; the object `Tr$A$B` is the model for the transition intensity from state `A` to state `B`.

1.6.2   The initial cohort

We now define an initial `Lexis` object of persons with all relevant covariates defined. Note that we use `subset` to get a `Lexis` object, this conserves the `time.scale` and `time.since` attributes which is needed for the simulation (the usual "`[`" operator does not preserve these attributes when you select columns):

```
> str( Si[NULL,1:9] )
Classes 'Lexis' and 'data.frame':        0 obs. of  9 variables:
 $ lex.id : int
 $ Per    : num
 $ Age    : num
 $ DMdur  : num
 $ t.Ins  : num
 $ lex.dur: num
 $ lex.Cst: Factor w/ 4 levels "DM","Ins","Dead",..:
 $ lex.Xst: Factor w/ 4 levels "DM","Ins","Dead",..:
 $ sex    : Factor w/ 2 levels "M","F":
 - attr(*, "breaks")=List of 4
  ..$ Per  : NULL
  ..$ Age  : NULL
  ..$ DMdur: num  0 0.5 1 1.5 2 2.5 3 3.5 4 4.5 ...
  ..$ t.Ins: NULL
 - attr(*, "time.scales")= chr  "Per" "Age" "DMdur" "t.Ins"
 - attr(*, "time.since")= chr  "" "" "" "Ins"

> ini <- subset(Si,FALSE,select=1:9)
> str( ini )

Classes 'Lexis' and 'data.frame':        0 obs. of  9 variables:
 $ lex.id : int
 $ Per    : num
 $ Age    : num
 $ DMdur  : num
 $ t.Ins  : num
 $ lex.dur: num
 $ lex.Cst: Factor w/ 4 levels "DM","Ins","Dead",..:
 $ lex.Xst: Factor w/ 4 levels "DM","Ins","Dead",..:
 $ sex    : Factor w/ 2 levels "M","F":
 - attr(*, "breaks")=List of 4
  ..$ Per  : NULL
  ..$ Age  : NULL
  ..$ DMdur: num  0 0.5 1 1.5 2 2.5 3 3.5 4 4.5 ...
  ..$ t.Ins: NULL
 - attr(*, "time.scales")= chr  "Per" "Age" "DMdur" "t.Ins"
 - attr(*, "time.since")= chr  "" "" "" "Ins"

> ini <- subset(Si,select=1:9)[NULL,]
> str( ini )

Classes 'Lexis' and 'data.frame':        0 obs. of  9 variables:
 $ lex.id : int
 $ Per    : num
 $ Age    : num
 $ DMdur  : num
 $ t.Ins  : num
 $ lex.dur: num
 $ lex.Cst: Factor w/ 4 levels "DM","Ins","Dead",..:
 $ lex.Xst: Factor w/ 4 levels "DM","Ins","Dead",..:
 $ sex    : Factor w/ 2 levels "M","F":
 - attr(*, "breaks")=List of 4
  ..$ Per  : NULL
  ..$ Age  : NULL
  ..$ DMdur: num  0 0.5 1 1.5 2 2.5 3 3.5 4 4.5 ...
  ..$ t.Ins: NULL
 - attr(*, "time.scales")= chr  "Per" "Age" "DMdur" "t.Ins"
 - attr(*, "time.since")= chr  "" "" "" "Ins"
```

We now have an empty `Lexis` object with attributes reflecting the timescales in multistate model we want to simulate, so we must now enter some data to represent the persons whose follow-up we want to simulate through the model; we set up an initial dataset with one man and one woman:

```
> ini[1:2,"lex.id"] <- 1:2
> ini[1:2,"lex.Cst"] <- "DM"
> ini[1:2,"Per"] <- 1995
> ini[1:2,"Age"] <- 60
> ini[1:2,"DMdur"] <- 5
> ini[1:2,"sex"] <- c("M","F")
> ini
  lex.id  Per Age DMdur t.Ins lex.dur lex.Cst lex.Xst sex
1      1 1995  60     5    NA      NA      DM    <NA>   M
2      2 1995  60     5    NA      NA      DM    <NA>   F
```

1.7   Simulation of the follow-up

Now we simulate 1000 of each of these persons using the estimated model. The `t.range`
argument gives the times range where the integrated intensities (cumulative rates) are
evaluated and where linear interpolation is used when simulating transition times. Note
that this must be given in the same units as `lex.dur` in the `Lexis` object used for fitting
the models for the transitions.

```
> set.seed( 52381764 )
> Nsim <- 1000
> system.time( simL <- simLexis( Tr,
+                                ini,
+                            t.range = 12,
+                                 N = Nsim ) )
   user  system elapsed
  4.112   0.012   4.124
```

The result is a `Lexis` object — a data frame representing the simulated follow-up of 2000
persons (1000 identical men and 1000 identical women) according to the rates we estimated
from the original dataset.

```
> summary( simL, by="sex" )
$M

Transitions:
     To
From   DM Ins Dead Dead(Ins)  Records:  Events: Risk time:  Persons:
  DM  300 388  312         0      1000      700    7425.96      1000
  Ins   0 273    0       115       388      115    2125.24       388
  Sum 300 661  312       115      1388      815    9551.20      1000

$F

Transitions:
     To
From   DM Ins Dead Dead(Ins)  Records:  Events: Risk time:  Persons:
  DM  428 346  226         0      1000      572    8528.06      1000
  Ins   0 265    0        81       346       81    1712.67       346
  Sum 428 611  226        81      1346      653   10240.73      1000
```

1.7.1   Using other models for simulation

1.7.1.1   Proportional hazards Poisson model

We fitted a proportional mortality model `All.Dead` (which fitted worse than the other
two), this is a model for *both* the transition from "DM" to "Death" *and* from "Ins" to

"Dead(Ins)", assuming that they are proportional. But it can easily be used in the simulation set-up, because the state is embedded in the model via the term `lex.Cst`, which is updated during the simulation.

Simulation using this instead just requires that we supply a different transition object:

```
> Tr.p <- list( "DM" = list( "Ins"      = DM.Ins,
+                            "Dead"      = All.Dead ),
+               "Ins" = list( "Dead(Ins)" = All.Dead ) )
> system.time( simP <- simLexis( Tr.p,
+                                ini,
+                                t.range = 12,
+                                     N = Nsim ) )
   user  system elapsed
  4.245   0.032   4.278

> summary( simP, by="sex" )

$M

Transitions:
     To
From   DM Ins Dead Dead(Ins)  Records:  Events: Risk time:  Persons:
  DM  326 405  269         0      1000      674    7526.98      1000
  Ins   0 226    0       179       405      179    2018.83       405
  Sum 326 631  269       179      1405      853    9545.81      1000

$F

Transitions:
     To
From   DM Ins Dead Dead(Ins)  Records:  Events: Risk time:  Persons:
  DM  432 363  205         0      1000      568    8494.55      1000
  Ins   0 243    0       120       363      120    1824.00       363
  Sum 432 606  205       120      1363      688   10318.55      1000
```

1.7.1.2 Proportional hazards Cox model

A third possibility would be to replace the two-time scale proportional mortality model by a one-time-scale Cox-model, using diabetes duration as time scale, and age at diagnosis of DM as (fixed) covariate:

```
> library( survival )
> Cox.Dead <- coxph( Surv( DMdur, DMdur+lex.dur,
+                          lex.Xst %in% c("Dead(Ins)","Dead")) ~
+               Ns( Age-DMdur, knots=ad.kn ) +
+               I(lex.Cst=="Ins") +
+               I(Per-2000) + sex,
+            data = Si )
> round( ci.exp( Cox.Dead ), 3 )
                              exp(Est.)   2.5%  97.5%
Ns(Age - DMdur, knots = ad.kn)1   4.172  3.535   4.923
Ns(Age - DMdur, knots = ad.kn)2   4.503  3.825   5.301
Ns(Age - DMdur, knots = ad.kn)3  16.076 14.086  18.347
Ns(Age - DMdur, knots = ad.kn)4   7.478  6.500   8.604
I(lex.Cst == "Ins")TRUE           2.170  1.948   2.418
I(Per - 2000)                     0.966  0.954   0.977
sexF                              0.667  0.616   0.723

> round( ci.exp( All.Dead ), 3 )

                         exp(Est.)   2.5%  97.5%
(Intercept)                  0.049  0.043   0.056
Ns(Age, knots = ad.kn)1      4.120  3.479   4.879
```

```
Ns(Age, knots = ad.kn)2        4.652  4.054  5.338
Ns(Age, knots = ad.kn)3       15.460 13.575 17.608
Ns(Age, knots = ad.kn)4        7.529  6.711  8.447
Ns(DMdur, knots = dd.kn)1      0.520  0.429  0.629
Ns(DMdur, knots = dd.kn)2      0.707  0.622  0.803
Ns(DMdur, knots = dd.kn)3      0.319  0.238  0.428
Ns(DMdur, knots = dd.kn)4      0.829  0.742  0.926
lex.CstIns                     2.168  1.946  2.414
I(Per - 2000)                  0.965  0.954  0.977
sexF                           0.665  0.614  0.720
```

Note that in order for this model to be usable for simulation, it is necessary that we use the components of the `Lexis` object to specify the survival. Each record in the data frame `Si` represents follow up from `DMdur` to `DMdur+lex.dur`, so the model is a Cox model with diabetes duration as underlying timescale and age at diagnosis, `Age-DMdur`, as covariate.

Also note that we used `I(lex.Cst=="Ins")` instead of just `lex.Cst`, because `coxph` assigns design matrix columns to all levels of `lex.Cst`, also those not present in data, which would give `NA`s among the parameter estimates and `NA`s as mortality outcomes.

We see that the effect of insulin and the other covariates are pretty much the same as in the two-time-scale model. We can simulate from this model too; there is no restrictions on what type of model can be used for different transitions

```
> Tr.c <- list( "DM" = list( "Ins"      = Tr$DM$Ins,
+                            "Dead"      = Cox.Dead  ),
+              "Ins" = list( "Dead(Ins)" = Cox.Dead ) )
> system.time( simC <- simLexis( Tr.c,
+                                ini,
+                              t.range = 12,
+                                  N = Nsim ) )
   user  system elapsed
  5.097   0.004   5.110

> summary( simC, by="sex" )

$M

Transitions:
     To
From   DM Ins Dead Dead(Ins)  Records:  Events: Risk time:  Persons:
  DM  362 418  220         0      1000      638    7509.31      1000
  Ins   0 280    0       138       418      138    2143.04       418
  Sum 362 698  220       138      1418      776    9652.35      1000

$F

Transitions:
     To
From   DM Ins Dead Dead(Ins)  Records:  Events: Risk time:  Persons:
  DM  490 324  186         0      1000      510    8696.24      1000
  Ins   0 263    0        61       324       61    1699.43       324
  Sum 490 587  186        61      1324      571   10395.67      1000
```

1.8   Reporting the simulation results

We can now tabulate the number of persons in each state at a predefined set of times on a given time scale. Note that in order for this to be sensible, the `from` argument would normally be equal to the starting time for the simulated individuals.

```
> system.time(
+ nSt <- nState( subset(simL,sex=="M"),
+                at=seq(0,11,0.2), from=1995, time.scale="Per" ) )
   user  system elapsed
  0.166   0.000   0.166

> nSt[1:10,]

        State
when       DM   Ins Dead Dead(Ins)
  1995   1000    0    0         0
  1995.2  987   12    1         0
  1995.4  967   22   10         1
  1995.6  944   36   19         1
  1995.8  929   46   24         1
  1996    911   52   35         2
  1996.2  891   64   43         2
  1996.4  873   74   48         5
  1996.6  864   79   51         6
  1996.8  852   85   57         6
```

We see that as time goes by, the 5000 men slowly move away from the starting state (`DM`).

Based on this table (`nSt` is a table) we can now compute the fractions in each state, or, rather more relevant, the cumulative fraction across the states in some specified order, so that a plot of the stacked probabilities can be made, using either the default rather colorful layout, or a more minimalistic version (both in figure 1.3):

```
> pM <- pState( nSt, perm=c(1,2,4,3) )
> head( pM )

        State
when       DM    Ins Dead(Ins) Dead
  1995   1.000 1.000     1.000    1
  1995.2 0.987 0.999     0.999    1
  1995.4 0.967 0.989     0.990    1
  1995.6 0.944 0.980     0.981    1
  1995.8 0.929 0.975     0.976    1
  1996   0.911 0.963     0.965    1

> par( mfrow=c(1,2), mar=c(3,3,1,1), mgp=c(3,1,0)/1.6 )
> plot( pM )
> plot( pM, border="black", col="transparent", lwd=3 )
> text( rep(as.numeric(rownames(pM)[nrow(pM)-1]),ncol(pM)),
+       pM[nrow(pM),]-diff(c(0,pM[nrow(pM),]))/5,
+       colnames( pM ), adj=1 )
> box( col="white", lwd=3 )
> box()
```

A more useful set-up of the graph would include a more through annotation and sensible choice of colors, as seen in figure 1.4:

```
> clr <- c("limegreen","orange")
> # expand with a lighter version of the two chosen colors
> clx <- c( clr, rgb( t( col2rgb( clr[2:1] )*2 + rep(255,3) ) / 3, max=255 ) )
> par( mfrow=c(1,2), las=1, mar=c(3,3,4,2), mgp=c(3,1,0)/1.6 )
> # Men
> plot( pM, col=clx )
> lines( as.numeric(rownames(pM)), pM[,2], lwd=3 )
> mtext( "60 year old male, diagnosed 1990, aged 55", side=3, line=2.5, adj=0, col=gray(0.6) )
> mtext( "Survival curve", side=3, line=1.5, adj=0 )
> mtext( "DM, no insulin   DM, Insulin", side=3, line=0.5, adj=0, col=clr[1] )
> mtext( "DM, no insulin", side=3, line=0.5, adj=0, col=clr[2] )
> axis( side=4 )
> axis( side=4, at=1:19/20, labels=FALSE )
> axis( side=4, at=1:99/100, labels=FALSE, tcl=-0.3 )
```
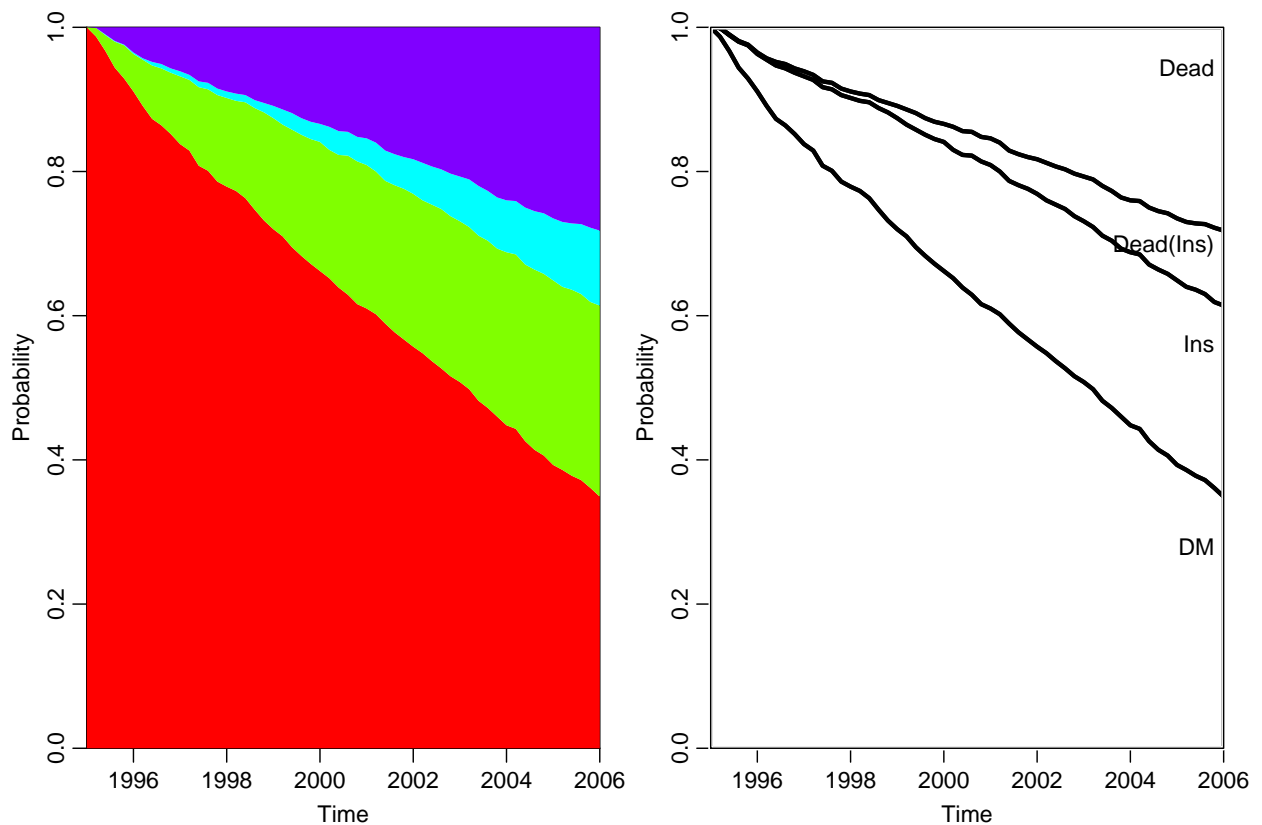
Figure 1.3: *Default layout of the* `plot.pState` *graph (left), and a version with the state probabilites as lines and annotation of states.*

```
> # Women
> pF <- pState( nState( subset(simL,sex=="F"),
+                        at=seq(0,11,0.2),
+                        from=1995,
+                        time.scale="Per" ),
+               perm=c(1,2,4,3) )
> plot( pF, col=clx )
> lines( as.numeric(rownames(pF)), pF[,2], lwd=3 )
> mtext( "60 year old female, diagnosed 1990, aged 55", side=3, line=2.5, adj=0, col=gray(0.6) )
> mtext( "Survival curve", side=3, line=1.5, adj=0 )
> mtext( "DM, no insulin   DM, Insulin", side=3, line=0.5, adj=0, col=clr[1] )
> mtext( "DM, no insulin", side=3, line=0.5, adj=0, col=clr[2] )
> axis( side=4 )
> axis( side=4, at=1:19/20, labels=FALSE )
> axis( side=4, at=1:99/100, labels=FALSE, tcl=-0.3 )
```

If we instead wanted to show the results on the age-scale, we would use age as timescale when constructing the probabilities; otherwise the code is pretty much the same as before (Figure 1.5):

```
> par( mfrow=c(1,2), las=1, mar=c(3,3,4,2), mgp=c(3,1,0)/1.6 )
> # Men
> pM <- pState( nState( subset(simL,sex=="M"),
+                        at=seq(0,11,0.2),
+                        from=60,
+                        time.scale="Age" ),
+               perm=c(1,2,4,3) )
> plot( pM, col=clx, xlab="Age" )
> lines( as.numeric(rownames(pM)), pM[,2], lwd=3 )
```
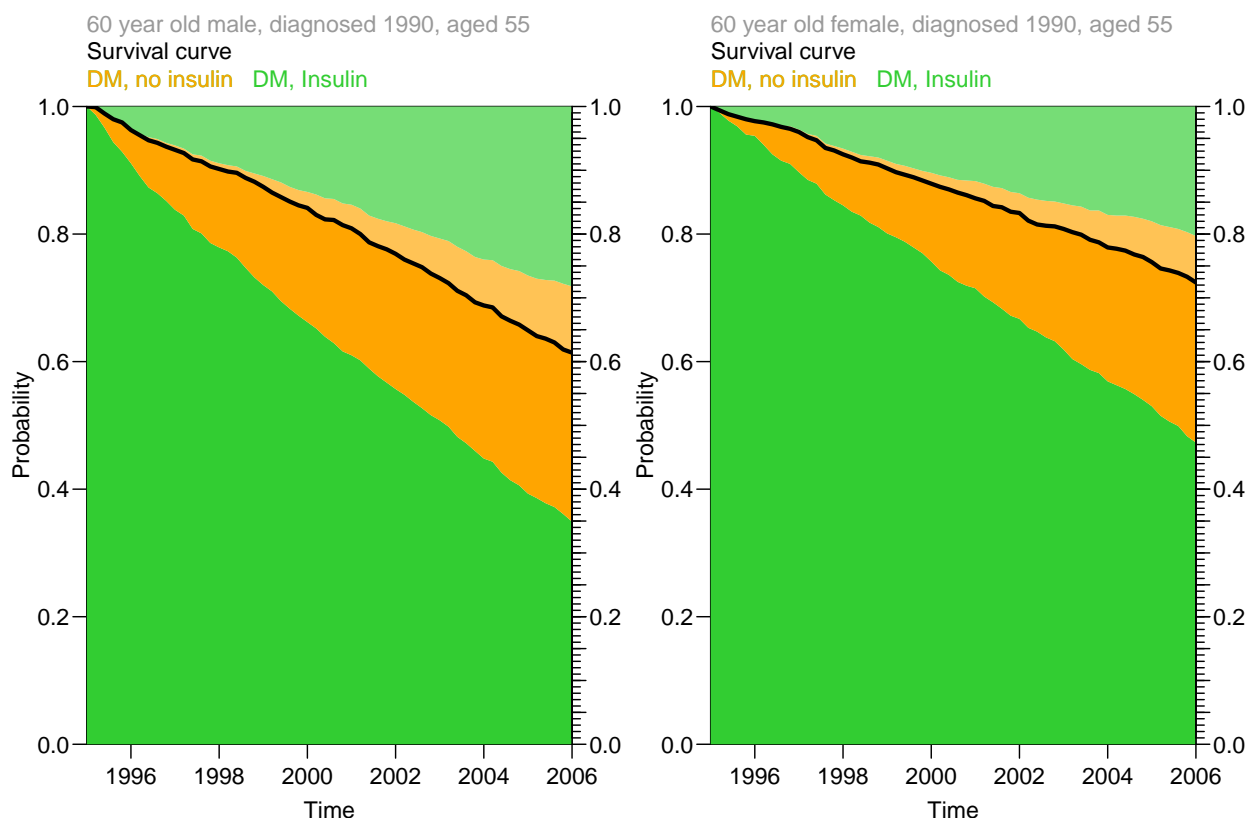
Figure 1.4: `plot.pState` *graphs where persons ever on insulin are given in orange and persons never on insulin in green, and the overall survival (dead over the line) as a black line.*

```
> mtext( "60 year old male, diagnosed 1990, aged 55", side=3, line=2.5, adj=0, col=gray(0.6) )
> mtext( "Survival curve", side=3, line=1.5, adj=0 )
> mtext( "DM, no insulin   DM, Insulin", side=3, line=0.5, adj=0, col=clr[1] )
> mtext( "DM, no insulin", side=3, line=0.5, adj=0, col=clr[2] )
> axis( side=4 )
> axis( side=4, at=1:19/20, labels=FALSE )
> axis( side=4, at=1:19/20, labels=FALSE, tcl=-0.4 )
> axis( side=4, at=1:99/100, labels=FALSE, tcl=-0.3 )
> # Women
> pF <- pState( nState( subset(simL,sex=="F"),
+                       at=seq(0,11,0.2),
+                       from=60,
+                       time.scale="Age" ),
+             perm=c(1,2,4,3) )
> plot( pF, col=clx, xlab="Age" )
> lines( as.numeric(rownames(pF)), pF[,2], lwd=3 )
> mtext( "60 year old female, diagnosed 1990, aged 55", side=3, line=2.5, adj=0, col=gray(0.6) )
> mtext( "Survival curve", side=3, line=1.5, adj=0 )
> mtext( "DM, no insulin   DM, Insulin", side=3, line=0.5, adj=0, col=clr[1] )
> mtext( "DM, no insulin", side=3, line=0.5, adj=0, col=clr[2] )
> axis( side=4 )
> axis( side=4, at=1:9/10, labels=FALSE )
> axis( side=4, at=1:19/20, labels=FALSE, tcl=-0.4 )
> axis( side=4, at=1:99/100, labels=FALSE, tcl=-0.3 )
```

Note the several statements with `axis(side=4,...`; they are nesessary to get the fine tick-marks in the right hand side of the plots that you will need in order to read off the probabilities at 2006 (or 71 years).
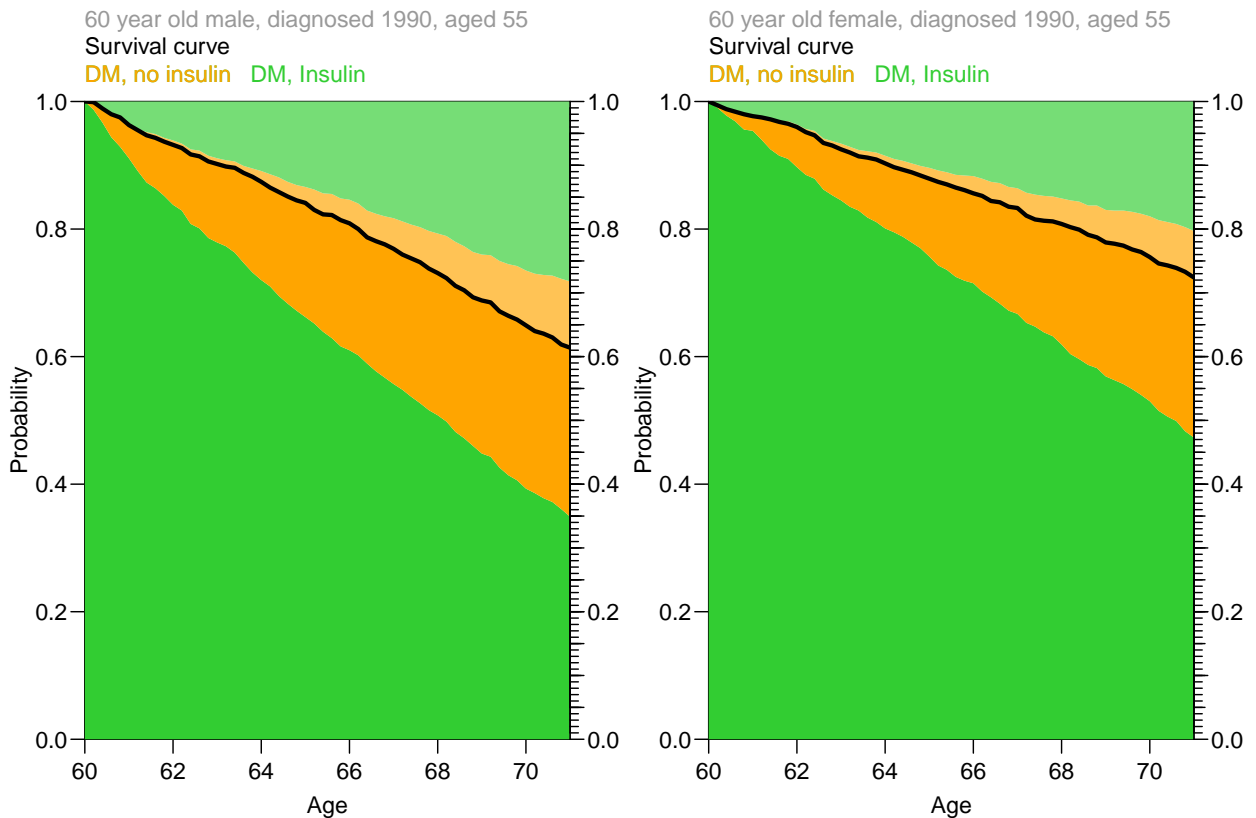
Figure 1.5: `plot.pState` _graphs where persons ever on insulin are given in orange and persons never on insulin in green, and the overall survival (dead over the line) as a black line._

1.8.1 Comparing predictions from different models

We have actually fitted different models for the transitions, and we have simulated Lexis objects from all three approaches, so we can plot these predictions on top of each other:

```
> PrM  <- pState( nState( subset(simP,sex=="M"),
+                         at=seq(0,11,0.2),
+                         from=60,
+                         time.scale="Age" ),
+                 perm=c(1,2,4,3) )
> PrF  <- pState( nState( subset(simP,sex=="F"),
+                         at=seq(0,11,0.2),
+                         from=60,
+                         time.scale="Age" ),
+                 perm=c(1,2,4,3) )
> CoxM <- pState( nState( subset(simC,sex=="M"),
+                         at=seq(0,11,0.2),
+                         from=60,
+                         time.scale="Age" ),
+                 perm=c(1,2,4,3) )
> CoxF <- pState( nState( subset(simC,sex=="F"),
+                         at=seq(0,11,0.2),
+                         from=60,
+                         time.scale="Age" ),
+                 perm=c(1,2,4,3) )
> par( mfrow=c(1,2), mar=c(3,3,1,1), mgp=c(3,1,0)/1.6 )
>  plot(  pM, border="black", col="transparent", lwd=3 )
> lines(  PrM, border="blue" , col="transparent", lwd=3 )
```

```
> lines( CoxM, border="red"  , col="transparent", lwd=3 )
> text( 60.5, 0.05, "M" )
> box( lwd=3 )
>  plot(   pF, border="black", col="transparent", lwd=3 )
> lines(  PrF, border="blue" , col="transparent", lwd=3 )
> lines( CoxF, border="red"  , col="transparent", lwd=3 )
> text( 60.5, 0.05, "F" )
> box( lwd=3 )
```
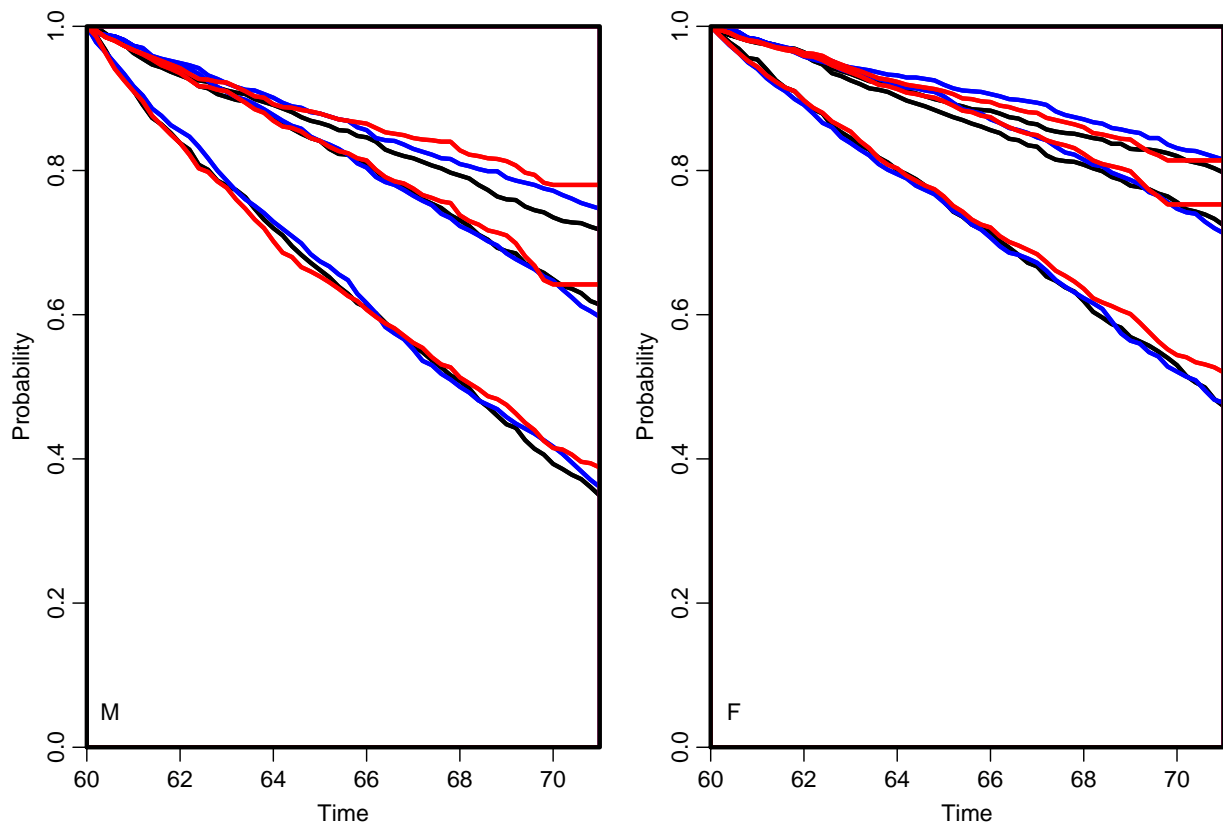

Figure 1.6: *Comparison of the simulated state occupancy probabilities using separate Poisson models for the mortality rates with and without insulin (black) and using proportional hazards Poisson models (blue) and Cox-models with diabetes duration as timescale and age at diabetes diagnosis as covariate (red).*

From figure 1.6 it is clear that the two proportional hazards models (blue and red curves) produce pretty much the same estimates of the state occupancy probabilites over time, but also that they relative to the model with separately estimated transition intensities overestimates the probability of being alive without insulin and underestimates the probabilities of being dead without insulin. However both the overall survival, and the fraction of persons on insulin are quite well in agreement with the more elaborate model. Thus the proportional hazards models overestimate the relative mortality of the insulin treated diabetes patients relative to the non-insulin treated.

Interestingly, we also see a bump in th estimated probabilities from the Cox-model based model, but this is entirely an artifact that comes from the estimation method for the baseline hazard of the Cox-model that lets the (cumulative) hazard have large jumps at event times where the risk set is small. So also here it shows up that an assumption of continuous underlying hazards leads to more credible estimates.

Chapter 2

Simulation of transitions in multistate models

2.1 Theory

Suppose that the rate functions for the transitions out of the current state to, say, 3 different states are λ_1, λ_2 and λ_3, and the corresponding cumulative rates are Λ_1, Λ_2 and Λ_3, and we want to simulate an exit time and an exit state (that is either 1, 2 or 3). This can be done in two slightly different ways:

1. First time, then state:

   (a) Compute the survival function, $S(t) = \exp\big(-\Lambda_1(t) - \Lambda_2(t) - \Lambda_3(t)\big)$

   (b) Simulate a random U(0,1) variate, u, say.

   (c) The simulated exit time is then the solution t_u to the equation
   $S(t_u) = u \quad \Leftrightarrow \quad \sum_j \Lambda_j(t_u) = -\log(u).$

   (d) A simulated transition at t_u is then found by simulating a random draw from the multinomial distribution with probabilities $p_i = \lambda_i(t_u)/\sum_j \lambda_j(t_u)$.

2. Separate cumulative incidences:

   (a) Simulate 3 independent U(0,1) random variates u_1, u_2 and u_3.

   (b) Solve the equations $\Lambda_i(t_i) = -\log(u_i), i = 1, 2, 3$ and get (t_1, t_2, t_3).

   (c) The simulated survival time is then $\min(t_1, t_2, t_3)$, and the simulated transition is to the state corresponding to this, that is $k \in \{1, 2, 3\}$, where $t_k = \min(t_1, t_2, t_3)$

The intuitive argument is that with three possible transition there are 3 independent processes running, but only the first transition is observed. The latter approach is used in the implementation in `simLexis`.

The formal argument for the equality of the two approaches goes as follows:

1. Equality of the transition times:

(a) In the first approach we simulate from a distribution with cumulative rate $\Lambda_1(t) + \Lambda_2(t) + \Lambda_3(t)$, hence from a distribution with survival function:

$$S(t) = \exp\big(-(\Lambda_1(t) + \Lambda_2(t) + \Lambda_3(t))\big)$$
$$= \exp\big(-\Lambda_1(t)\big) \times \exp\big(-\Lambda_2(t)\big) \times \exp\big(-\Lambda_3(t)\big)$$

(b) In the second approach we choose the smallest of three independent survival times, with survival functions $\exp(-\Lambda_i), i = 1, 2, 3$. Now, the survival function for the minimum of three independent survival times is:

$$S_{\min}(t) = \mathrm{P}\left\{\min(t_1, t_2, t_3) > t\right\}$$
$$= \mathrm{P}\left\{t_1 > t\right\} \times \mathrm{P}\left\{t_2 > t\right\} \times \mathrm{P}\left\{t_3 > t\right\}$$
$$= \exp\big(-\Lambda_1(t)\big) \times \exp\big(-\Lambda_2(t)\big) \times \exp\big(-\Lambda_3(t)\big)$$

which is the same survival function as derived above.

2. Type of transition:

(a) In the first instance the probability of a transition to state i, conditional on the transition time being t, is as known from standard probability theory: $\lambda_i(t)/\big(\lambda_1(t) + \lambda_2(t) + \lambda_3(t)\big)$.

(b) In the second approach we choose the transition corresponding to the the smallest of the transition times. So when we condition on the event that a transition takes place at time t, we have to show that the conditional probability that the smallest of the three simulated transition times was actually the ith, is as above.

But conditional on *survival* till t, the probabilities that events of type $1, 2, 3$ takes place in the interval $(t, t + \mathrm{d}t)$ are $\lambda_1(t)\,\mathrm{d}t$, $\lambda_2(t)\,\mathrm{d}t$ and $\lambda_1(t)\,\mathrm{d}t$, respectively (assuming that the probability of more than one event in the interval of length $\mathrm{d}t$ is 0). Hence the conditional probabilities *given a transition time* in $(t, t + \mathrm{d}t)$ is:

$$\frac{\lambda_i(t)\,\mathrm{d}t}{\lambda_1(t)\,\mathrm{d}t + \lambda_2(t)\,\mathrm{d}t + \lambda_3(t)\,\mathrm{d}t} = \frac{\lambda_i(t)}{\lambda_1(t) + \lambda_2(t) + \lambda_3(t)}$$

— exactly as above.

2.2   Components of `simLexis`

This section explains the actually existing code for `simLexis`, as it is in the current version of `Epi`. The function `simLexis` takes a `Lexis` object as input. This defines the initial state(s) and times of the start for a number of persons. Since the purpose is to simulate a history through the estimated multistate model, the input values of the outcome variables `lex.Xst` and `lex.dur` are ignored — the aim is to simulate values for them.

Note however that the attribute `time.since` must be present in the object. This is used for initializing timescales defined as time since entry into a particular state, it is a character vector of the same length as the `time.scales` attribute, with value equal to a state name if the corresponding time scale is defined as time since entry into that state. In this example the 4th timescale is time since entry into the "Ins" state, and hence:

```
> cbind(
+ attr( ini, "time.scale" ),
+ attr( ini, "time.since" ) )
     [,1]    [,2]
[1,] "Per"   ""
[2,] "Age"   ""
[3,] "DMdur" ""
[4,] "t.Ins" "Ins"
```

Lexis objects will have this attribute set for time scales created using **cutLexis**.

The other necessary argument is a transition object **Tr**, which is a list of lists. The elements of the lists should be **glm** objects derived by fitting Poisson models to a **Lexis** object representing follow-up data in a multistate model. It is assumed (but not checked) that timescales enter in the model via the timescales of the **Lexis** object. Formally, there are no assumptions about how **lex.dur** enters in the model.

Optional arguments are **t.range**, **n.int** or **time.pts**, specifying the times after entry at which the cumulative rates will be computed (the maximum of which will be taken as the censoring time), and **N** a scalar or numerical vector of the number of persons with a given initial state each record of the **init** object should represent.

The central part of the functions uses a **do.call** / **lapply** / **split** construction to do simulations for different initial states. This is the construction in the middle that calls **simX**. **simLexis** also calls **get.next** which is further detailed below.

```
> simLexis
function (Tr, init, N = 1, lex.id, t.range = 20, n.int = 101,
    time.pts = seq(0, t.range, length.out = n.int))
{
    if (time.pts[1] != 0)
        stop("First time point must be 0, time.pts[1:3]= ", time.pts[1:3])
    if (!missing(N)) {
        if (length(N) == 1)
            init <- init[rep(1:nrow(init), each = N), ]
        else init <- init[rep(1:nrow(init), N), ]
    }
    if (!missing(lex.id)) {
        if (length(lex.id) == nrow(init))
            init$lex.id <- lex.id
        else init$lex.id <- 1:nrow(init)
    }
    else init$lex.id <- 1:nrow(init)
    if (is.null(tS <- attr(init, "time.scales")))
        stop("No time.scales attribute for init")
    if (is.null(tF <- attr(init, "time.since"))) {
        attr(init, "time.since") <- tF <- rep("", tS)
        warning("'time.since' attribute set to blanks")
    }
    np <- length(time.pts)
    tr.st <- names(Tr)
    sf <- NULL
    nxt <- init[init$lex.Cst %in% tr.st, ]
    if (nrow(nxt) < nrow(init)) {
        tt <- table(init$lex.Cst)
        tt <- tt[tt > 0]
        nt <- length(tt)
        warning("\nSome initiators start in a absorbing state\n",
            "Initiator states represented are: ", paste(rbind(names(tt),
                rep(":", nt), paste(tt), rep(" ", nt)), collapse = ""),
            "\n", "Transient states are: ", paste(names(Tr),
                coll = " "))
        if (nrow(nxt) == 0)
            stop("\nNo initiators in transient states!")
```

```
    }
    while (nrow(nxt) > 0) {
        nx <- do.call(rbind.data.frame, lapply(split(nxt, nxt$lex.Cst),
            simX, Tr, time.pts, tS))
        sf <- rbind.data.frame(sf, nx)
        nxt <- get.next(nx, tr.st, tS, tF)
    }
    sf$lex.Xst <- factor(sf$lex.Xst, levels = levels(sf$lex.Cst))
    sf$lex.Xst[is.na(sf$lex.Xst)] <- sf$lex.Cst[is.na(sf$lex.Xst)]
    nord <- match(c("lex.id", tS, "lex.dur", "lex.Cst", "lex.Xst"),
        names(sf))
    noth <- setdiff(1:ncol(sf), nord)
    sf <- sf[order(sf$lex.id, sf[, tS[1]]), c(nord, noth)]
    rownames(sf) <- NULL
    attr(sf, "time.scales") <- tS
    attr(sf, "time.since") <- tF
    chop.lex(sf, tS, max(time.pts))
}
<environment: namespace:Epi>
```

2.2.1   `simX`

`simX` is called by `simLexis` and simulates transition-times and -types for a set of patients
assumed to be in the same state. It is called from `simLexis` with a data frame as
argument, uses the state in `lex.Cst` to select the relevant component of `Tr` and compute
predicted cumulative intensities for all states reachable from this state.

   Note that it is here the switch between `glm`, `coxph` and objects of class `function` is made.

   The dataset on which this prediction is done has `length(time.pts)` rows per person.

```
> Epi:::simX
function (nd, Tr, time.pts, tS)
{
    np <- length(time.pts)
    nr <- nrow(nd)
    if (nr == 0)
        return(NULL)
    cst <- as.character(unique(nd$lex.Cst))
    if (length(cst) > 1)
        stop("More than one lex.Cst present:\n", cst, "\n")
    prfrm <- nd[rep(1:nr, each = np), ]
    prfrm[, tS] <- prfrm[, tS] + rep(time.pts, nr)
    prfrm$lex.dur <- il <- min(diff(time.pts))
    prfrp <- prfrm
    prfrp[, tS] <- prfrp[, tS] + il/2
    rt <- data.frame(lex.id = prfrm$lex.id)
    for (i in 1:length(Tr[[cst]])) {
        if (inherits(Tr[[cst]][[i]], "glm"))
            rt <- cbind(rt, predict(Tr[[cst]][[i]], type = "response",
                newdata = prfrp))
        else if (inherits(Tr[[cst]][[i]], "coxph"))
            rt <- cbind(rt, predict(Tr[[cst]][[i]], type = "expected",
                newdata = prfrm))
        else if (is.function(Tr[[cst]][[i]]))
            rt <- cbind(rt, Tr[[cst]][[i]](prfrm))
        else stop("Invalid object supplied as transition, elements of the list must be either:\n",
            "- a glm(poisson) object fitted to a Lexis object\n",
            "- a coxph object fitted to a Lexis object\n", "- a function that takes a Lexis object a
            "  average rates for each record in the same units as lex.dur.")
    }
    names(rt)[-1] <- names(Tr[[cst]])
    xx <- match(c("lex.dur", "lex.Xst"), names(nd))
    if (any(!is.na(xx)))
```

```
        nd <- nd[, -xx[!is.na(xx)]]
    merge(nd, do.call(rbind, lapply(split(rt, rt$lex.id), sim1,
        time.pts)), by = "lex.id")
}
<environment: namespace:Epi>
```

As we see, `simX` calls `sim1` which simulates the transition for one person.

2.2.2   sim1

The predicted cumulative intensities are fed, person by person, to `sim1` — again via a `do.call` / `lapply` / `split` construction — and the resulting time and state is appended to the `nd` data frame. This way we have simulated *one* transition (time and state) for each person:

```
> Epi:::sim1
function (rt, time.pts)
{
    ci <- apply(rbind(0, rt[, -1, drop = FALSE]), 2, cumsum)[1:nrow(rt),
        , drop = FALSE]
    tt <- uu <- -log(runif(ncol(ci)))
    for (i in 1:ncol(ci)) tt[i] <- lint(ci[, i], time.pts, uu[i])
    data.frame(lex.id = rt[1, 1], lex.dur = min(tt, na.rm = TRUE),
        lex.Xst = factor(if (min(tt) < max(time.pts))
            colnames(ci)[tt == min(tt)]
        else NA))
}
<environment: namespace:Epi>
```

The `sim1` function uses `lint` to do linear interpolation.

2.2.3   lint

We do not use `approx` to do the linear interpolation, because this function does not do the right thing if the cumulative incidences (`ci`) are constant across a number of times. Therefore we have a custom made linear interpolator that does the interpolation exploiting the fact the the `ci` is non-decreasing and `tt` is strictly monotonously increasing:

```
> Epi:::lint
function (ci, tt, u)
{
    if (any(diff(ci) < 0) | any(diff(tt) < 0))
        stop("Non-icreasing arguments")
    c.u <- min(c(ci[ci > u], max(ci)))
    c.l <- max(c(ci[ci < u], min(ci)))
    t.u <- min(c(tt[ci > u], max(tt)))
    t.l <- max(c(tt[ci < u], min(tt)))
    ifelse(c.u == c.l, t.l, t.l + (u - c.l)/(c.u - c.l) * (t.u -
        t.l))
}
<environment: namespace:Epi>
```

2.2.4   get.next

We must repeat the simulation operation on those that have a simulated entry to a transient state, and also make sure that any time scales defined as time since entry to one

of these states be initialized to 0 before a call to `simX` is made for these persons. This accomplished by the function `get.next`:

```
> Epi:::get.next
function (sf, tr.st, tS, tF)
{
    if (nrow(sf) == 0)
        return(sf)
    nxt <- sf[sf$lex.Xst %in% tr.st, ]
    if (nrow(nxt) == 0)
        return(nxt)
    nxt[, tS] <- nxt[, tS] + nxt$lex.dur
    wh <- tF
    for (i in 1:length(wh)) if (wh[i] != "")
        nxt[nxt$lex.Xst == wh[i], tS[i]] <- 0
    nxt$lex.Cst <- nxt$lex.Xst
    return(nxt)
}
<environment: namespace:Epi>
```

2.2.5 `chop.lex`

The operation so far has censored individuals `max(time.pts)` after *each* new entry to a transient state. In order to groom the output data we use `chop.lex` to censor all persons at the same designated time after *initial* entry.

```
> Epi:::chop.lex
function (obj, tS, cens)
{
    zz <- entry(obj, 1, by.id = TRUE)
    ww <- merge(obj, data.frame(lex.id = as.numeric(names(zz)),
        cens = zz + cens))
    ww <- ww[ww[, tS[1]] < ww$cens, ]
    x.dur <- pmin(ww$lex.dur, ww[, "cens"] - ww[, tS[1]])
    ww$lex.Xst[x.dur < ww$lex.dur] <- ww$lex.Cst[x.dur < ww$lex.dur]
    ww$lex.dur <- pmin(x.dur, ww$lex.dur)
    ww
}
<environment: namespace:Epi>
```

2.3   Probabilities from simulated `Lexis` objects

Once we have simulated a Lexis object we will of course want to use it for estimating probabilities, so basically we will want to enumerate the number of persons in each state at a pre-specified set of time points.

2.3.1 `nState`

Since we are dealing with multistate model with potentially multiple time scales, it is necessary to define the timescale (`time.scale`), the starting point on this timescale (`from`) and the points after this where we compute the number of occupants in each state, (`at`).

```
> nState
```

```
function (obj, at, from, time.scale = 1)
{
    tS <- check.time.scale(obj, time.scale)
    TT <- tmat(obj)
    absorb <- rownames(TT)[apply(!is.na(TT), 1, sum) == 0]
    transient <- setdiff(rownames(TT), absorb)
    tab.frm <- obj[rep(1:nrow(obj), each = length(at)), c(tS,
        "lex.dur", "lex.Cst", "lex.Xst")]
    tab.frm$when <- rep(at, nrow(obj)) + from
    tab.tr <- tab.frm[tab.frm[, tS] <= tab.frm$when & tab.frm[,
        tS] + tab.frm$lex.dur > tab.frm$when, ]
    tab.tr$State <- tab.tr$lex.Cst
    tab.ab <- tab.frm[tab.frm[, tS] + tab.frm$lex.dur <= tab.frm$when &
        tab.frm$lex.Xst %in% absorb, ]
    tab.ab$State <- tab.ab$lex.Xst
    with(rbind(tab.ab, tab.tr), table(when, State))
}
<environment: namespace:Epi>
```

2.3.2   `pState`, `plot.pState` and `lines.pState`

In order to plot probabilities of state-occupancy it is useful to compute cumulative probabilities across states in any given order; this is done by the function `pState`, which returns a matrix of class `pState`:

```
> pState
function (nSt, perm = 1:ncol(nSt))
{
    tt <- t(apply(nSt[, perm], 1, cumsum))
    tt <- sweep(tt, 1, tt[, ncol(tt)], "/")
    class(tt) <- c("pState", "matrix")
    tt
}
<environment: namespace:Epi>
```

There is also a `plot` and `lines` method for the resulting `pState` objects:

```
> plot.pState
function (x, col = rainbow(ncol(x)), border = "transparent",
    xlab = "Time", ylim = 0:1, ylab = "Probability", ...)
{
    matplot(as.numeric(rownames(x)), x, type = "n", ylim = ylim,
        yaxs = "i", xaxs = "i", xlab = xlab, ylab = ylab, ...)
    lines.pState(x, col = col, border = border, ...)
}
<environment: namespace:Epi>
> lines.pState
function (x, col = rainbow(ncol(x)), border = "transparent",
    ...)
{
    nc <- ncol(x)
    col <- rep(col, nc)[1:nc]
    border <- rep(border, nc)[1:nc]
    pSt <- cbind(0, x)
    for (i in 2:ncol(pSt)) {
        polygon(c(as.numeric(rownames(pSt)), rev(as.numeric(rownames(pSt)))),
            c(pSt[, i], rev(pSt[, i - 1])), col = col[i - 1],
            border = border[i - 1], ...)
    }
}
<environment: namespace:Epi>
```

Bibliography

[1] B. Carstensen and M. Plummer. Using Lexis objects for multi-state models in R. *Journal of Statistical Software*, 38(6):1–18, 1 2011.

[2] S. Iacobelli and B. Carstensen. Multiple time scales in multi-state models. *Stat Med*, 32(30):5315–5327, Dec 2013.

[3] M. Plummer and B. Carstensen. Lexis: An R class for epidemiological studies with long-term follow-up. *Journal of Statistical Software*, 38(5):1–12, 1 2011.